


**DataStax**

# ACID transactions in Apache Cassandra®

Cassandra Day - Version 0.1.0 BETA



# Background

If you haven't been on the dev list for a while

**Cassandra is great, but...**

**I need transactions**

# Lightweight Transactions

Cassandra 2.0 - 2013

## Pros

- Paxos - Established protocol
- CAS functionality
- Guarantees exclusive operation

## Cons

- Only one partition
- Serialized operations - nope
- Was slower (Fixed in V2)

```
INSERT INTO cycling.cyclist_name (id, lastname, firstname)
VALUES (4647f6d3-7bd2-4085-8d6c-1229351b5498, 'KNETEMANN', 'Roxxane')
IF NOT EXISTS;
```

```
UPDATE cycling.cyclist_name
SET firstname = 'Roxane'
WHERE id = 4647f6d3-7bd2-4085-8d6c-1229351b5498
IF firstname = 'Roxxane';
```

# CEP-15: General Purpose Transactions

## Goals

- General purpose transactions (may operate over any keys in the database at once)
- Strict-serializable isolation
- Optimal latency: one wide area round-trip for all transactions under normal conditions
- Optimal failure tolerance: latency and performance should be unaffected by any minority of replica failures
- Scalability: there should be no bottleneck introduced
- Should have no intrinsic limit to transaction complexity
- Must work on commodity hardware
- Must support live migration from Paxos

TL;DR ACID Transactions in Cassandra

DS

# Developer Impact


The people who count

# Cassandra Relationship With Developers

Easy



Hard



- Scale
- Resilience
- Distribution

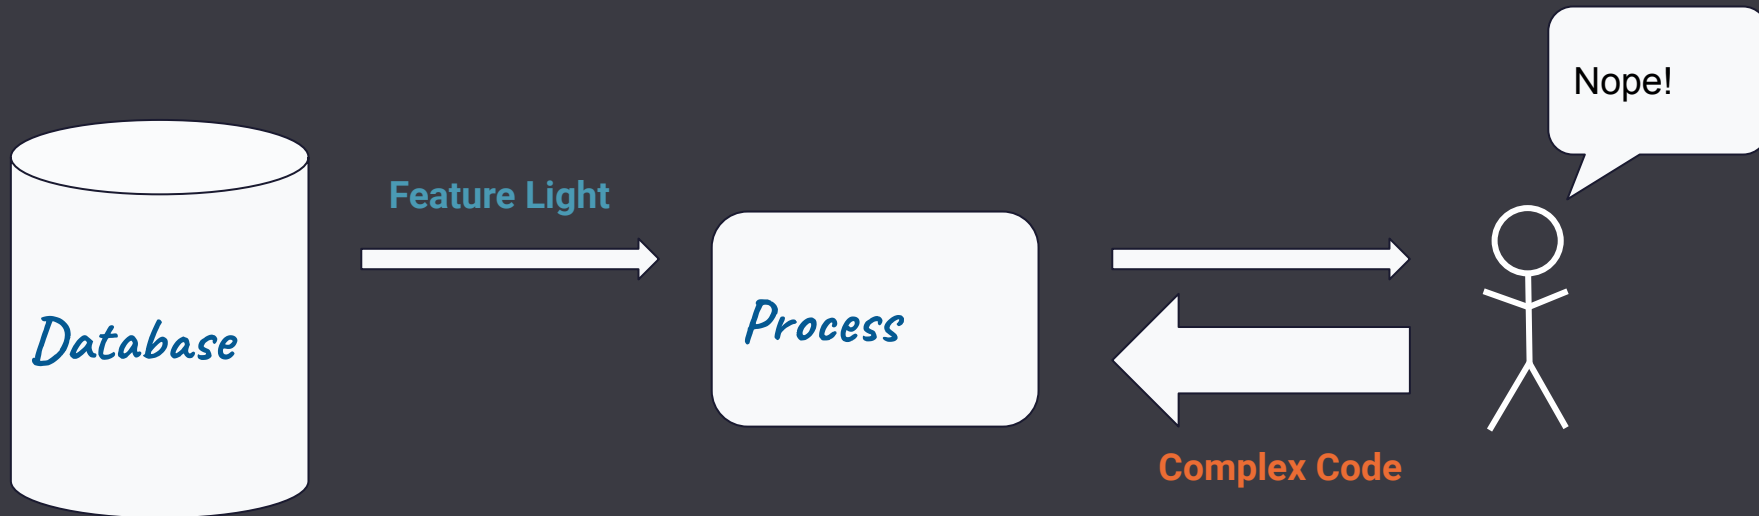
Built-in

- Atomicity
- Serialized changes
- Complicated State

Data Modeling  
& Code



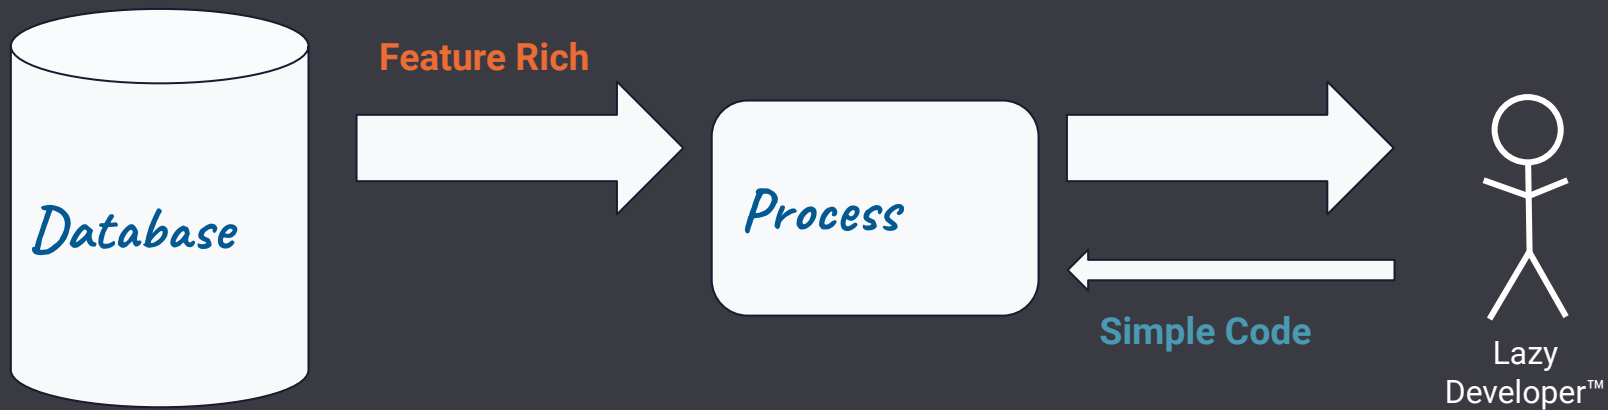
# Developer Experience



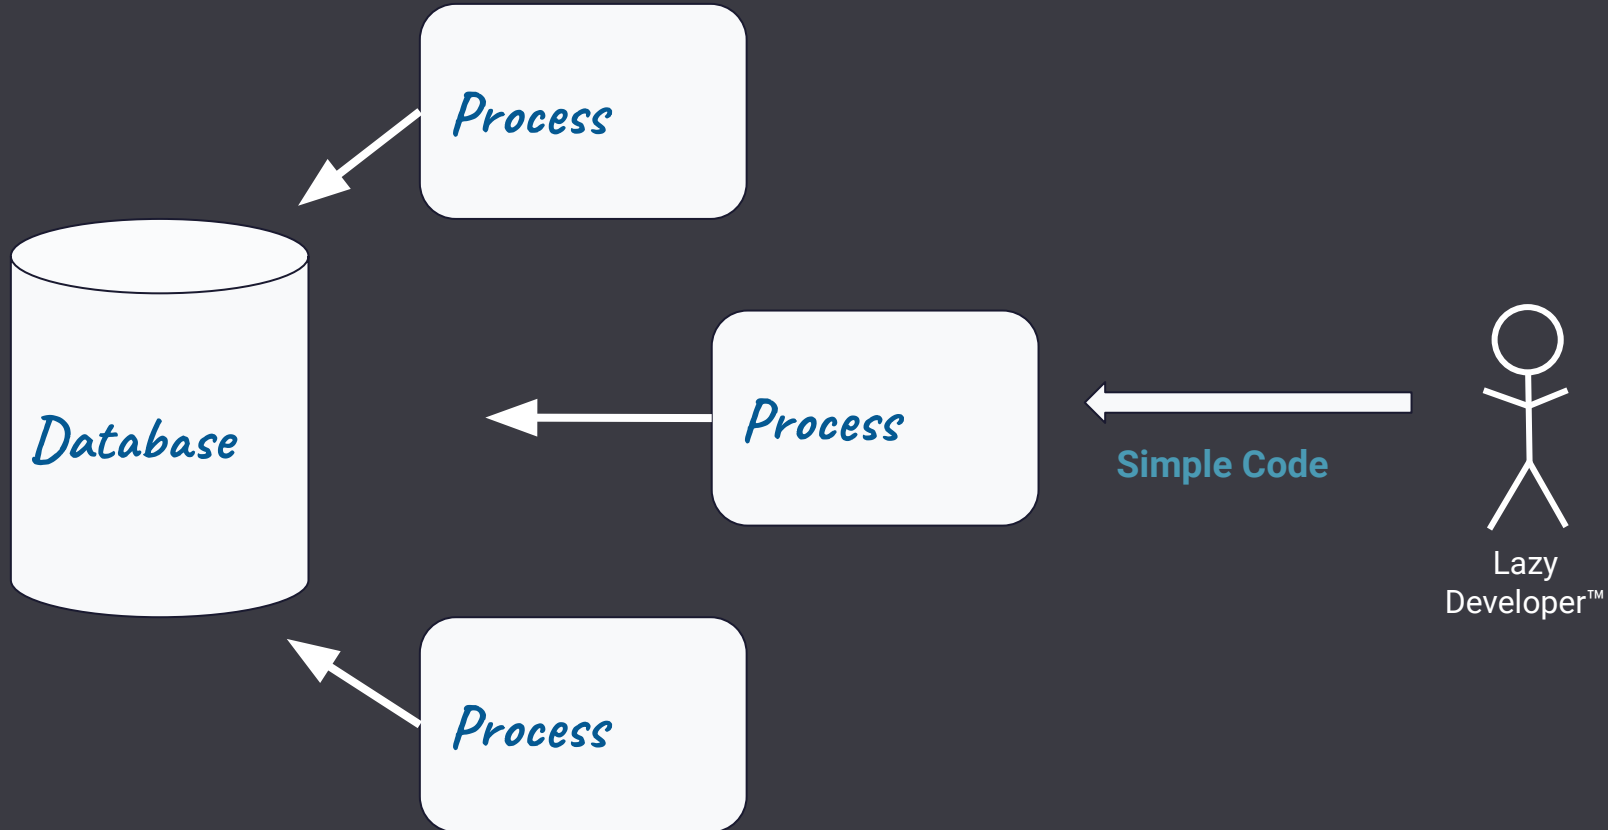
# Developer Experience




# Developer Experience



# Observer Reference Frame – Non Exclusive



# Observer Reference Frame – Exclusive



<https://pathelland.substack.com/p/i-am-so-glad-im-uncoordinated>

## Scattered Thoughts on Distributed Systems

# I Am SO Glad I'm Uncoordinated!

Technology trends have evolved to provide an abundance of CPU, Memory, Storage, and Networking. Coordination, formerly dirt cheap, has become the precious commodity.



**Pat Helland**

Feb 27, 2021



8




**ACM Keynote Video:** [https://www.youtube.com/watch?v=E6JUA\\_XH\\_tE](https://www.youtube.com/watch?v=E6JUA_XH_tE)

# Cost of Latency

CPU	Cheap
Network	Cheap
Memory	Cheap
Storage	Cheap
Coordination	Expensive(last frontier)


# Cost of Transaction

Single System




Cheap

Distributed System



Expensive





DS


# What's different



## Discussion about Consensus Protocols

Nerds may spontaneously erupt in debate, conjecture and questionable behavior.

# Paxos




- Origin of most consensus protocols
- **Proposer**: I want to do something
- **Acceptor**: Ok or Nope
- Do that until you have majority
- Network round trips add up
- Used in LWT(multi-paxos)

1989



# Raft

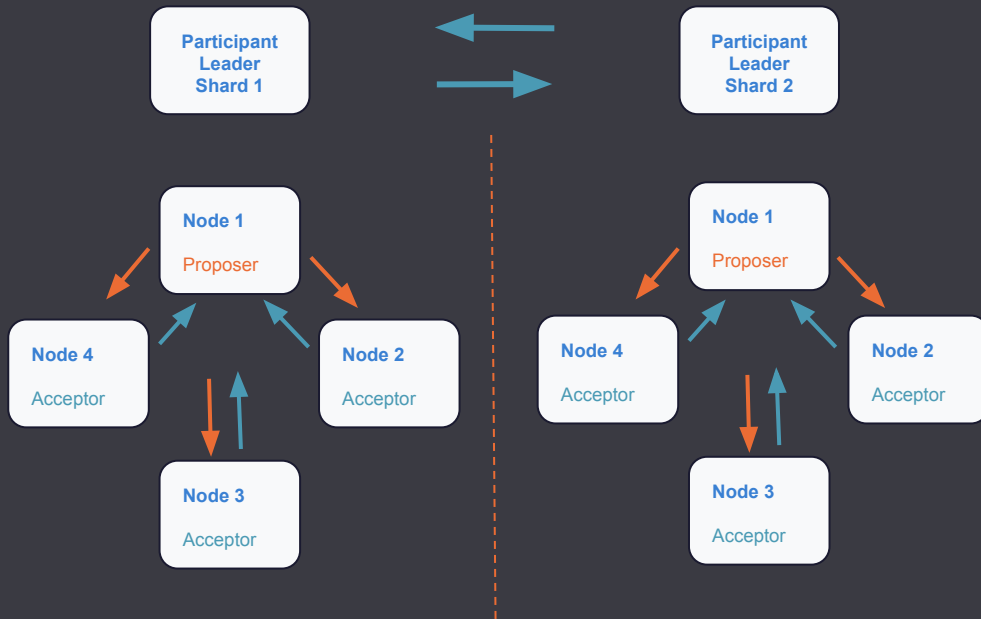


- Leader election to eliminate round trips
- **Leader**: All changes through me
- **Follower**: I trust dear leader

## Bad for Cassandra

- Failure Modes lead to latency
- Multi-datacenter? Nope

# Spanner

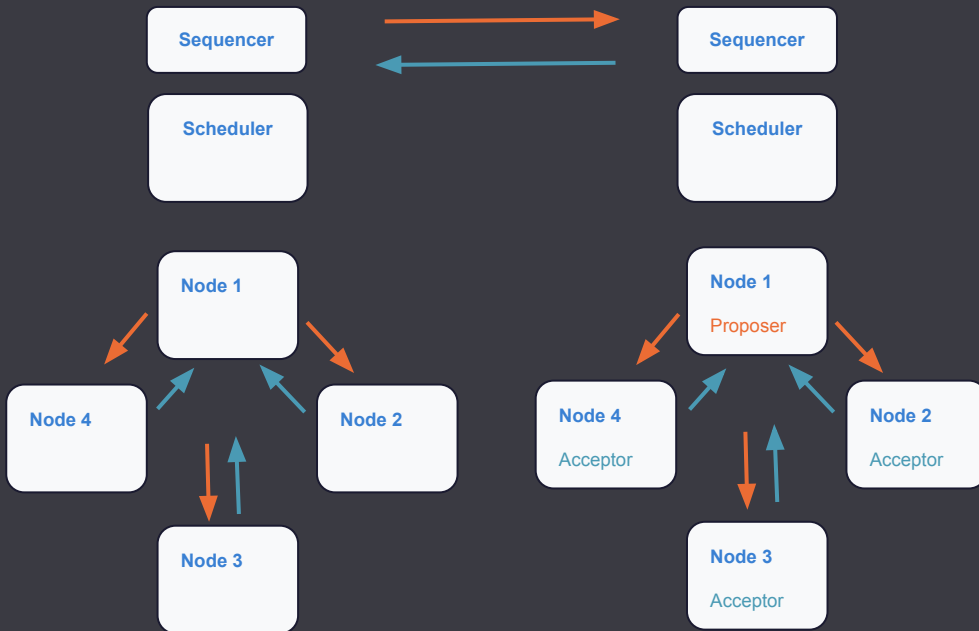


- One paxos group per shard
- Single row and single shard: Paxos
- Multi-row: Leaders coordinate for 2pc
- All depends on TrueTime™

## Bad for Cassandra

- Depends on TrueTime™
- Many hops for one insert

# Calvin




- Improvement in hops from Spanner
- Global Consensus vs Sharded
- Sequencer addresses clock skew
- Scheduler eliminates the TPC

## Bad for Cassandra

- Complex failure modes
- Latency on multi-datacenter

# Accord



- Every node has a Reorder Buffer
- Clock skew is cool
- Leaderless timestamp protocol
- Fast Path Electorates: Fault tolerance
- TL;DR One Round Trip - ish

## Good for Cassandra

- Leaderless
- Scales like Cassandra
- Failure modes match



DS

# Usage

# Transaction Syntax – Boundaries

```
BEGIN TRANSACTION
```

```
  LET <tuple> = ( SELECT <column1>,<column2>.. FROM <table> WHERE <condition>;
```

```
  SELECT <return_column> FROM <table> WHERE <condition>;
```

```
  IF <tuple_condition> THEN
```

```
    UPDATE | INSERT | DELETE..
```

```
  END IF
```

```
COMMIT TRANSACTION;
```

- Everything inside happens or it doesn't
- All statements are isolated
- All mutations are atomic
- No rollbacks (yet)

# Transaction Syntax – State collection

```
BEGIN TRANSACTION

  LET <tuple> = (SELECT <column1>,<column2>.. FROM <table> WHERE <condition>);

  SELECT <return_column> FROM <table> WHERE <condition>;

  IF <tuple_condition> THEN
    UPDATE|INSERT|DELETE..
  END IF
COMMIT TRANSACTION;
```

- Gather state for use in the conditional below
- Tuple stores one or more columns of data

# Transaction Syntax – Return value

```
BEGIN TRANSACTION

  LET <tuple> = ( SELECT <column1>,<column2>.. FROM <table> WHERE <condition>;

  SELECT <return_column> FROM <table> WHERE <condition>;

  IF <tuple_condition> THEN
    UPDATE|INSERT|DELETE..
  END IF
COMMIT TRANSACTION;
```

- Return state from before transaction

# Transaction Syntax – Conditional mutation

```
BEGIN TRANSACTION  
    LET <tuple> = (SELECT <column1>,<column2>.. FROM <table> WHERE <condition>);  
  
    SELECT <return_column> FROM <table> WHERE <condition>;  
  
    IF <tuple_condition> THEN  
        UPDATE | INSERT | DELETE..  
    END IF  
COMMIT TRANSACTION;
```

- Condition test (=, !=, >, <, <=, >=)
- Introduction of NULL, NOT NULL test
- False condition avoids updates

# New Use Cases

# Bank Transaction

Multi-Partition Exclusive Operation

# Bank Transaction – Setup

## Table

```
CREATE TABLE ks.accounts (  
  account_holder text,  
  account_balance decimal,  
  PRIMARY KEY (account_holder)  
);
```

## Data

```
INSERT INTO ks.accounts(account_holder, account_balance) VALUES ('bob', 100);  
INSERT INTO ks.accounts(account_holder, account_balance) VALUES ('alice', 100);
```



# Bank Transaction – Begin

```
BEGIN TRANSACTION
```

```
// Get the balance from Alice's account and store as a Tuple
```

```
LET fromBalance = (SELECT account_balance FROM ks.accounts WHERE account_holder='alice');
```

```
// Return the balance before update after transaction complete
```

```
SELECT account_balance FROM ks.accounts WHERE account_holder='alice';
```

```
// If Alice's account balance is greater than $20, move $ 20 to Bob
```

```
IF fromBalance.account_balance >= 20 THEN
```

```
    UPDATE ks.accounts SET account_balance -= 20 WHERE account_holder='alice';
```

```
    UPDATE ks.accounts SET account_balance +=20 WHERE account_holder='bob';
```

```
END IF
```

```
COMMIT TRANSACTION;
```

# Bank Transaction – Pre-condition

```
BEGIN TRANSACTION

// Get the balance from Alice's account and store as a Tuple
LET fromBalance = (SELECT account_balance FROM ks.accounts WHERE account_holder='alice');

// Return the balance before update after transaction complete
SELECT account_balance FROM ks.accounts WHERE account_holder='alice';

// If Alice's account balance is greater than zero, move $ 20 to Bob
IF fromBalance.account_balance >= 20 THEN
    UPDATE ks.accounts SET account_balance -= 20 WHERE account_holder='alice';
    UPDATE ks.accounts SET account_balance +=20 WHERE account_holder='bob';
END IF

COMMIT TRANSACTION;
```

# Bank Transaction – Output previous state

```
BEGIN TRANSACTION

// Get the balance from Alice's account and store as a Tuple
LET fromBalance = (SELECT account_balance FROM ks.accounts WHERE account_holder='alice');

// Return the balance before update after transaction complete
SELECT account_balance FROM ks.accounts WHERE account_holder='alice';

// If Alice's account balance is greater than zero, move $ 20 to Bob
IF fromBalance.account_balance >= 20 THEN
    UPDATE ks.accounts SET account_balance -= 20 WHERE account_holder='alice';
    UPDATE ks.accounts SET account_balance +=20 WHERE account_holder='bob';
END IF

COMMIT TRANSACTION;
```

# Bank Transaction – Conditional Statement

```
BEGIN TRANSACTION

// Get the balance from Alice's account and store as a Tuple
LET fromBalance = (SELECT account_balance FROM ks.accounts WHERE account_holder='alice');

// Return the balance before update after transaction complete
SELECT account_balance FROM ks.accounts WHERE account_holder='alice';

// If Alice's account balance is greater than zero, move $ 20 to Bob
IF fromBalance.account_balance >= 20 THEN
    UPDATE ks.accounts SET account_balance -= 20 WHERE account_holder='alice';
    UPDATE ks.accounts SET account_balance +=20 WHERE account_holder='bob';
END IF

COMMIT TRANSACTION;
```

# Bank Transaction – Commit

```
BEGIN TRANSACTION

// Get the balance from Alice's account and store as a Tuple
LET fromBalance = (SELECT account_balance FROM ks.accounts WHERE account_holder='alice');

// Return the balance before update after transaction complete
SELECT account_balance FROM ks.accounts WHERE account_holder='alice';

// If Alice's account balance is greater than zero, move $ 20 to Bob
IF fromBalance.account_balance >= 20 THEN
    UPDATE ks.accounts SET account_balance -= 20 WHERE account_holder='alice';
    UPDATE ks.accounts SET account_balance +=20 WHERE account_holder='bob';
END IF

COMMIT TRANSACTION;
```

# Inventory Management

Multi-Table/Multi-Partition Exclusive Update

# Inventory Management – Setup

## Tables

```
CREATE TABLE ks.products (  
    item text,  
    inventory_count int,  
    PRIMARY KEY (item)  
);
```

```
CREATE TABLE ks.shopping_cart (  
    user_name text,  
    item text,  
    item_count int,  
    PRIMARY KEY (user_name, item)  
);
```

## Data

```
INSERT INTO ks.products(item, inventory_count) VALUES ('PlayStation 5', 100);
```

# Inventory Management – Pre-Condition

```
BEGIN TRANSACTION

// Find out how many PlayStations are left
LET inventory = (SELECT inventory_count FROM ks.products WHERE item='PlayStation 5');

// Return the inventory count before deducting
SELECT item, inventory_count FROM ks.products WHERE item='PlayStation 5';

// Take a PlayStation out of inventory and put in users shopping cart
IF inventory.inventory_count >0 THEN
    UPDATE ks.products SET inventory_count -=1 WHERE item='PlayStation 5';
    INSERT INTO ks.shopping_cart(user_name, item, item_count)VALUES ('patrick', 'PlayStation 5', 1);
END IF

COMMIT TRANSACTION;
```



# Inventory Management – Output Previous State

```
BEGIN TRANSACTION

// Find out how many PlayStations are left
LET inventory = (SELECT inventory_count FROM ks.products WHERE item='PlayStation 5');

// Return the inventory count before deducting
SELECT item, inventory_count FROM ks.products WHERE item='PlayStation 5';

// Take a PlayStation out of inventory and put in users shopping cart
IF inventory.inventory_count >0 THEN
    UPDATE ks.products SET inventory_count -=1 WHERE item='PlayStation 5';
    INSERT INTO ks.shopping_cart(user_name, item, item_count)VALUES ('patrick', 'PlayStation 5', 1);
END IF

COMMIT TRANSACTION;
```

# Inventory Management – Conditional Statement

```
BEGIN TRANSACTION

// Find out how many PlayStations are left
LET inventory = (SELECT inventory_count FROM ks.products WHERE item='PlayStation 5');

// Return the inventory count before deducting
SELECT item, inventory_count FROM ks.products WHERE item='PlayStation 5';

// Take a PlayStation out of inventory and put in users shopping cart
IF inventory.inventory_count > 0 THEN
    UPDATE ks.products SET inventory_count -= 1 WHERE item='PlayStation 5';
    INSERT INTO ks.shopping_cart(user_name, item, item_count) VALUES ('patrick', 'PlayStation 5', 1);
END IF

COMMIT TRANSACTION;
```

# Real Atomic Batch

Foreign Key Management

# Real Atomic Batch – Setup

## Tables

```
CREATE TABLE ks.user (  
  user_id UUID,  
  email text,  
  country text,  
  city text,  
  PRIMARY KEY (user_id)  
);
```

```
CREATE TABLE ks.user_by_email (  
  email text,  
  user_id UUID,  
  PRIMARY KEY (email)  
);
```

```
CREATE TABLE ks.user_by_location (  
  country text,  
  city text,  
  user_id UUID,  
  PRIMARY KEY ((country, city), user_id)  
);
```

# Real Atomic Batch – Existence Check

```
BEGIN TRANSACTION

// Find any existing users with same email

LET existCheck = (SELECT user_id FROM ks.user_by_email WHERE email='patrick@datastax.com');

// If email isn't in use, then add the new user
IF existCheck IS NULL THEN

    INSERT INTO ks.user(user_id, email, country, city)
    VALUES (94813846-4366-11ed-b878-0242ac120002, 'patrick@datastax.com', 'US', 'Windsor');

    INSERT INTO ks.user_by_email(email, user_id)
    VALUES ('patrick@datastax.com', 94813846-4366-11ed-b878-0242ac120002);

    INSERT INTO ks.user_by_location(country, city, user_id)
    VALUES ('US', 'Windsor', 94813846-4366-11ed-b878-0242ac120002);

END IF

COMMIT TRANSACTION ;
```

# Real Atomic Batch – Execution

```
BEGIN TRANSACTION

// Find any existing users with same email
LET existCheck = ($SELECT user_id FROM ks.user_by_email WHERE email='patrick@datastax.com');

// If email isn't in use, then add the new user
IF existCheck IS NULL THEN

    INSERT INTO ks.user(user_id, email, country, city)
    VALUES (94813846-4366-11ed-b878-0242ac120002, 'patrick@datastax.com', 'US', 'Windsor');

    INSERT INTO ks.user_by_email(email, user_id)
    VALUES ('patrick@datastax.com', 94813846-4366-11ed-b878-0242ac120002);

    INSERT INTO ks.user_by_location(country, city, user_id)
    VALUES ('US', 'Windsor', 94813846-4366-11ed-b878-0242ac120002);

END IF

COMMIT TRANSACTION ;
```

# Usable Counters

Imagine a world without counter columns

# Usable Counters – Setup

## Table

```
CREATE TABLE ks.products (  
    item text,  
    inventory_count decimal,  
    PRIMARY KEY (item)  
);
```



# Usable Counters – Operations

## Set

```
BEGIN TRANSACTION
    UPDATE ks.products SET inventory_count = 100 WHERE item='PlayStation 5';
COMMIT TRANSACTION;
```

## Increment

```
BEGIN TRANSACTION
    UPDATE ks.products SET inventory_count += 1 WHERE item='PlayStation 5';
COMMIT TRANSACTION;
```

## Decrement

```
BEGIN TRANSACTION
    UPDATE ks.products SET inventory_count -= 1 WHERE item='PlayStation 5';
COMMIT TRANSACTION;
```

## Other Things (Haven't Worked Them Out Yet)

Document Locking

Lease Management

Match Making (gaming)


# Final Thoughts

## What's next?

- Performance testing. TPC-C?
- Gather user feedback
- We need to find limits and communicate early

# This Will Change Cassandra in Profound Ways

# Save the Date!



# CASSANDRA SUMMIT

MARCH 13-14, 2023 • SAN JOSE, CA

- Training day March 12
- In-person + Virtual
- CFP Open
- Reg coming soon

<http://cassandrasummit.org>

DS

# Thank You!