

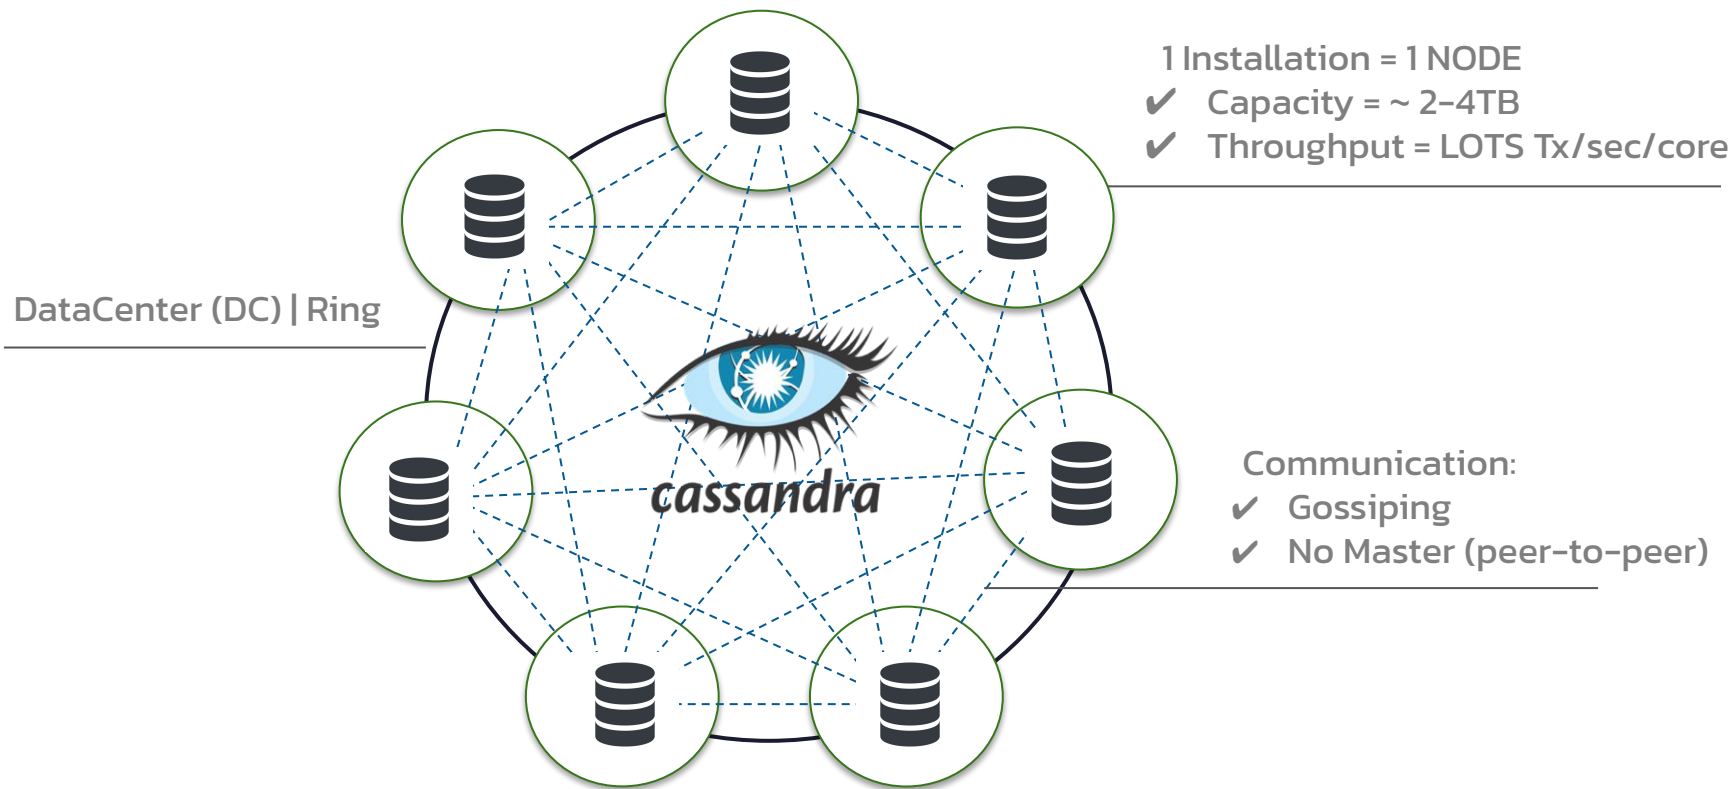


**Cassandra Day**  
2022

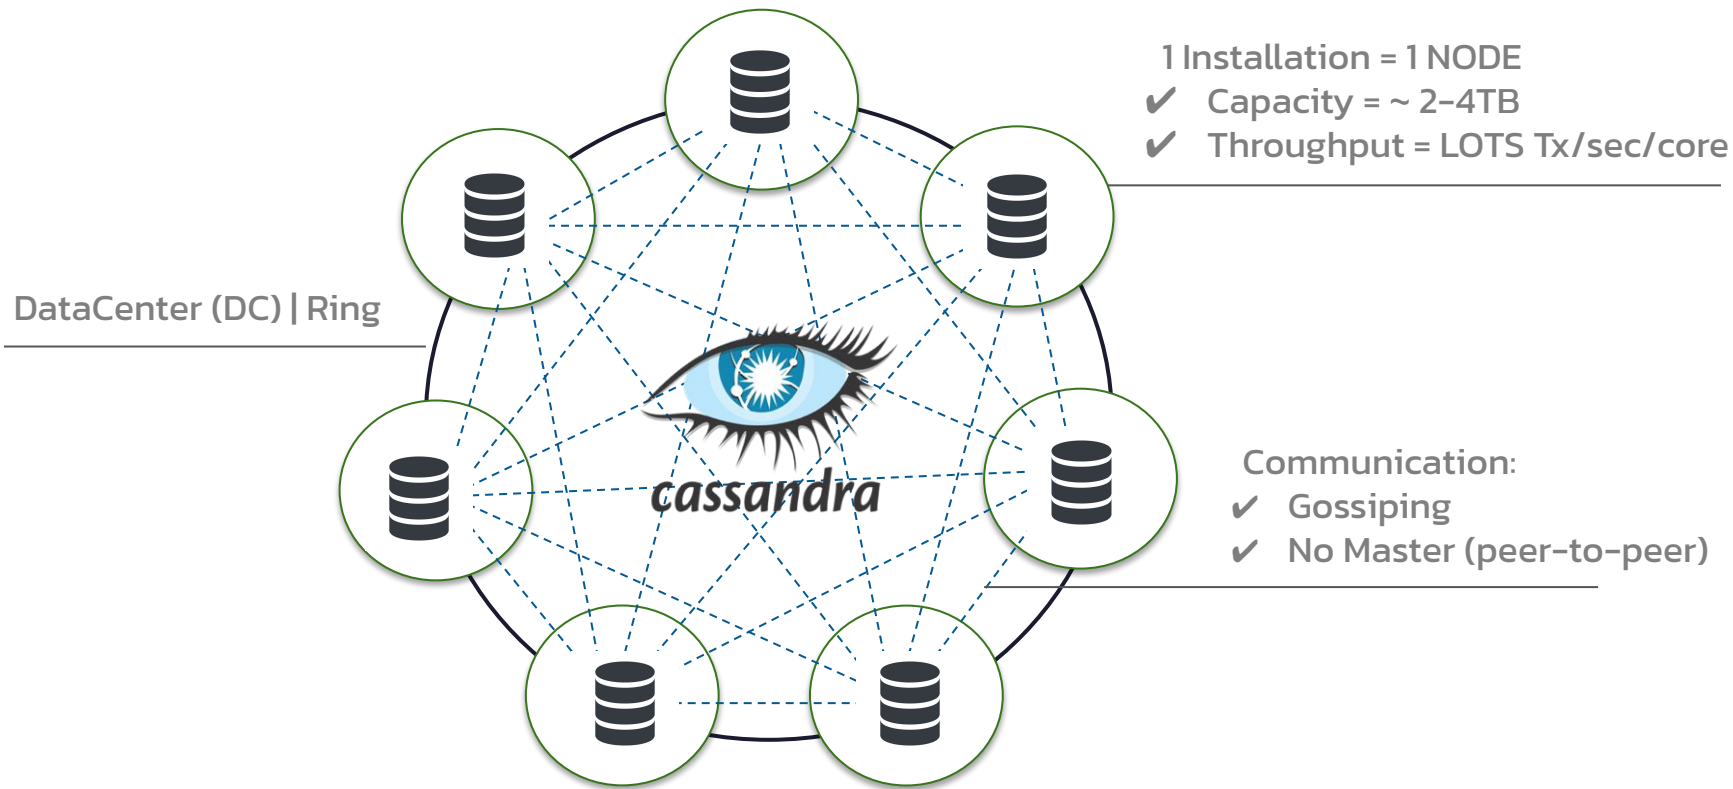
# Apache Cassandra™ in a Nutshell

**Sponsored by DataStax**

# NoSQL Distributed Database



# ~~NoSQL~~ Specialized, Purpose-built Distributed Database




# Apache Cassandra in a nutshell



1. Big Data Ready
2. Linear Scalability
3. Read / Write Performance
4. Highest Availability
5. Self-Healing and Automation
6. Geographical Distribution
7. Platform Agnostic
8. Vendor Independent

# 1. Data is Distributed

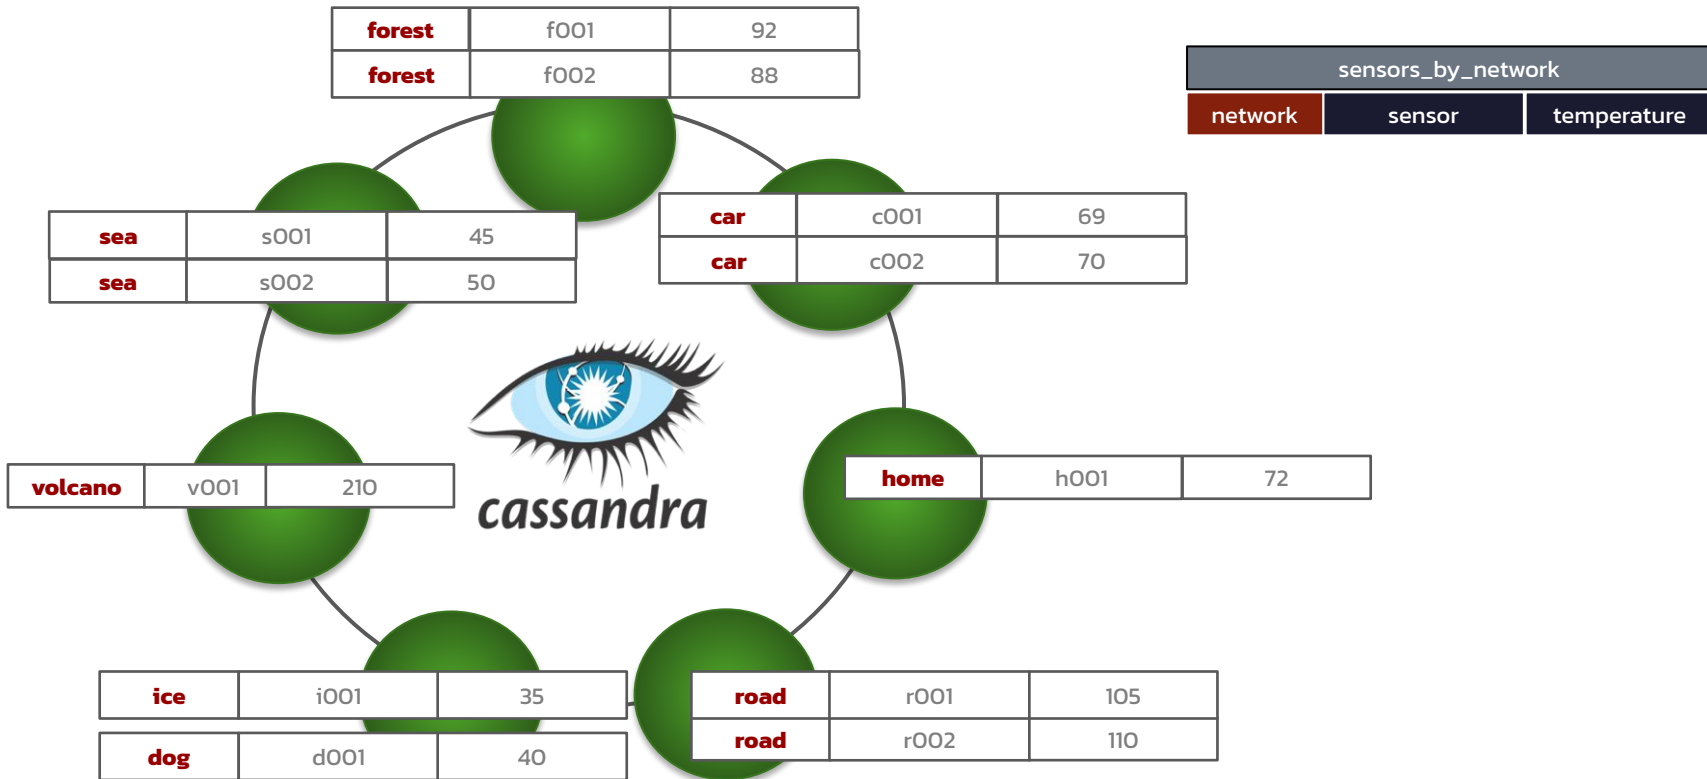


sensors_by_network		
network	sensor	temperature
forest	f001	92
forest	f002	88
volcano	v001	210
sea	s001	45
sea	s002	50
home	h001	72
car	c001	69
car	c002	70
dog	d001	40
road	r001	105
road	r002	110
ice	i001	35

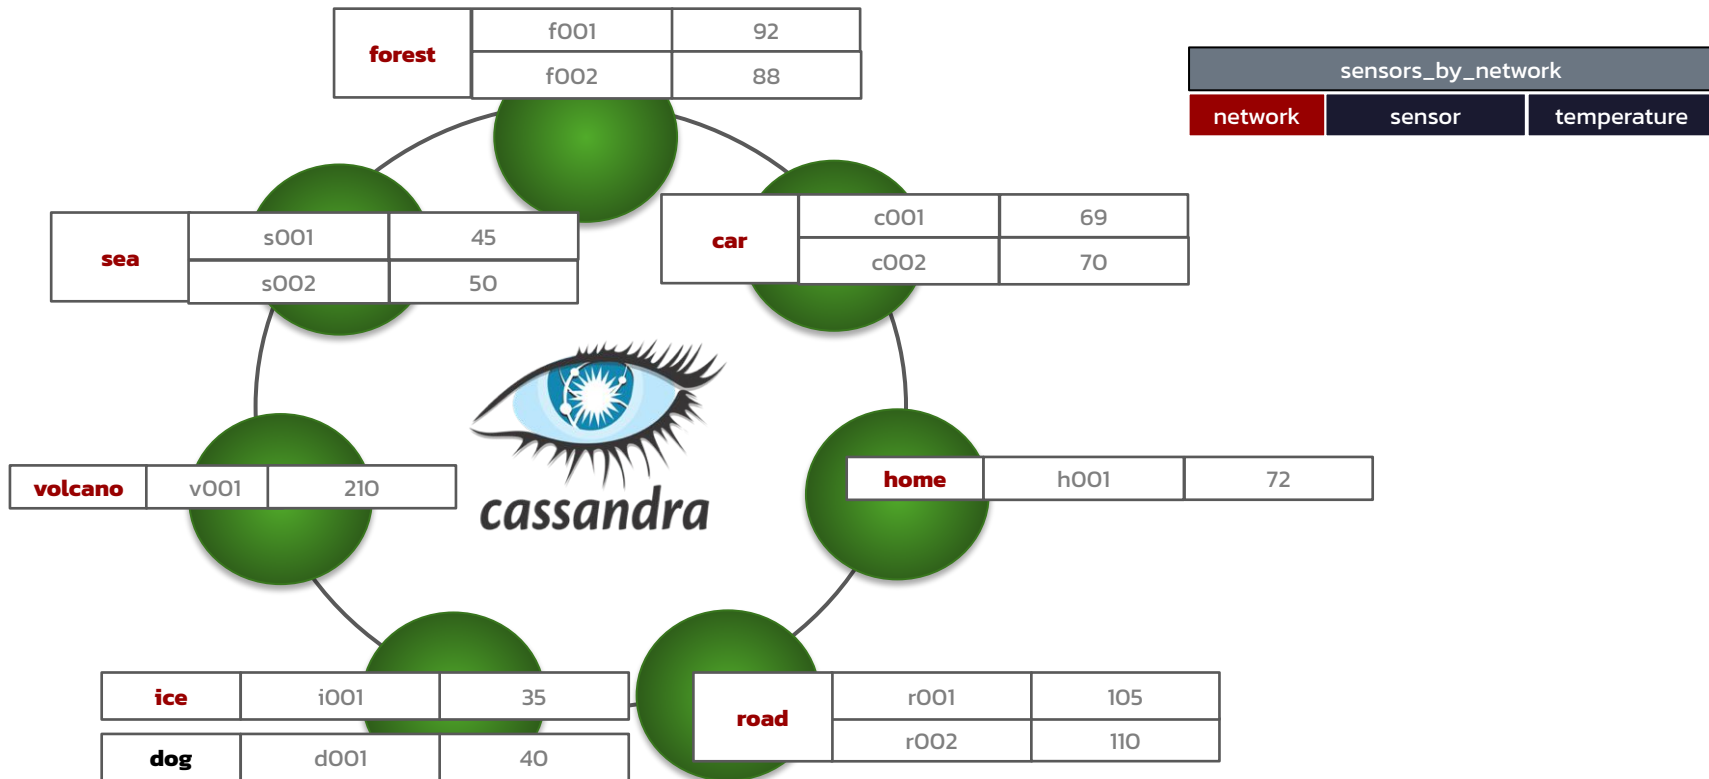
**Partition Key**

**Primary Key**

# 1. Data is Distributed



# 1. Data is Distributed




# DEMO ?



## 2. Linear Scalability

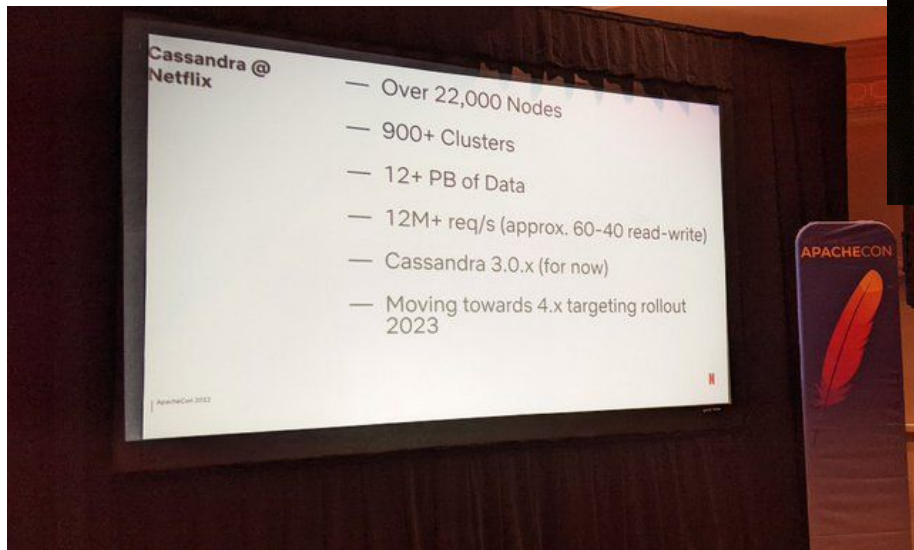
- For volume or velocity, there are no limitations
- **Linear** - No overhead on new nodes, scales with your needs\*



## 2. Linear Scalability



And many others...




**Hacker News** [new](#) | [past](#) | [comments](#) | [ask](#) | [show](#) | [jobs](#) | [submit](#)

31. ▲ **"Mother Will Be Pleased": How It Feels to Be Run over (1900)** ([publicdomainreview.org](#))  
41 points by samclemons 12 hours ago | [hide](#) | 13 comments
32. ▲ **Cassandra at Apple: 1000s of Clusters, 300k Nodes, 100 PB** ([twitter.com/erickramirezau](#))  
258 points by mfiguiere 17 hours ago | [hide](#) | 156 comments
33. ▲ **DNS based ad/abuse domain blocking** ([github.com/openwrt](#))  
14 points by behnamoh 7 hours ago | [hide](#) | [discuss](#)
34. ▲ **Technical milestone reached: global earth system simulations with 1.2 km resolu** ([mpg.de](#))  
149 points by jokabrink 15 hours ago | [hide](#) | 73 comments
35. ▲ **Why Disconnecting from Your Phone Is So Good for You – and How to Do It** ([time.com](#))  
67 points by ppjim 7 hours ago | [hide](#) | 30 comments
36. ▲ **Happy 20th birthday Jira You suck so bad** ([rida.dev](#))  
70 points by julian\_digital 2 hours ago | [hide](#) | 38 comments

# Geographical Distribution

Cassandra's trademark is multi-datacenter deployments, granting you an exceptional capability for disaster tolerance while keeping your data close to your clients - worldwide.


All DCs are active (available for both writes and reads)!



# High Availability


Replication, Decentralisation, and Topology-Aware Placement Strategy take care of possible downtimes:

- Multiple Live Replicas
- No Single Point of Failure
- Network topology-aware data placement
- Client-side Smart Reconnection and Strong Retry Mechanism



# Platform Agnostic

Apache Cassandra is **not bound to any platform** or service provider, helping you build hybrid-cloud and multi-cloud solutions with ease.



# Vendor Independent

Cassandra doesn't belong to any of commercial vendors but controlled by a non-profit Open Source **Apache Software Foundation**, already familiar to you by *Hadoop, Spark, Kafka, Zookeeper, Maven* and many other projects.



# Apache Cassandra

## Built as a Technical Necessity



of Fortune 100  
companies  
have adopted

### High Availability

#### Always On

Every second of downtime  
translates into lost revenue

### Linear Scalability

#### Hyper-Scalability

Millions of operations per day,  
hour, or second

### Low Latency

#### Faster Pace

Every millisecond of latency  
has consequence

### Global Distribution

#### Data Everywhere

On-premises, hybrid, multi-cloud,  
centralized, or edge



**Cassandra Day**  
2022

**Thank you!**

**Sponsored by DataStax**