



**Cassandra Day**  
2022

# Future of Cassandra

**Sponsored by DataStax**

# Cassandra 4.0 (current)

25% performance improvement

Faster big clusters

Introduction of virtual tables for observability

Transient Replicas for storage optimisation

Audit logging + full query logging

Support for JDK 11

# Cassandra 4.1 (upcoming)

Pluggable extensions – storage, auth, encryption

Operational guardrails

Improved configuration format

Deny-listing partitions

Paxos v2 consensus protocol

CQL improvements

Consistent hint windows


## C\* 4.1 – Pluggable extensions

- Pluggable memtables
- Pluggable network encryption
- Pluggable authentication



## C\* 4.1 – Improved configuration

- Standardised names
- Flexible units in values
- Compatibility-mode JVM flags
- Use of virtual tables for verification



## C\* 4.1 – Guardrails

- Prevents wrong usage
- Hard and soft limits
- Override for experts
- Based on work in Astra DB



## C\* 4.1 – Deny-listing records

- Disables "bad" partitions
  - very large partitions
  - unintended data
  - attack vectors
- Prevents partitions from being accessed to minimise negative impact to cluster



# Cassandra 5.0 (future)

ACID transactions with Accord consensus

Storage-Attached Index

Transactional cluster metadata

Trie memtables

Dynamic data masking

Support for JDK 17



## C\* 5.0 – ACID transactions

- Fully-compliant ACID
- Globally available consensus with Accord protocol
- Statements executed **atomically** in **isolation** from other processes



# C\* 5.0 – Example transaction

```
BEGIN TRANSACTION
```

```
    LET fromBalance = (SELECT balance FROM accounts  
                        WHERE holder='alice');
```

```
    IF fromBalance.balance >= 20 THEN
```

```
        UPDATE accounts SET balance -= 20 WHERE holder='alice';
```


```
        UPDATE accounts SET balance += 20 WHERE holder='bob';
```

```
    END IF
```

```
COMMIT TRANSACTION ;
```


# C\* 5.0 – Storage Attached Index

- Replaces legacy secondary indexing and SASI
- Highly-scalable
- Globally-distributed index
- Faster than existing implementations
- Uses significantly less disk space



## C\* 5.0 – Transactional cluster metadata

- Fast, unlimited concurrent schema changes
- Better elasticity with fast scale up and down
- Automatic rebalancing of nodes
- PLUS more to come



# C\* 5.0 – Trie memtables

- Memtable implementation based on trie data structure
- Better memory management due to more compact storage
- Can be placed off-heap, GC-friendly when on-heap
- Fast lookups resulting in faster writes + faster memtable reads



# We need you!



- Anyone can contribute
- Mentors will guide you through the process
- Documentation is an easy way to get started
- Career booster + get hired!
- [cassandra.apache.org](https://cassandra.apache.org)





# CASSANDRA SUMMIT

MARCH 13-14, 2023 • SAN JOSE, CA

---

## SAVE THE DATE

---

McEnery Convention Center  
San Jose, CA + Virtual

# Be a contributor today!



 Watch tag







**Cassandra Day**  
2022

**Thank you!**

**Sponsored by DataStax**