



Cassandra Day  
2022

# Cassandra 4.1 Guardrails

**Sponsored by DataStax**

# Apache Cassandra Guardrails

**A framework that allows to enforce certain system-wide rules, soft and hard limits.**




- First implemented in DSE 6.8
- Merged into Cassandra this year
- Released with C\* 4.1
- Minimum-to-zero overhead
- Designed to easily implement new and/or custom rules



→ Ask Aleks to explain why it happens! ←

- Too many tables
- Or secondary indexes
- Queries touching too many partitions
- Queries using replica-side filtering
- Size of non-frozen collections
- Etc.

# Configuration

- *cassandra.yaml* 
- Dynamically via JMX 
- Virtual tables 

```
tables_warn_threshold: 5
tables_fail_threshold: 10
secondary_indexes_per_table_warn_threshold: 5
secondary_indexes_per_table_fail_threshold: 10
allow_filtering_enabled: false
partition_keys_in_select_warn_threshold: 10
partition_keys_in_select_fail_threshold: 20
collection_size_warn_threshold: 10MiB
collection_size_fail_threshold: 20MiB
```

# Soft and Hard Limits

```
CREATE TABLE table81 (.....);
```

## **Warning :**

Guardrail tables violated: Creating table t81, current number of tables exceeds warning threshold of 80.

```
CREATE TABLE k.t101 (.....);
```

InvalidRequest: **Error** from server:  
code=2200 [Invalid query] message="Guardrail  
tables violated: Cannot have more than 100  
tables, aborting the creation of table t101"

# CQL-Layer Guardrails 1/2

- Implemented on CQL level (quick)
- Don't involve Storage level
- Don't involve additional replicas
- "Foreground"
- Support both soft and hard limits

```
cqlsh> SELECT * FROM table WHERE v=0  
ALLOW FILTERING;
```

**InvalidRequest: Error from server:  
code=2200 [Invalid query]  
message="Guardrail allow\_filtering  
violated: Querying with ALLOW FILTERING  
is not allowed"**



# Background Guardrails 2/2

- Run in background
- Aren't associated with any query
- Used when CQL-specific guardrail would be slow or not related to CQL
- "Background"
- Support mostly soft limits
- **Watch the logs!**

- Disk Space Usage
- Number of items in a non-frozen collections

# Already Available

- Number of user keyspaces
- Number of user tables
- Number of columns per table
- Number of secondary indexes per table
- Number of materialized tables per table
- Number of fields per user-defined type
- Number of items in a collection
- Number of partition keys selected by an IN restriction
- Number of partition keys selected by the cartesian product of multiple IN restrictions.
- Allowed table properties
- Allowed read consistency levels
- Allowed write consistency levels
- Allowed write consistency levels
- Collections size
- Query page size
- Minimum replication factor
- Data disk usage, defined either as a percentage or as an absolute size
- Whether user-defined timestamps are allowed
- Whether GROUP BY queries are allowed
- Whether the creation of secondary indexes is allowed
- Whether the creation of uncompressed tables is allowed
- Whether querying with ALLOW FILTERING is allowed
- Whether dropping or truncating a table is allowed



To be Available

And more to come!

Psssssst. You can add yours, it's easy!

# Implementation and Migration

**Be careful planning to enable guardrails. Your app may totally fail!**

1. First start with soft limits!
2. Watch the logs!
3. Test PRs (new code) against it!



The background is a solid green color. Overlaid on this are several sets of thin, light green lines that flow in a wavy, organic pattern across the frame. These lines are densely packed in some areas and more sparse in others, creating a sense of movement and depth. Small, light green dots are scattered throughout the composition, often appearing along the paths of the flowing lines.

# DEMO



# Thank you!

**Sponsored by DataStax**