



**Cassandra Day**  
2022

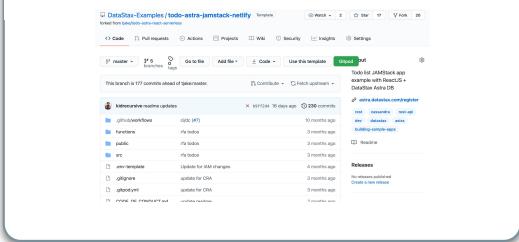
Hands-on Workshop

# Cassandra Fundamentals

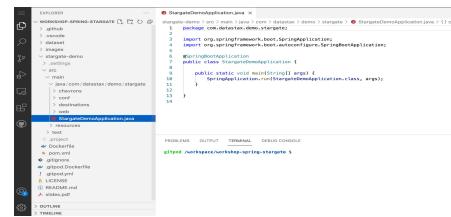
**Sponsored by DataStax**

Nothing to install !

## Source code + exercises + slides



## IDE



## Database + Api + Streaming



DataStax  
**Astra DB**

Hands-On Housekeeping

572 and counting



## Intro to Cassandra Workshop

Awarded to **YOUR NAME** • attendee@workshop.com

Issued on Jun 16, 2021

Offered By  
[DataStax Developers](#)

Upgrade Complete! This badge is to certify successful completion of the DataStax Cassandra Workshop: "Fundamentals of Apache Cassandra".



Verified  
Last verified by Badgr on Jun 17, 2021

[Re-verify Badge](#)

### EARNING CRITERIA

Recipients must complete the earning criteria to earn this Badge

To earn this badge, individuals must complete the following during the [Intro to Cassandra Workshop](#):

- Attend the lecture
- Complete the practical steps by doing all required



Achievement Unlocked !

# 01



Why Cassandra ?

# 02

Introduction to  
Apache Cassandra™

# 03

Distributed Architecture

# 04

CAP Theorem

# 05

Tables  
and Partitions

# 06

What's next?  
Homework, Next Session

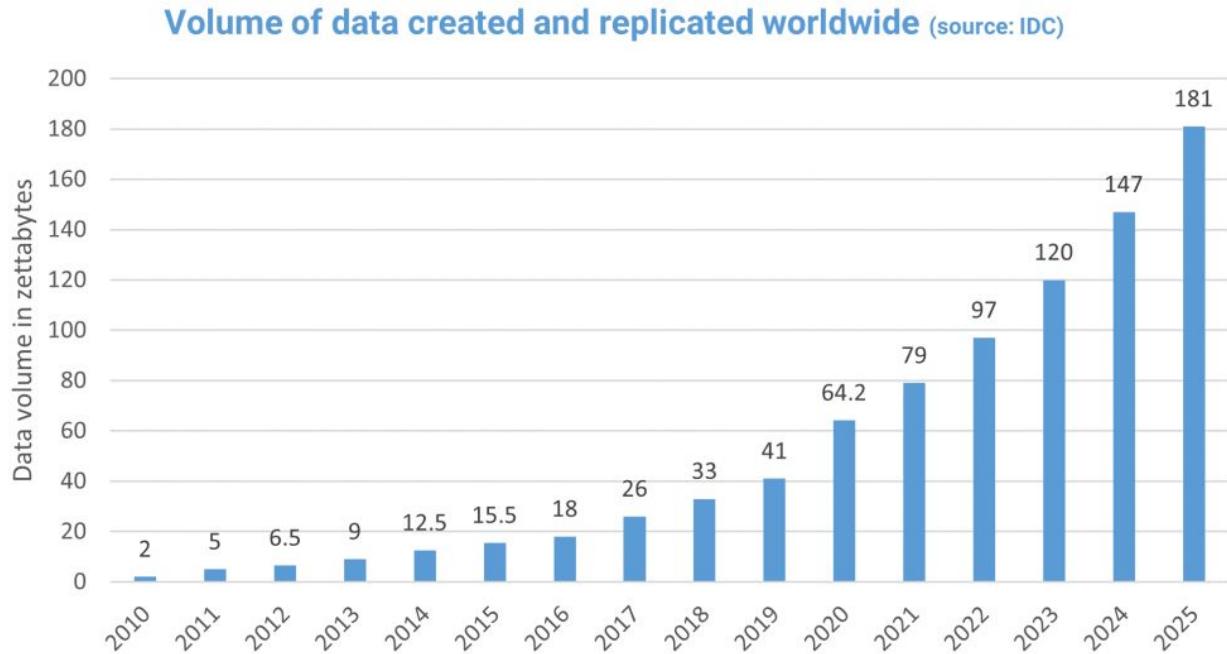


Agenda

# RDBMS were designed to run on a single machine

WHAT'S A ZETTABYTE?	
1 kilobyte	1,000,000,000,000,000,000,000
1 megabyte	1,000,000,000,000,000,000,000
1 gigabyte	1,000,000,000,000,000,000,000
1 terabyte	1,000,000,000,000,000,000,000
1 petabyte	1,000,000,000,000,000,000,000
1 exabyte	1,000,000,000,000,000,000,000
1 zettabyte	1,000,000,000,000,000,000,000

SOURCE: CISCO



Relation Databases have **limited scalability**

## OLTP: Online Transaction Processing

Fast Queries  
“Customer-facing”  
High number of transactions  
Usually Hot / Live Data  
High SLA Requirements  
(Response Time / Availability)

## OLAP: Online Analytical Processing

Complex Queries  
Historical  
High volume of data  
Often “Cold Data”  
Used by / for Analytics



Response Time



Relational Databases are **versatile**

**In the year 2008 Facebook reached their first 100,000,000 users.**



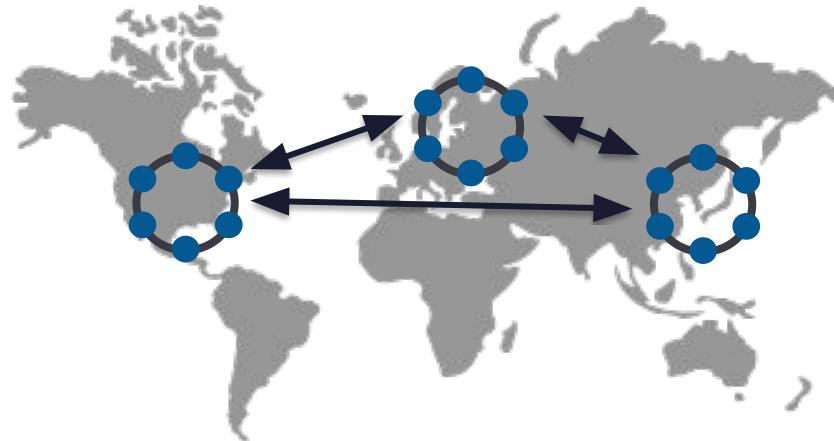
**Read this number again**



- **Users around all the world**
- **Fast Internet Connection**
- **A LOT of data**
- **High Availability requirements**



**Creation of Cassandra**



Distributed Decentralised OLTP Database

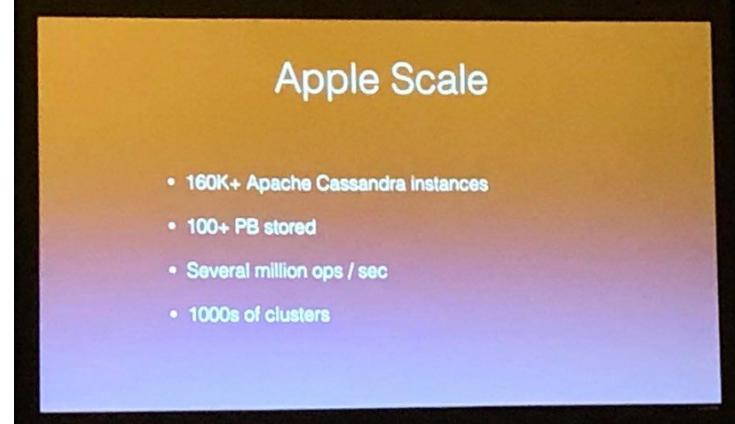
## Apache Cassandra @ Netflix

- . 98% of streaming data is stored in Apache Cassandra
- . Data ranges from customer details to viewing history to billing and payments
- . Foundational datastore for serving millions of operations per second

- 30 million ops/sec on most active single cluster
- 500 TB most dense single cluster
- 9216 CPUs in biggest cluster

O(100) Clusters  
O(10000) Instances  
O(10,000,000) Replications per second  
O(100,000,000) Operations per second  
O(1,000,000,000,000,000) Petabytes of data

[dtsx.io/cassandra-at-netflix](http://dtsx.io/cassandra-at-netflix)



And many others...



Cassandra Biggest Users (and Developers)

# Apache Cassandra

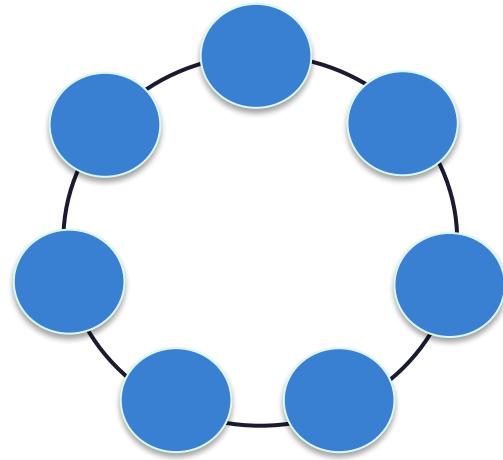
NoSQL Distributed Decentralised Database Management System



- Big Data Ready
- Read / Write Performance
- Linear Scalability
- Highest Availability
- Self-Healing and Automation
- Geographical Distribution
- Platform Agnostic
- Vendor Independent



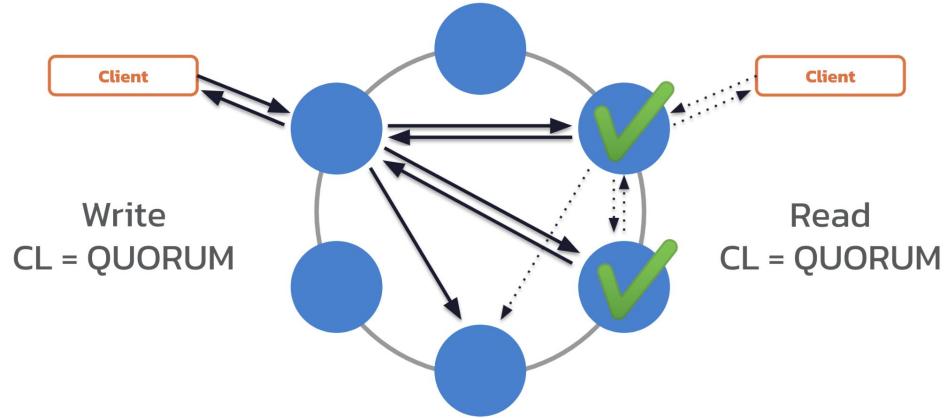
**Partitioning** over distributed architecture makes the database capable to handle data of any size: we mean petabytes scale. Need more volume? Add more nodes.



Big Data Ready

Even a single Cassandra node is very performant but a cluster consisting of multiple nodes and data centers brings throughput to the next level.

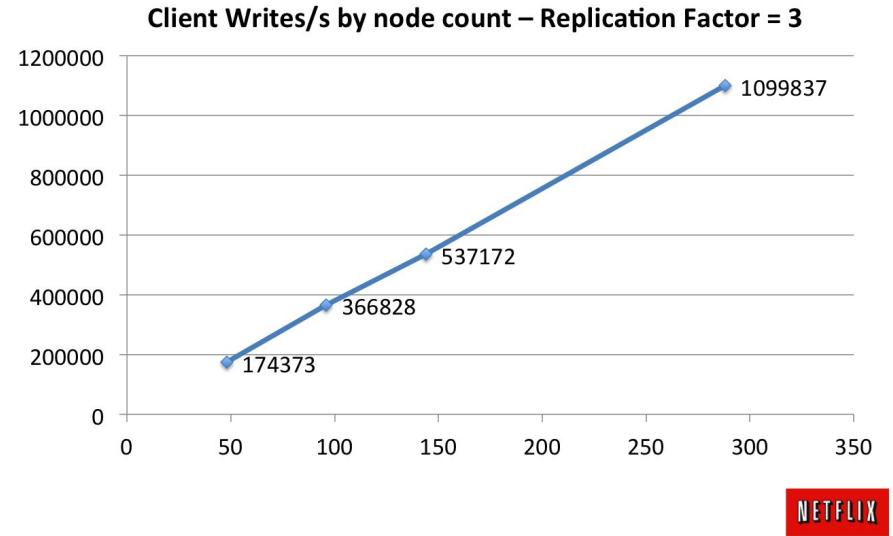
Decentralisation (**masterless architecture**) means that every node is able to deal with any request, read or write.



Read / Write Performance



- For volume or velocity, there are no limitations
- **Linear** - No overhead on new nodes, scales with your needs\*

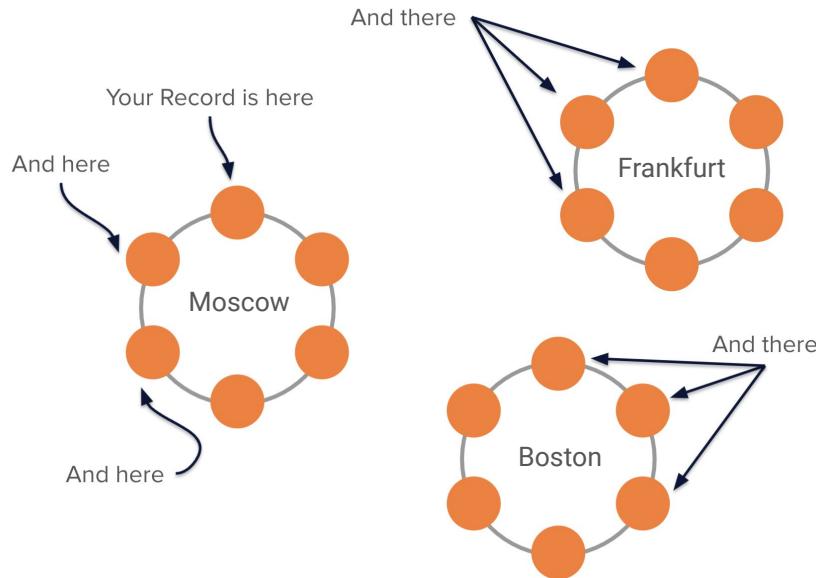


Linear Scalability



Replication, Decentralisation, and Topology-Aware Placement Strategy take care of possible downtimes:

- Multiple Live Replicas
- No Single Point of Failure
- Network topology-aware data placement
- Client-side Smart Reconnection and Strong Retry Mechanism



Highest Availability



Operations for a huge cluster can be exhausting so Apache Cassandra clusters are smart and able to scale, change data placement and recover automatically.

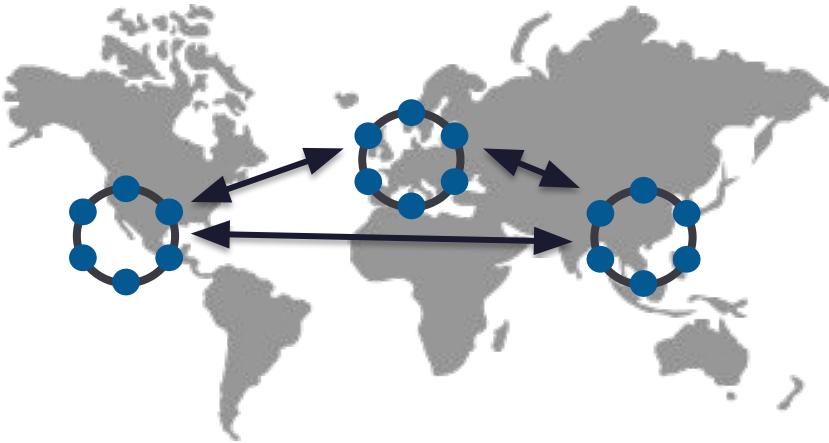


**Self-Healing and Automation**



Cassandra's trademark is multi-datacenter deployments, granting you an exceptional capability for disaster tolerance while keeping your data close to your clients - worldwide.

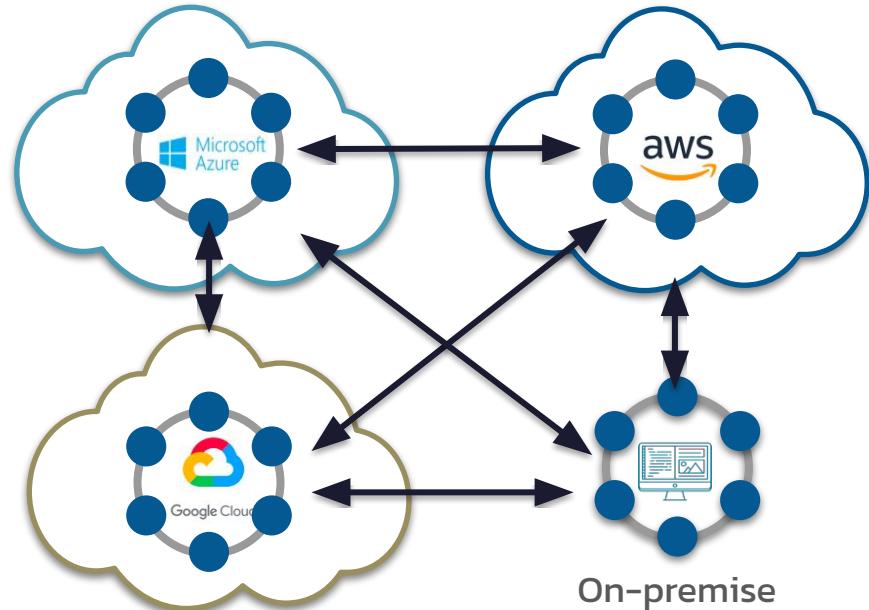
All DCs are active (available for both writes and reads)!



## Geographical Distribution



Apache Cassandra is **not bound to any platform** or service provider, helping you build hybrid-cloud and multi-cloud solutions with ease.



Platform Agnostic

Cassandra doesn't belong to any of commercial vendors but controlled by a non-profit Open Source **Apache Software Foundation**, already familiar to you by *Hadoop*, *Spark*, *Kafka*, *Zookeeper*, *Maven* and many other projects.



Vendor Independent



# 01



Why Cassandra ?

# 02

Introduction to  
Apache Cassandra™

# 03

Distributed Architecture

# 04

CAP Theorem

# 05

Tables  
and Partitions

# 06

What's next?  
Homework, Next Session



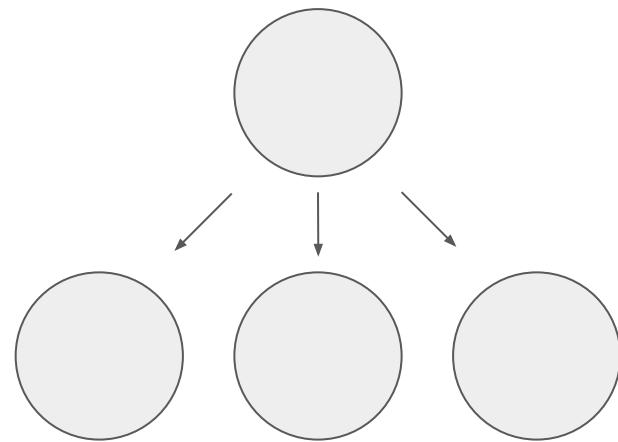
Agenda

# All servers are created equal

*Cassandra Holy Book, P.I Paragraph I*



Leader Server (Write/Read)

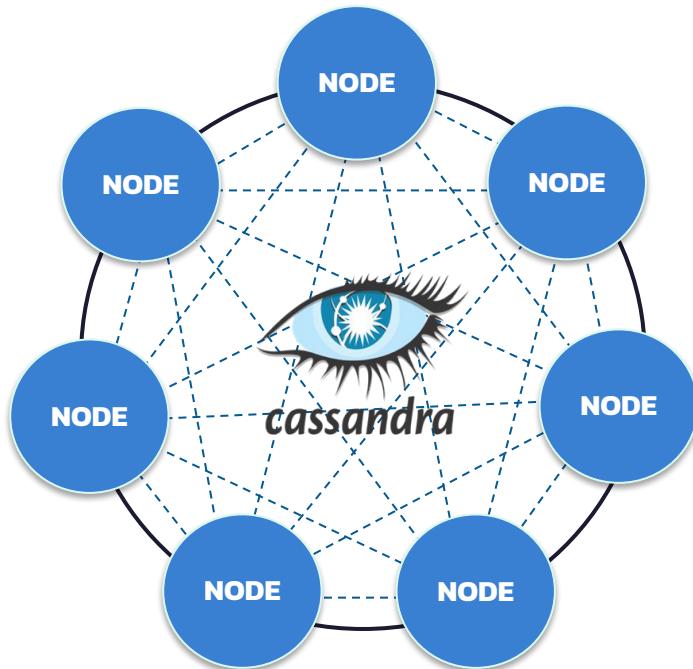


1. Single Point of Failure
2. Hard to scale for writes
3. Application needs to know where to write



Traditional Architecture





1. NO Single Point of Failure
2. Scales for writes and reads
3. Application can contact any node  
(in case of failure - just contact next one)



Master-less (Peer-to-Peer) Architecture

# Data is Partitioned



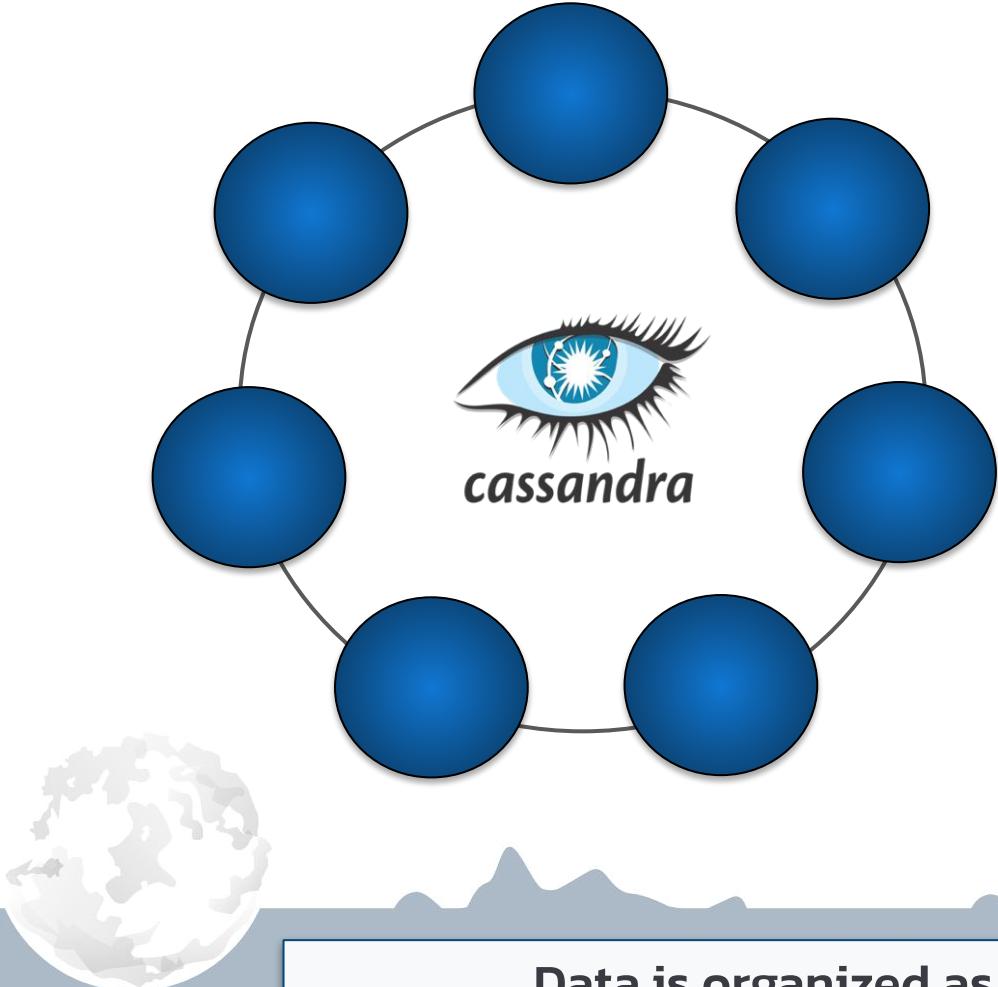
```
CREATE TABLE sensors_by_network (
    network      text,
    sensor       text,
    temperature integer,
    ...
);
```

sensors_by_network		
network	sensor	temperature
alabama	a001	92
alabama	a002	88
california	c001	210
delaware	d001	45
delaware	d002	50
florida	f001	72
georgia	g001	69
georgia	g002	70
hawaii	h001	40
idaho	i001	105
idaho	i002	110
kansas	k001	35



Let's speak data

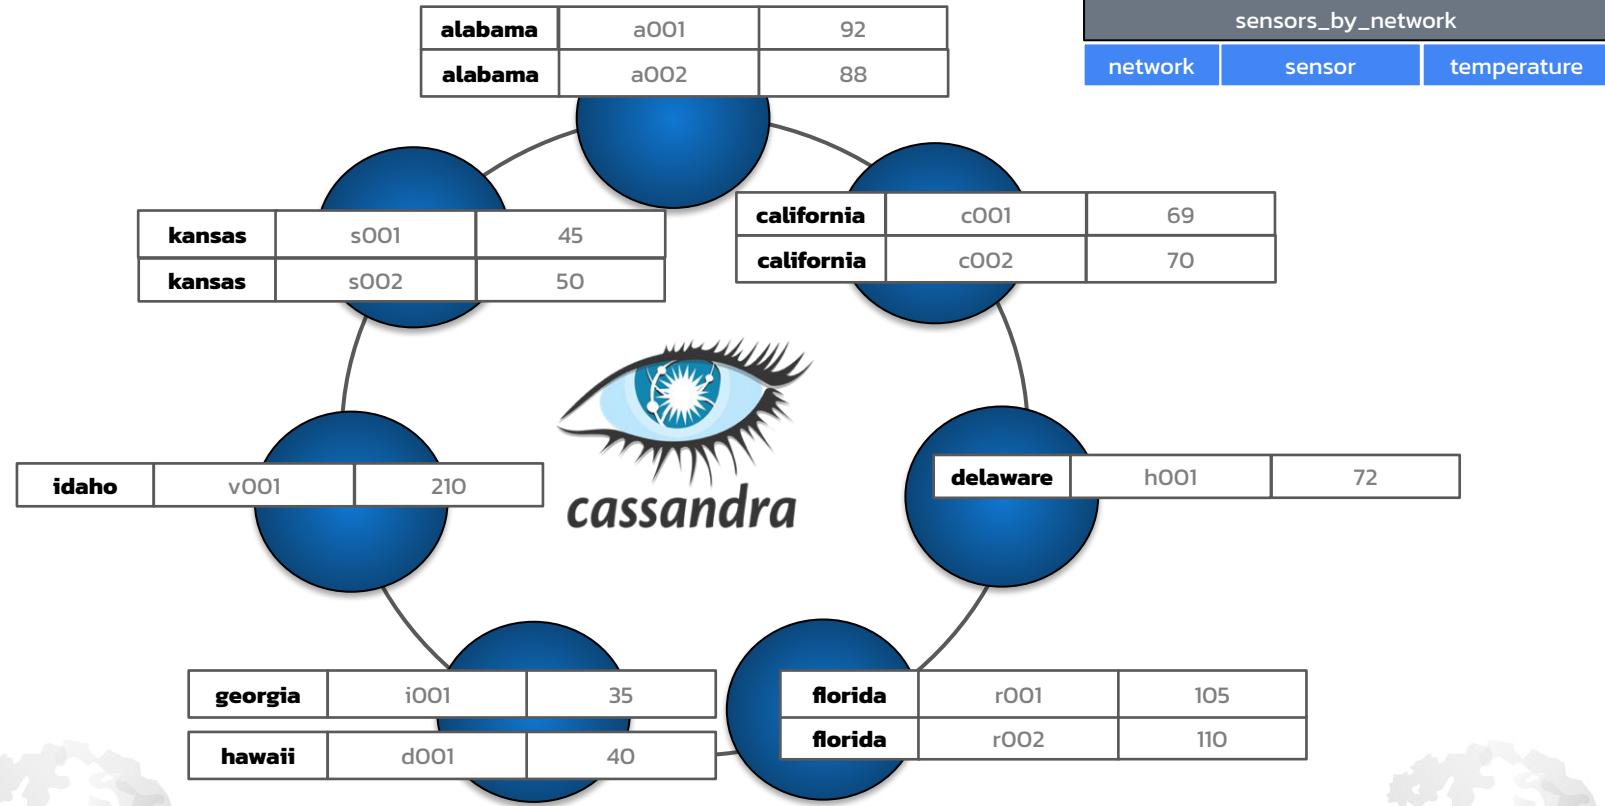




sensors_by_network		
network	sensor	temperature
<b>alabama</b>	a001	92
	a002	88
<b>california</b>	c001	210
<b>delaware</b>	d001	45
<b>delaware</b>	d002	50
<b>florida</b>	f001	72
<b>georgia</b>	g001	69
<b>georgia</b>	g002	70
<b>hawaii</b>	h001	40
<b>idaho</b>	i001	105
<b>idaho</b>	i002	110
<b>kansas</b>	k001	35

Partition Key

Data is organized as Tables



Data is organized as Distributed tables

```
CREATE TABLE sensor_data.sensors_by_network (
    network      text,
    sensor       text,
    temperature integer,
    PRIMARY KEY ((network), sensor)
);
```

Table



Partition key



Key-based Partitioning

network	sensor	Value
alabama	f001	92
alabama	f002	88
idaho	s001	45
idaho	s002	50
hawaii	r001	105
hawaii	r002	110

Partition Keys



Network	Sensor	Value
59	f001	92
59	f002	88
12	s001	45
12	s002	50
45	r001	105
45	r002	110

Tokens

Cassandra Nodes



Partitioning and Token Ranges

Why partitioning?  
Because scaling doesn't have to be [s]hard!

Big Data doesn't fit to a single server, splitting it into chunks we can easily spread them over dozens, hundreds or even thousands of servers, adding more if needed.



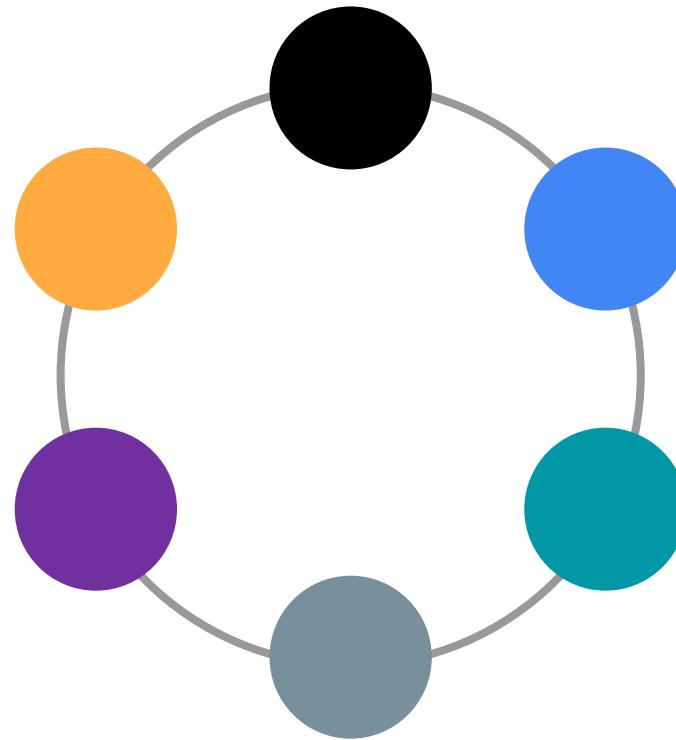


# Data is Replicated



$RF = ?$

Replication Factor  
means the number  
of nodes used to  
store each partition



**Replication Factor**

```
CREATE KEYSPACE sensor_data  
  WITH REPLICATION = {  
    'class' : 'NetworkTopologyStrategy',  
    'us-west-1' : 3,  
    'eu-east-2' : 5  
};
```

keyspace

replication strategy



Replication factor by data center

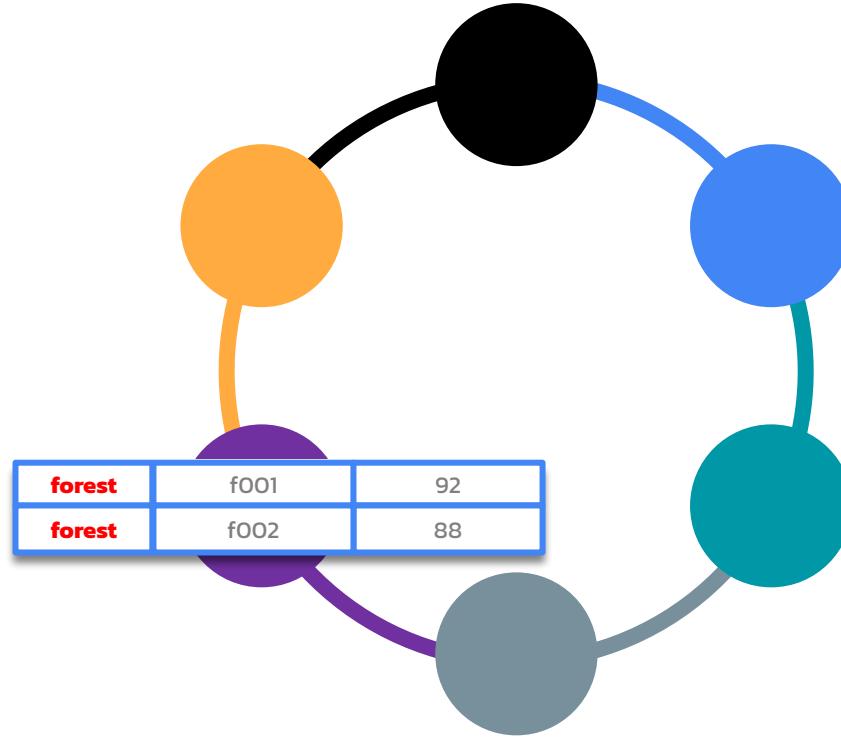


Data is Replicated



**RF = 1**

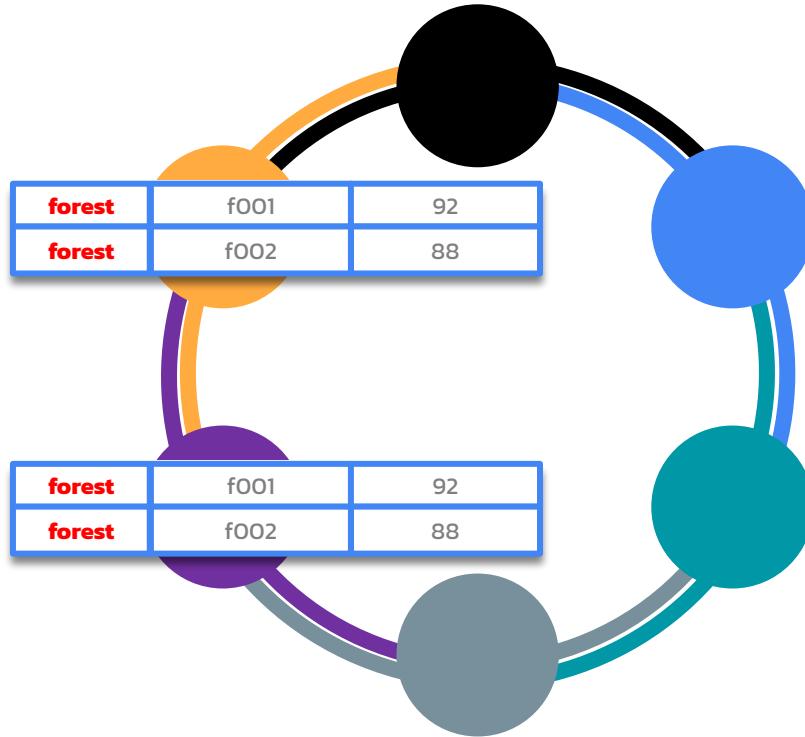
Replication Factor 1  
means that every  
partition is stored  
on 1 node



## Replication Factor

**RF = 2**

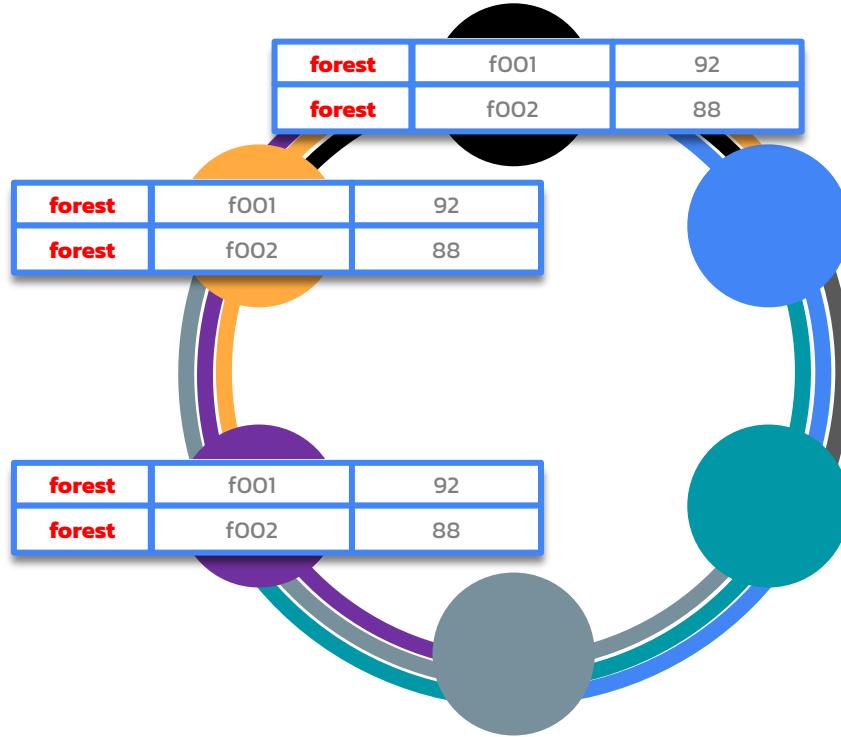
Replication Factor 2  
means that every  
partition is stored  
on 2 nodes



**Replication Factor**

**RF = 3**

Replication Factor 3  
means that every  
partition is stored  
on 3 nodes



**Replication Factor**

Replication Factor = 1

1-25

26-50

51-75

76-100

Replication Factor = 2

1-25,  
26-50

26-50,  
51-75

51-75,  
76-100

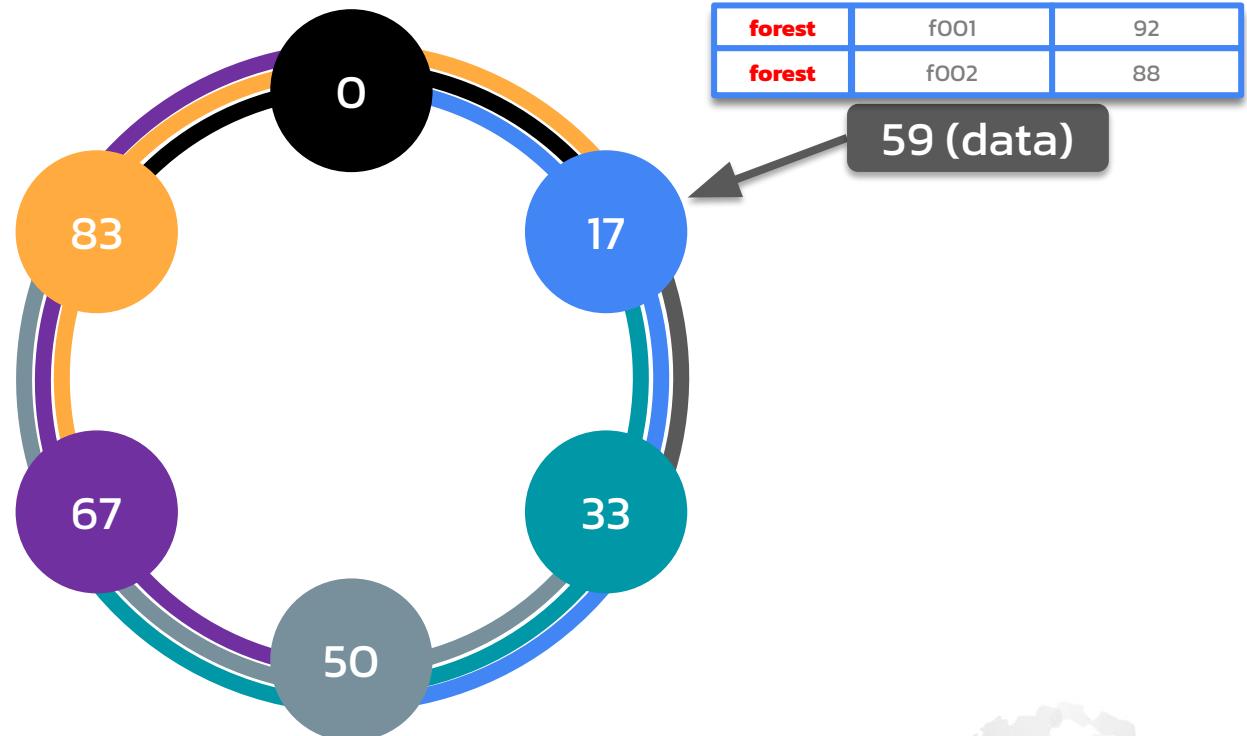
76-100,  
1-25

And so on...



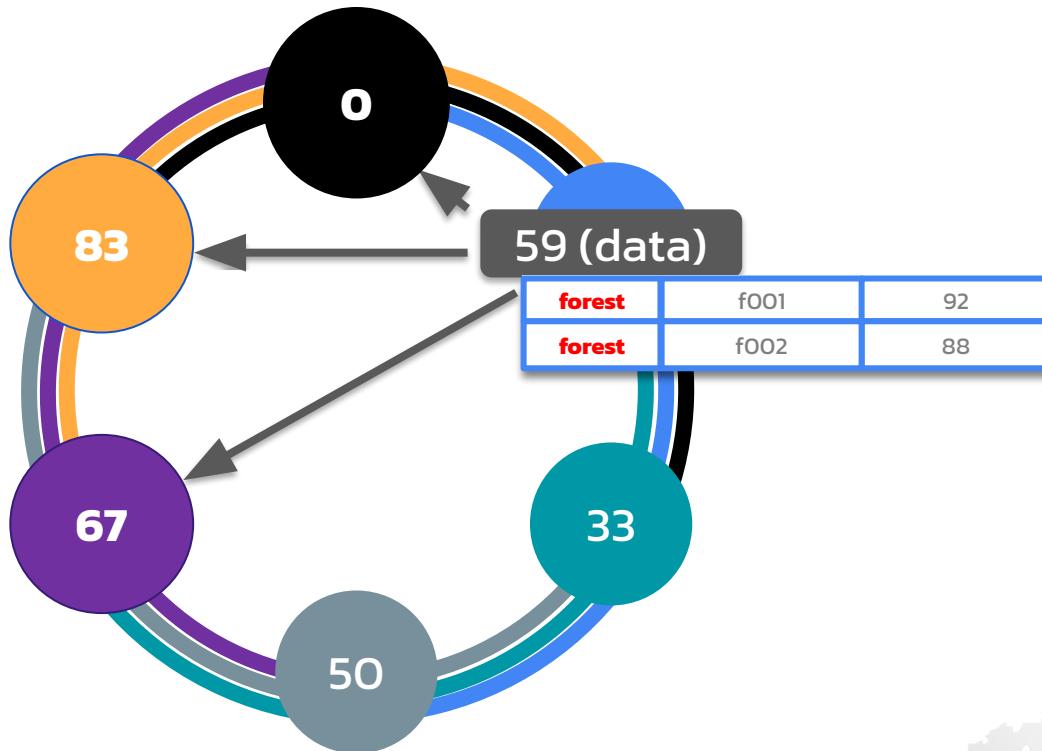
Partitioning + Replication

RF = 3



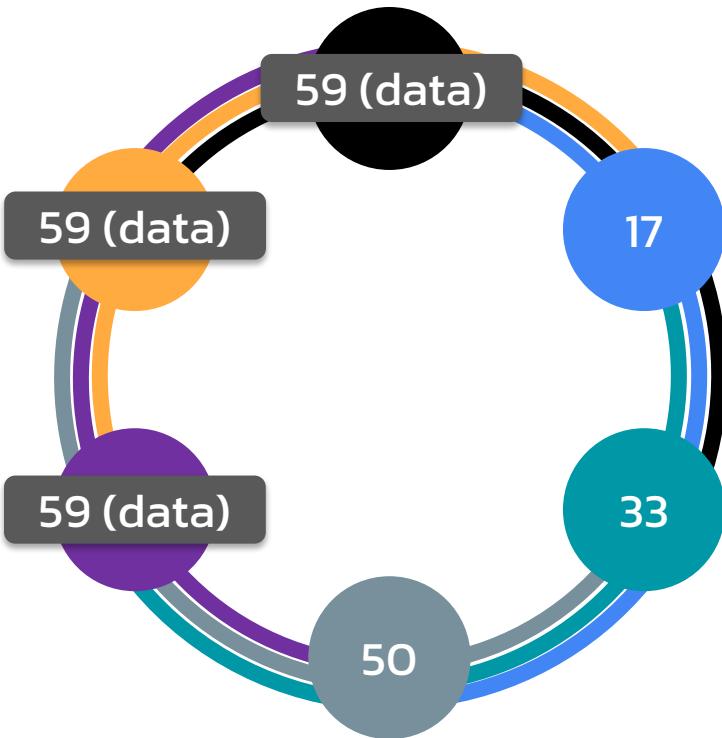
Partitioning + Replication

RF = 3



Partitioning + Replication

RF = 3



Partitioning + Replication

## **IMPORTANT**

Each Cassandra node and even each Cassandra driver knows Data Allocation in a cluster (it's called Token-Aware), so your application can contact literally ANY server and still get the answer.

(but in normal circumstances a driver will contact the replica node)

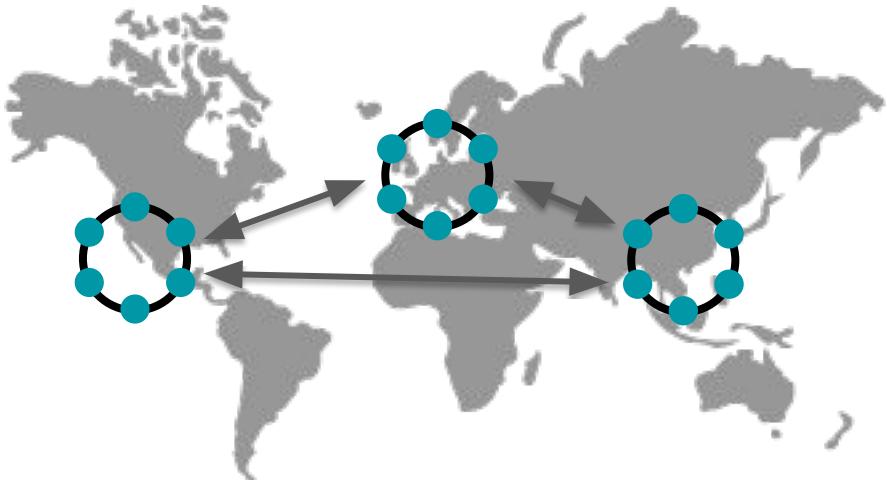




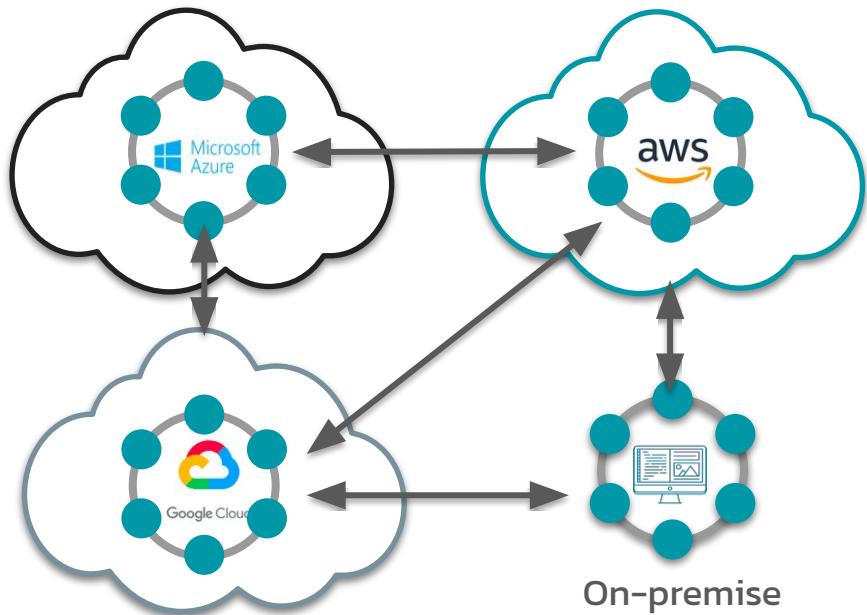
Data is GLOBALLY distributed



## Geographical Distribution



## Hybrid-Cloud and Multi-Cloud



Data is globally distributed



# Lab 1

## Database Setup

[github.com/datastaxdevs/  
workshop-cassandra-fundamentals](https://github.com/datastaxdevs/workshop-cassandra-fundamentals)



- ✓ Create a new database
- ✓ Wake up an existing hibernated database





#### Free to Use

Up to 80GB storage and/or 20 million operations monthly.



#### Serverless

Lower your costs by running Cassandra clusters only when needed.



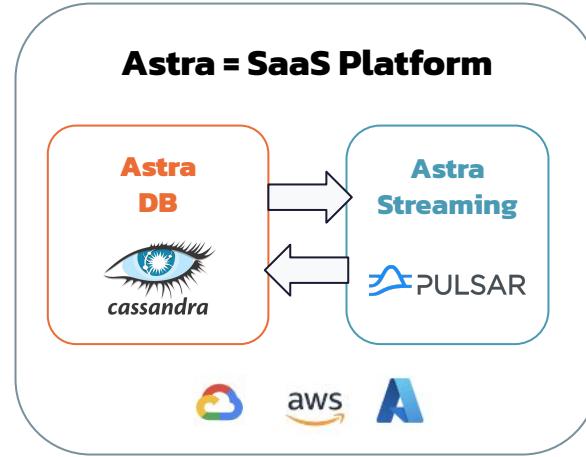
#### No Operations

Eliminate the overhead to install, operate, and scale Cassandra.



#### Data APIs

Work natively with Document (JSON), REST, GraphQL and gRPC APIs.



#### Global Scale

Put your data where you need it without compromising performance, availability or accessibility.



#### End-to-End Security

Secure connect with VPC peering and Private Link. Bring your own encryption key management. SAML SSO secure account access.



#### Zero Lock-in

Deploy on AWS, GCP or Azure and keep compatibility with open-source Cassandra.



#### Relational Indexes

Storage Attached Indexing (SAI) lets you query tables using any columns.



Astra = Cassandra As a Service++



# 01



Why Cassandra ?

# 02

Introduction to  
Apache Cassandra™

# 03

Distributed Architecture

# 04

CAP Theorem

# 05

Tables  
and Partitions

# 06

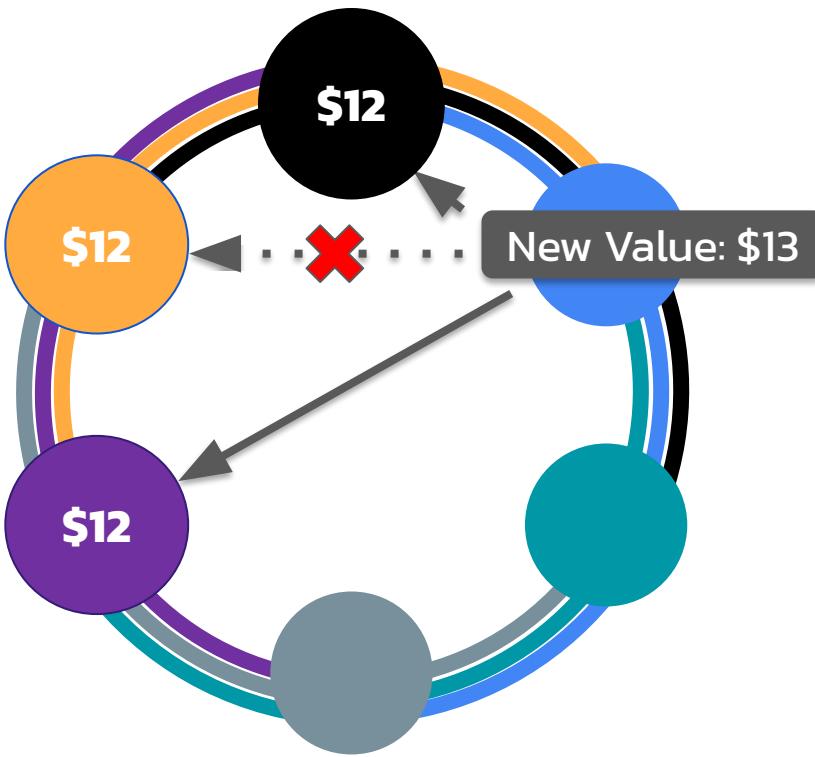
What's next?  
Homework, Next Session



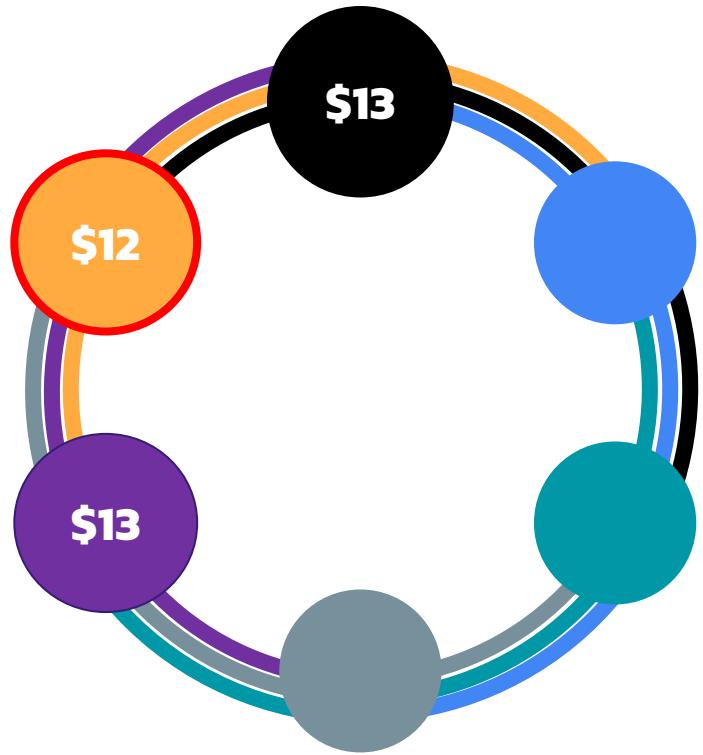
Agenda



# What is the biggest problem of replication?

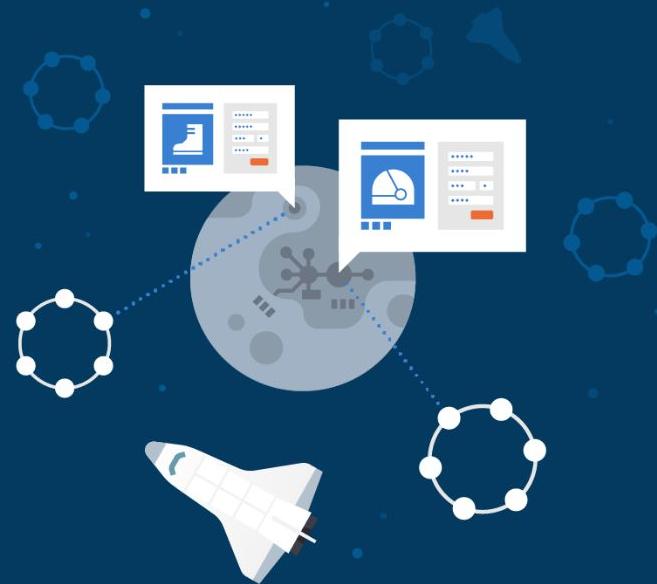


Potential Inconsistency

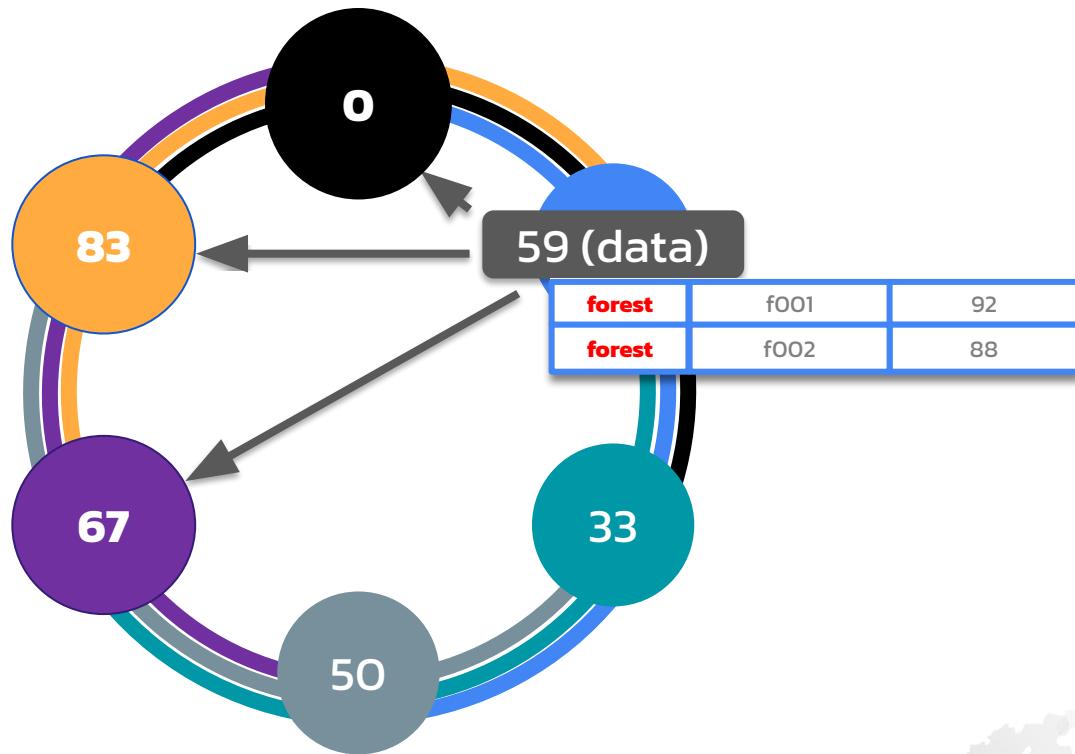


**Potential Inconsistency**

# Cassandra Layered Self-Defence

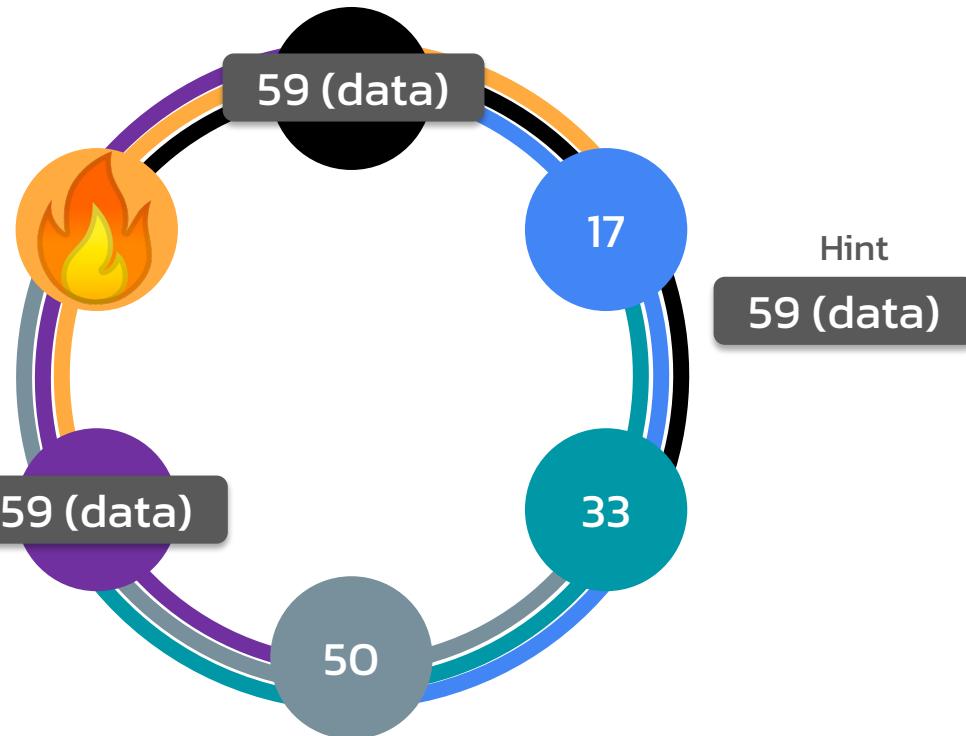


RF = 3



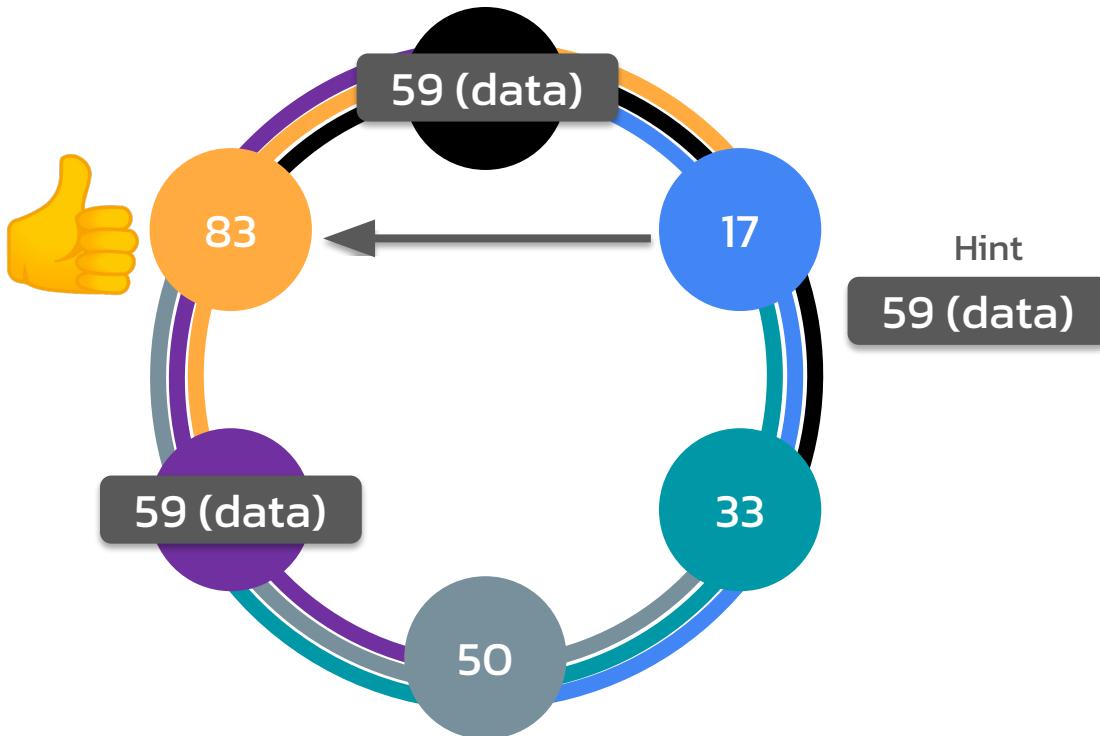
## Hinted Hand-offs

RF = 3



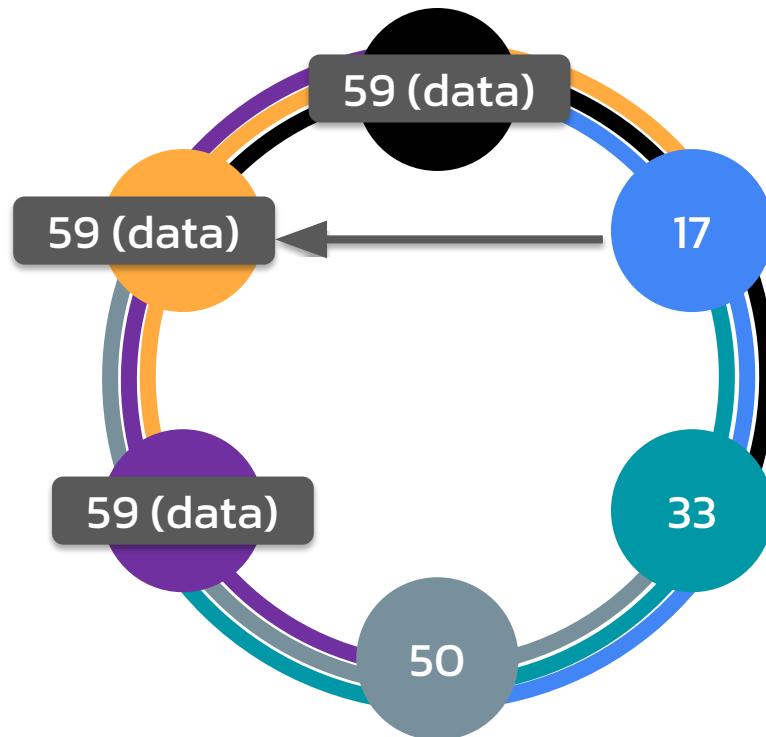
## Hinted Hand-offs

RF = 3



## Hinted Hand-offs

RF = 3



## Hinted Hand-offs

Besides **hinted hand-offs**, very are few more tools to keep consistency under control:

- Consistency Levels      *Developer job*
- Repair on Read            *Cassandra job*
- Repairs                    *Admin job*



**Consistency**



# 01



Why Cassandra ?

# 02

Introduction to  
Apache Cassandra™

# 03

Distributed Architecture

# 04

CAP Theorem

# 05

Tables  
and Partitions

# 06

What's next?  
Homework, Next Session



Agenda

CAP Theorem operates three features:

1. Availability
2. Consistency
3. Partition Tolerance



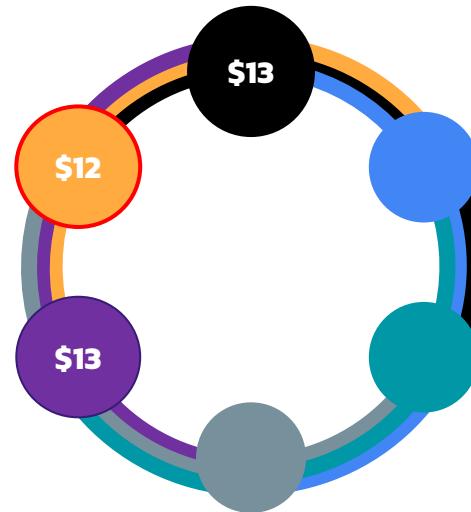
**Availability** basically means “Uptime”. You ask the question, you get the answer. If failure of a single or even multiple servers doesn’t lead to no response (no downtime), your system is **available**.



## CAP Theorem: Availability



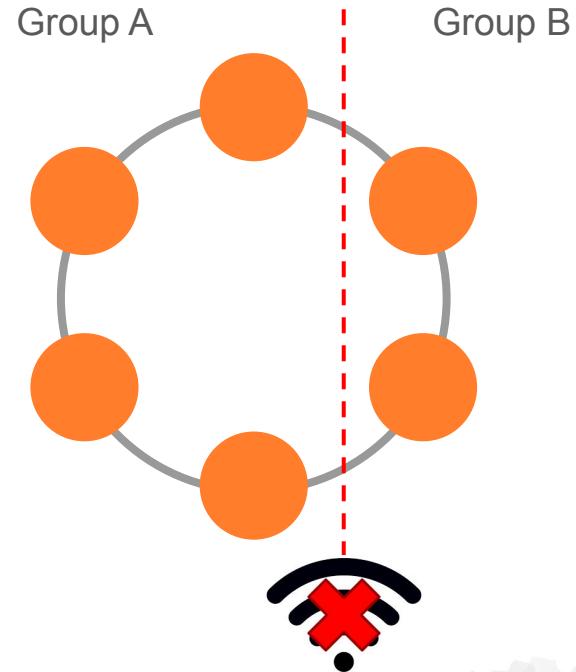
**Consistency** means “no stale data”. You ask for something, you get the most recent value. If you get outdated information, your system is **inconsistent**.



CAP Theorem: Consistency

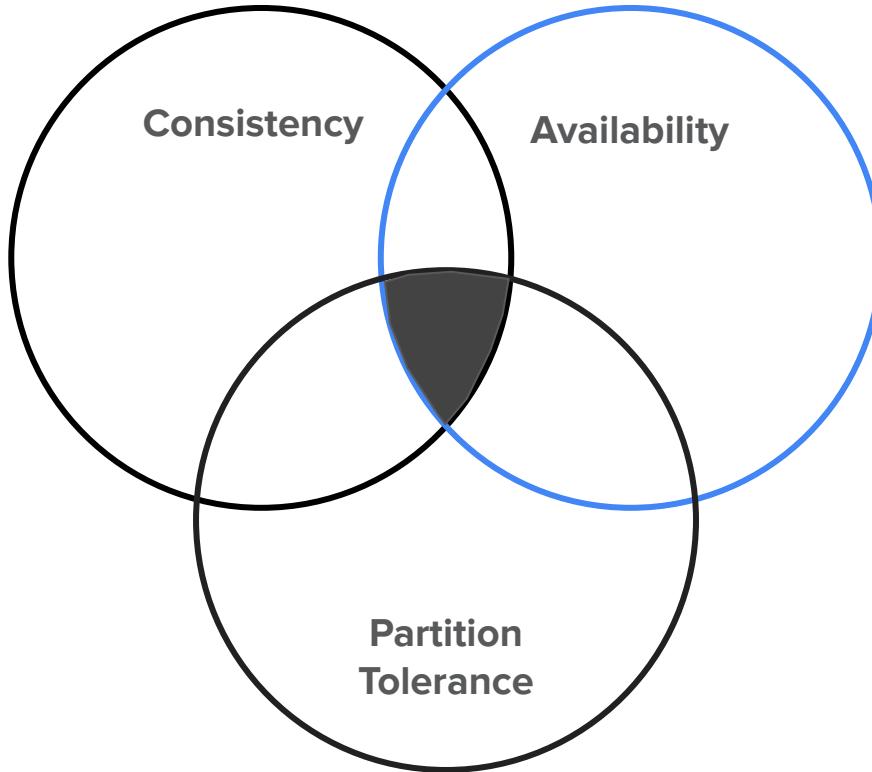


**Partition Tolerance** is the ability of a distributed system to survive “network partitioning”. Network partitioning means that the part of the servers can not reach the second part.



## CAP Theorem: Partition Tolerance

In the distributed environment **in case of emergency** you can have only two guaranteed qualities out of three :(



## CAP Theorem

Cassandra is configurable consistent. In any moment of the time, for any particular query you can set the Consistency Level you require to have. It defines how many **CONFIRMATIONS** you'll wait before the response is dispatched;

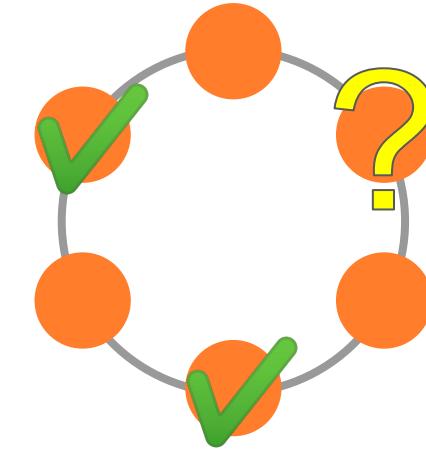
```
PreparedStatement pstmt = session.prepare(  
    "INSERT INTO product (sku, description) VALUES (?, ?)"  
);  
pstmt.setConsistencyLevel(ConsistencyLevel.ONE);
```

```
cqlsh> CONSISTENCY
```

```
Current consistency level is QUORUM.
```

```
cqlsh> CONSISTENCY ALL
```

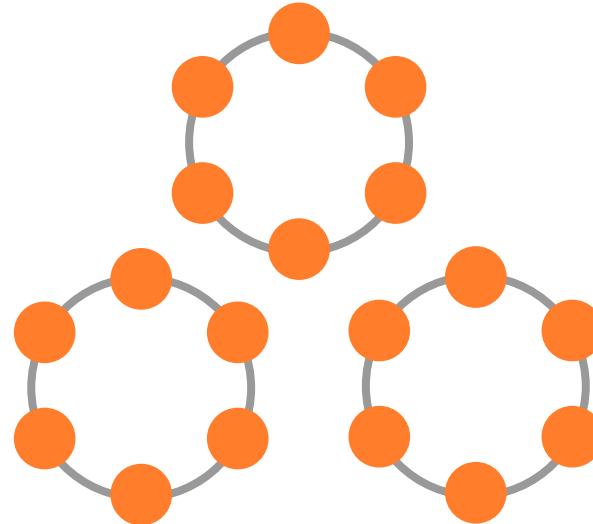
```
Consistency level set to ALL.
```



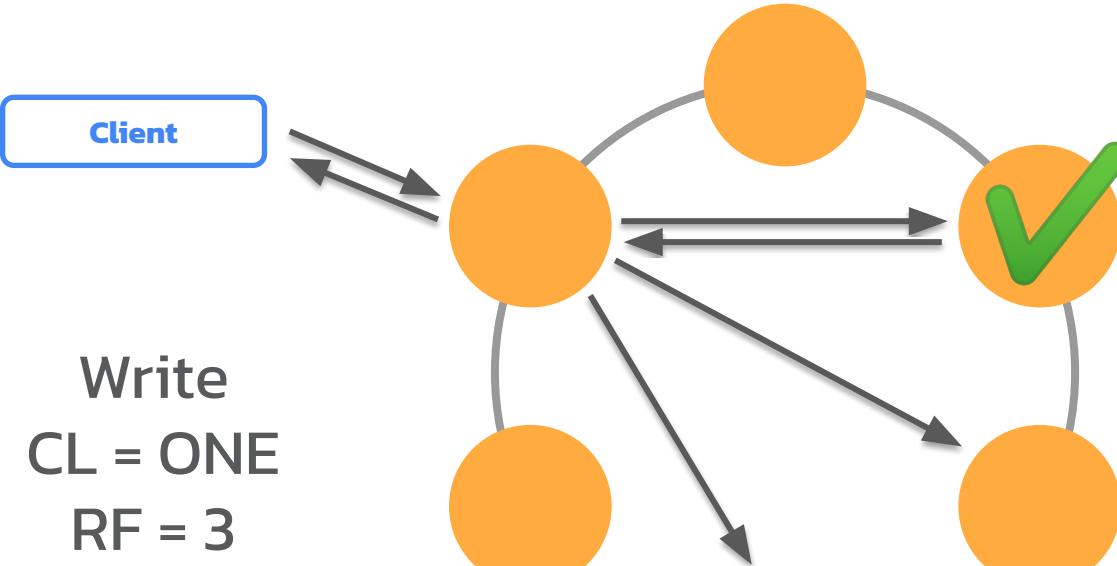
Is Cassandra AP or CP?



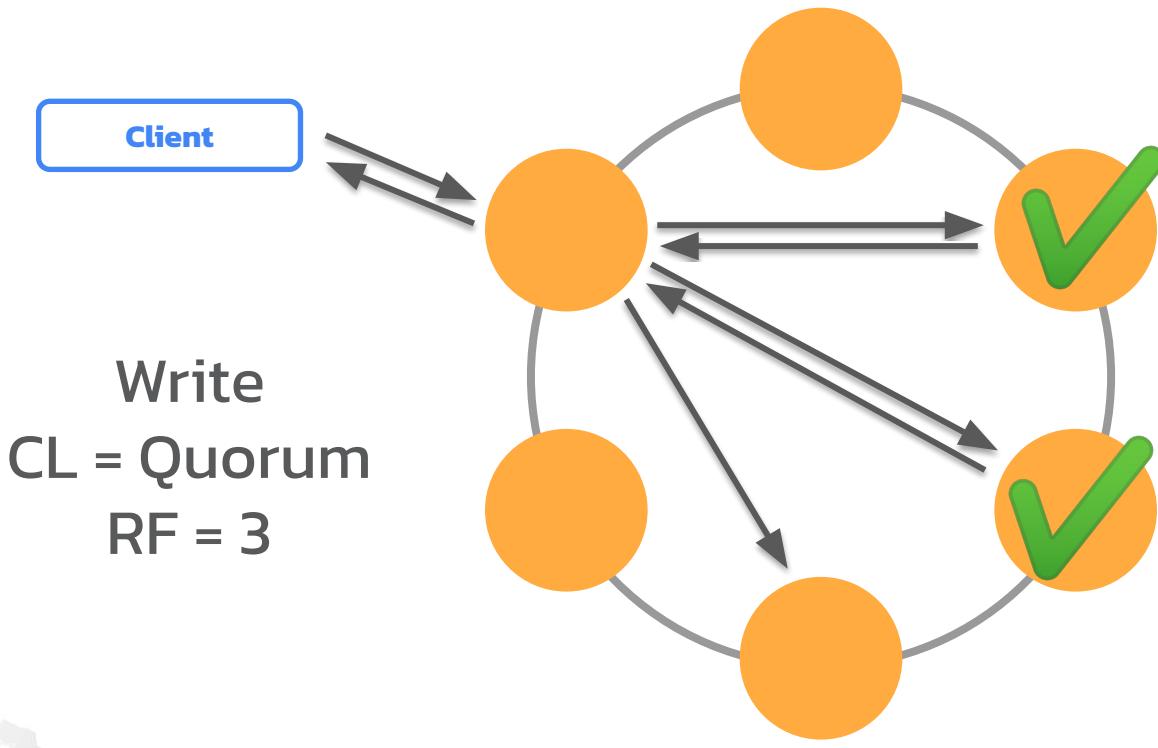
- ANY
- ONE
- LOCAL\_ONE
- TWO, THREE
- QUORUM
- LOCAL\_QUORUM
- EACH\_QUORUM
- ALL



## Query Consistency Levels



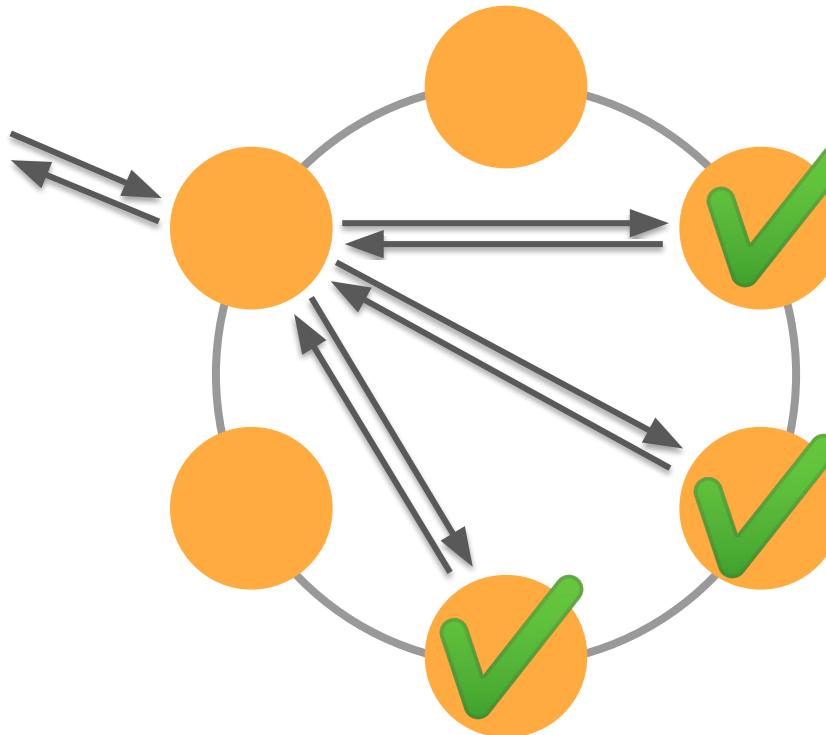
Consistency Level One



Consistency Level Quorum

**Client**

**Write**  
**CL = ALL**  
**RF = 3**



**Consistency Level ALL**



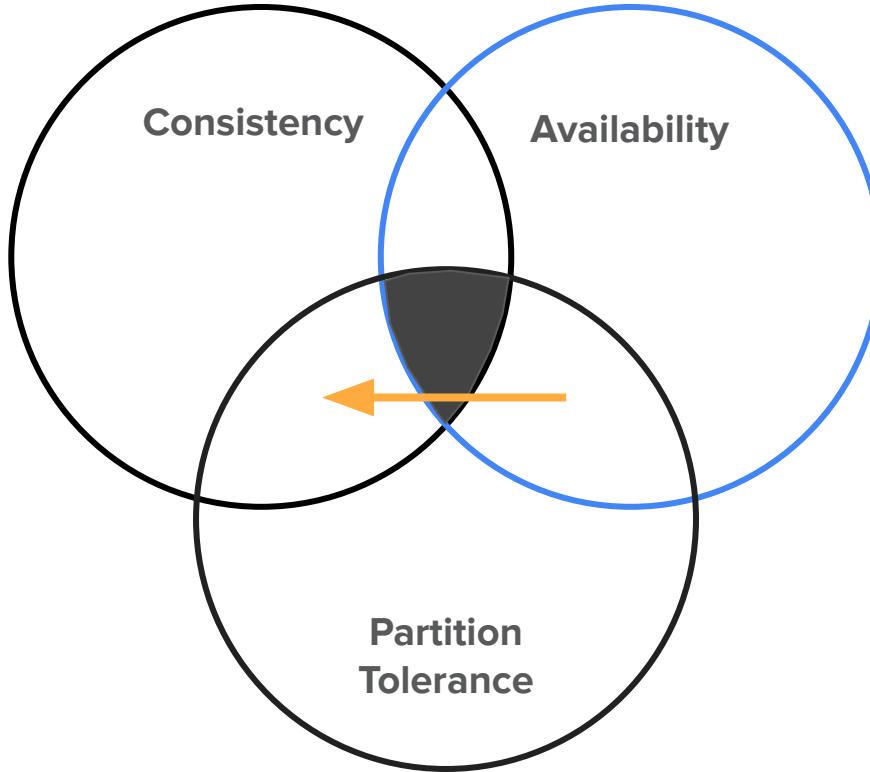
Client

Writes  
CL = ALL  
RF = 3

IT'S A TRAP!

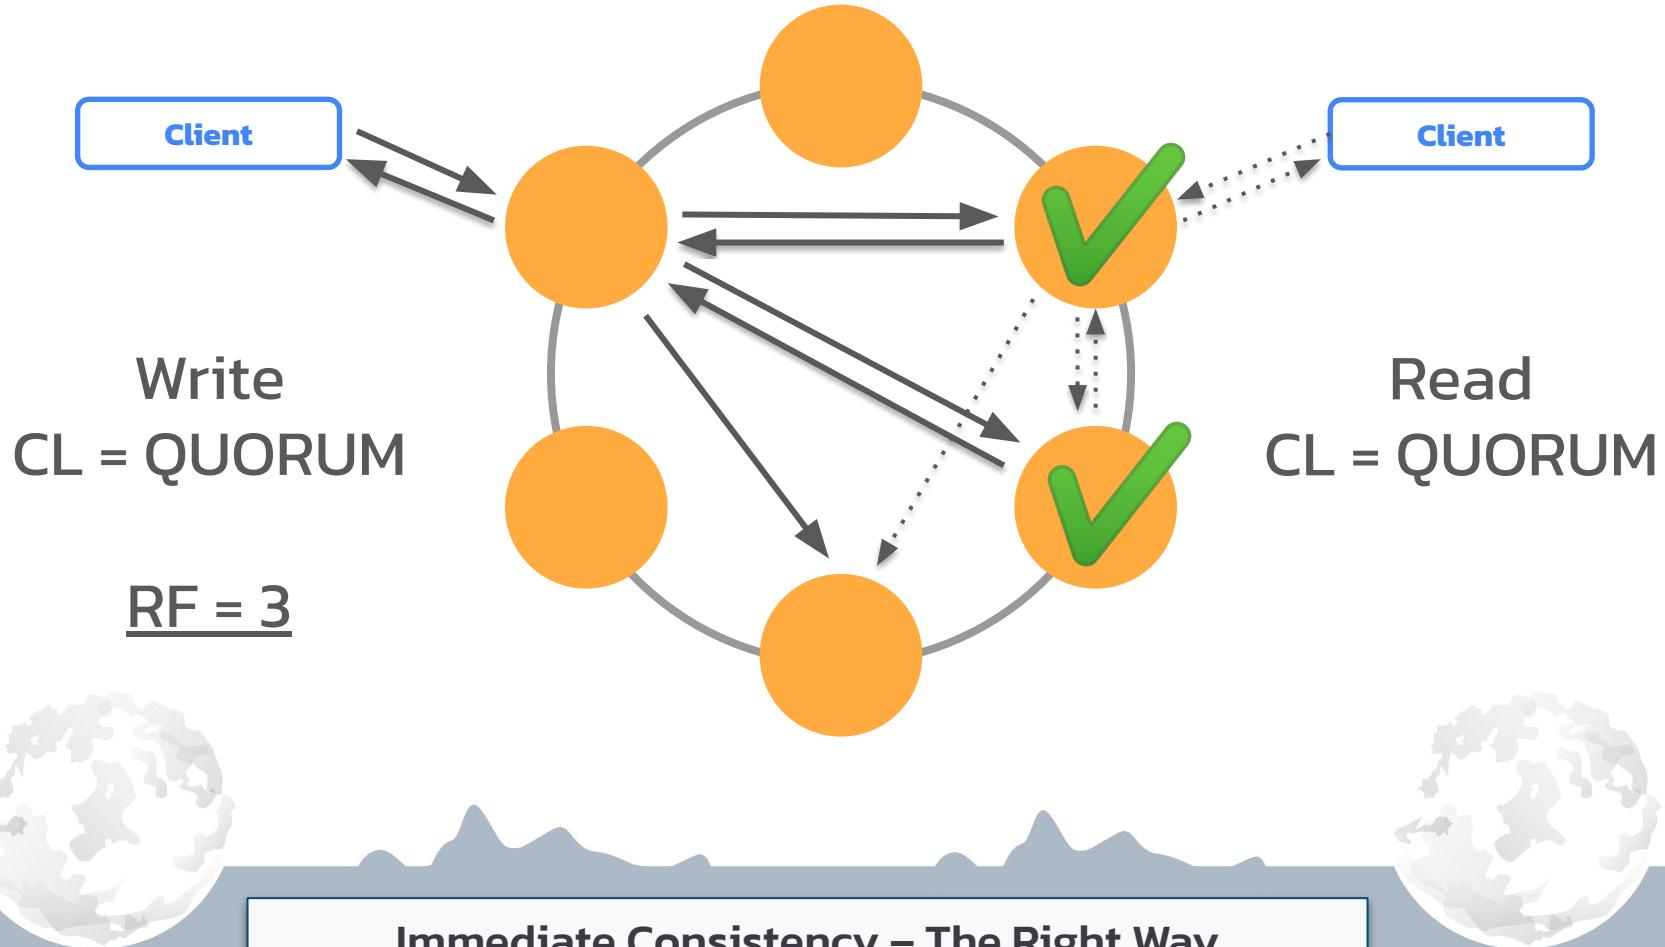
CAP THEOREM IS STILL HERE!

Consistency Level ALL



**CL = ALL**

CL Write + CL Read > RF → Immediate Consistency





# Lab 2

## Create Tables

[github.com/datastaxdevs/  
workshop-cassandra-fundamentals](https://github.com/datastaxdevs/workshop-cassandra-fundamentals)



# 01



Why Cassandra ?

# 02

Introduction to  
Apache Cassandra™

# 03

Distributed Architecture

# 04

CAP Theorem

# 05

Tables  
and Partitions

# 06

What's next?  
Homework, Next Session



Agenda



An intersection of a row  
and a column, stores data.

**foo1**



Data Structure: a Cell





A single, structured  
data item in a table.

<b>forest</b>	<b>f001</b>	<b>92</b>
---------------	-------------	-----------



Data Structure: a Row



A group of rows having the same partition token, a base unit of access in Cassandra.

IMPORTANT: stored together, all the rows are guaranteed to be neighbors.

<b>forest</b>	<b>f001</b>	<b>92</b>
<b>forest</b>	<b>f002</b>	<b>88</b>
<b>forest</b>	<b>f003</b>	<b>90</b>
<b>sea</b>	<b>s001</b>	<b>45</b>



Data Structure: a Partition



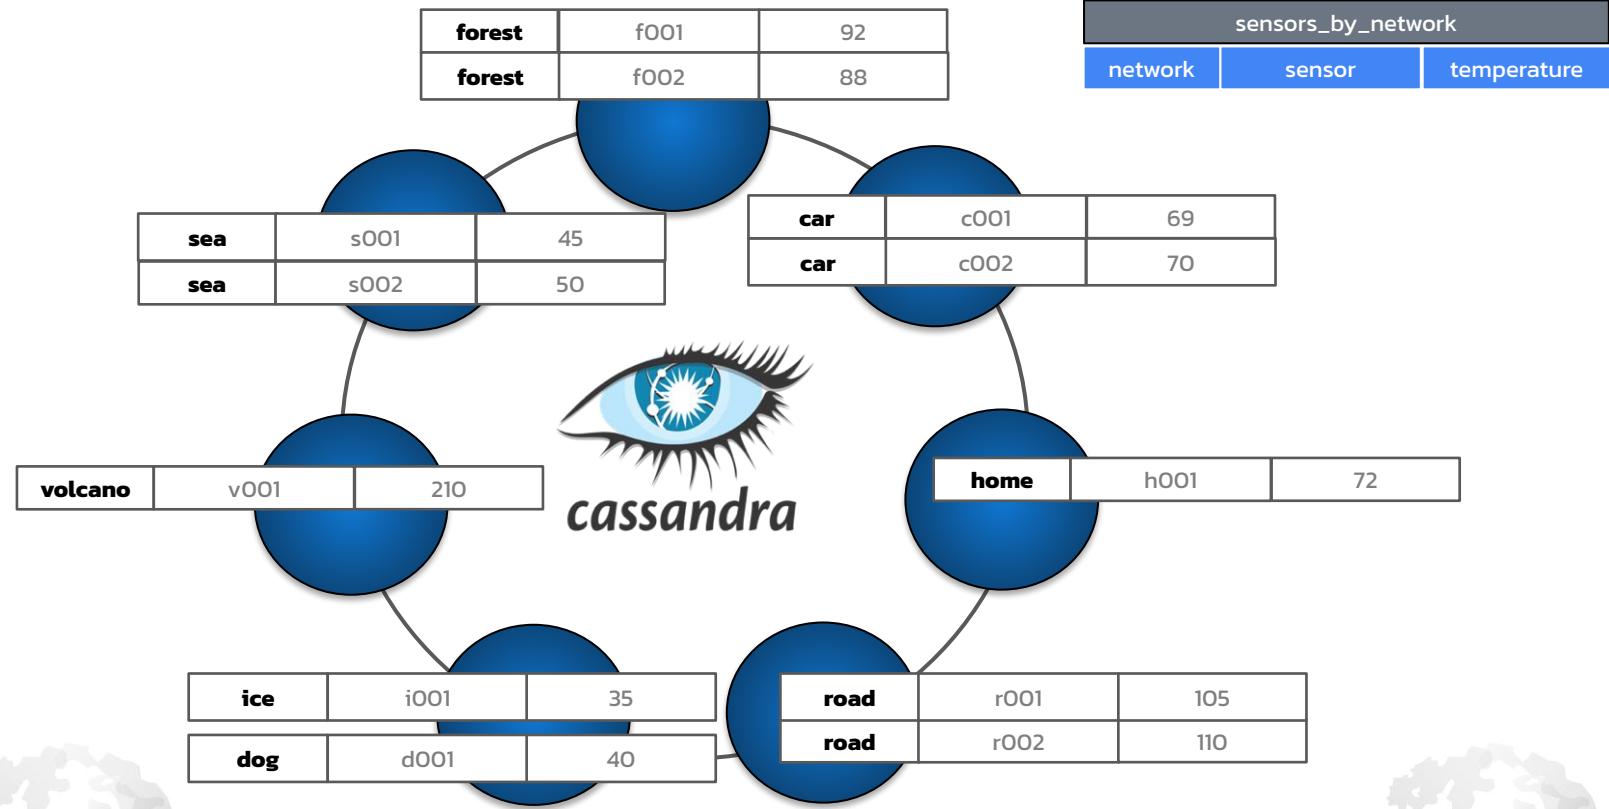
A group of columns and rows storing partitions.

network	sensor	temperature
<b>forest</b>	f001	92
<b>forest</b>	f002	88
<b>volcano</b>	v001	210
<b>sea</b>	s001	45
<b>sea</b>	s002	50
<b>home</b>	h001	72
<b>car</b>	c001	69
<b>car</b>	c002	70
<b>dog</b>	d001	40
<b>road</b>	r001	105
<b>road</b>	r002	110
<b>ice</b>	i001	35



## Data Structure: a Table





Data is organized as Distributed tables

An identifier for a partition.  
Consists of at least one column,  
may have more if needed.

## PARTITIONS TABLE.

Defines data distribution strategy

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
);
```

Partition key

Examples:

```
PRIMARY KEY ((sensor), timestamp);
```

```
PRIMARY KEY ((sensor), date, timestamp);
```

```
PRIMARY KEY ((sensor, timestamp));
```

Partition Key

Used to **ensure uniqueness** and **sorting order**. Optional.

Defines physical ordering on disk

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Clustering columns

PRIMARY KEY ((**sensor**));

Not Unique

PRIMARY KEY ((**sensor**), timestamp);

Not filter on dates

PRIMARY KEY ((**sensor**), value, timestamp);

Not sorted

PRIMARY KEY ((**sensor**), date, timestamp);



Clustering Columns

An identifier for a row. Consists of at least one Partition Key and zero or more Clustering Columns.

## UNIQUELY IDENTIFIES A ROW

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Primary key

Examples:

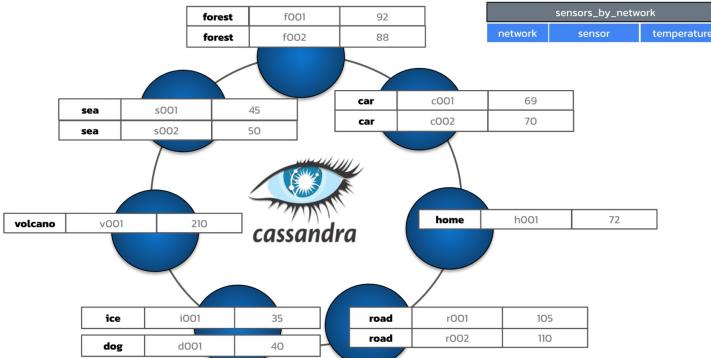
```
PRIMARY KEY ((sensor), timestamp);
```

```
PRIMARY KEY ((sensor, date), timestamp);
```

Primary Key

## Partition Keys

define data distribution over the cluster



## Clustering Columns

define how data is physically stored (written on disk)

```
{
  "partition": {
    "key": [ "y58dbfa0-b87f-11ea-ab9b-f55dd84a27b8" ],
    "position": 131
  },
  "rows": [
    {
      "type": "row",
      "position": 261,
      "liveness_info": { "tstamp": "2021-02-26T14:59:18.322385Z" },
      "cells": [
        { "name": "createdt", "value": "2020-06-28 11:00:00.000" },
        { "name": "publish_date", "value": "2020-06-28 11:00:00.000Z" },
        { "name": "title", "value": "Title02" },
        { "name": "user_email", "value": "placeholder@email.com" },
        { "name": "user_name", "value": "name01" }
      ]
    },
    {
      "type": "row",
      "position": 262,
      "liveness_info": { "tstamp": "2021-02-26T14:59:17.893212Z" },
      "cells": [
        { "name": "createdt", "value": "2020-06-28 11:00:00.000" },
        { "name": "publish_date", "value": "2020-06-28 11:00:00.000Z" },
        { "name": "title", "value": "Title01" },
        { "name": "user_email", "value": "placeholder@email.com" },
        { "name": "user_name", "value": "name01" }
      ]
    }
  ]
}
```

Primary Key = Ensure Uniqueness = PK + CC

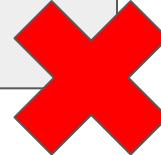
## Important:

Once created, the data model cannot be changed! You will need new tables and migration. Stay lazy, design it right in advance!

```
ALTER TABLE temperatures_by_sensor  
ADD season TEXT;
```



```
ALTER TABLE temperatures_by_sensor  
DROP PRIMARY KEY  
ADD PRIMARY KEY (sensor, timestamp)
```



Corollary: Schema Immutability



# Lab 3

## CRUD Operations

[github.com/datastaxdevs/  
workshop-cassandra-fundamentals](https://github.com/datastaxdevs/workshop-cassandra-fundamentals)

That's step 1 of your homework!



# 01



Why Cassandra ?

# 02

Introduction to  
Apache Cassandra™

# 03

Distributed Architecture

# 04

CAP Theorem

# 05

Tables  
and Partitions

# 06

What's next?  
Homework, Next Session



Agenda

# Homework



## Intro to Cassandra Homework

Welcome and thank you! Here you can submit your homework for the datastax developers "Intro to Cassandra" workshop. In case of any questions please contact organisers at <https://dtsx.io/aleks> or just send an email to [aleksandr.volochnev@datastax.com](mailto:aleksandr.volochnev@datastax.com)

- Workshop materials: <https://github.com/datastaxdevs/workshop-intro-to-cassandra>
- Discord chat: <https://dtsx.io/discord>

[cedrick.lunven@datastax.com](mailto:cedrick.lunven@datastax.com) [Switch account](#)



The name and photo associated with your Google account will be recorded when you upload files and submit this form. Only the email you enter is part of your response.

\* Required



# Thank you!



**Sponsored by DataStax**