




# Apache Cassandra® Hands-On Workshop

**Sponsored by DataStax**


# Hands-On Housekeeping

Nothing to install

## Source code + exercises + slides



## IDE



## Database + Api + Streaming



DataStax

Astra DB

**01**



**Getting started with  
Apache Cassandra®**

**02**

**Hands-on lab 1:  
Create a database**

**03**

**Data definition and  
manipulation with CQL**

**04**

**Hands-on lab 2:  
Create, populate and query**

**05**

**Application development  
in Java and Python**

**06**

**Hands-on lab 3:  
Create an app with C\* driver**

**The most scalable  
NoSQL database**

# Apache Cassandra

NoSQL Distributed Decentralised Database Management System



- Distributed, real-time, OLTP, NoSQL
- Linear Scalability
- High Availability
- Geographical Distribution
- Platform Agnostic
- Vendor Independent

# Cassandra Biggest Users (and Developers)

## Apache Cassandra @ Netflix

- 98% of streaming data is stored in Apache Cassandra
- Data ranges from customer details to viewing history to billing and payments
- Foundational datastore for serving millions of operations per second

- 30 million ops/sec on most active single cluster
- 500 TB most dense single cluster
- 9216 CPUs in biggest cluster

$O(100)$  Clusters  
 $O(10000)$  Instances  
 $O(10,000,000)$  Replications per second  
 $O(100,000,000)$  Operations per second  
 $O(1,000,000,000,000,000)$  Petabytes of data

[dtsx.io/cassandra-at-netflix](https://dtsx.io/cassandra-at-netflix)


## Apple Scale

- 160K+ Apache Cassandra instances
- 100+ PB stored
- Several million ops / sec
- 1000s of clusters



And many others...

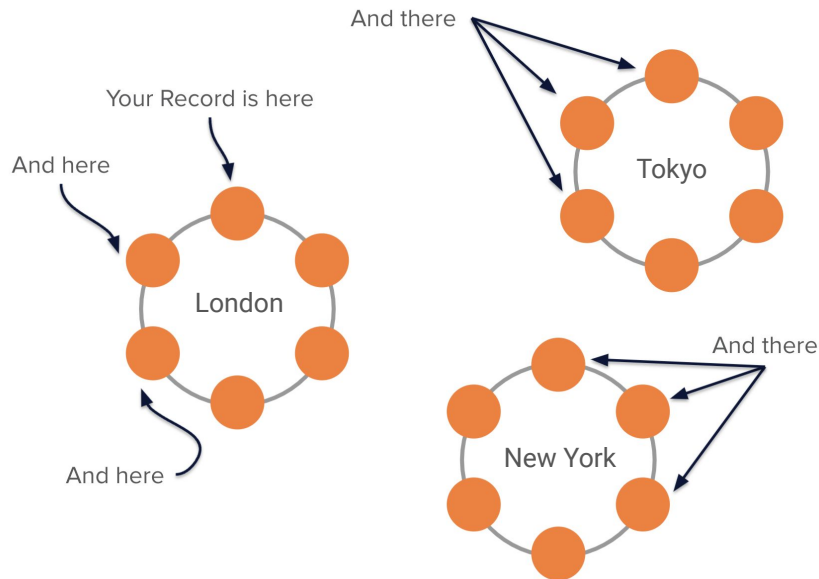
# Linear Scalability



# High Availability

Replication, Decentralisation, and Topology-Aware Placement Strategy take care of possible downtimes:

- Multiple Live Replicas
- No Single Point of Failure
- Network topology-aware data placement
- Client-side Smart Reconnection and Strong Retry Mechanism






# Geographical Distribution


Cassandra's trademark is multi-datacenter deployments, granting you an exceptional capability for disaster tolerance while keeping your data close to your clients - worldwide.

All DCs are active (available for both writes and reads)!



# Platform Agnostic

Apache Cassandra is **not bound to any platform** or service provider, helping you build hybrid-cloud and multi-cloud solutions with ease.




# Vendor Independent

Cassandra doesn't belong to any of commercial vendors but controlled by a non-profit Open Source **Apache Software Foundation**, already familiar to you by *Hadoop*, *Spark*, *Kafka*, *Zookeeper*, *Maven* and many other projects.



# **Data distribution, replication and consistency**

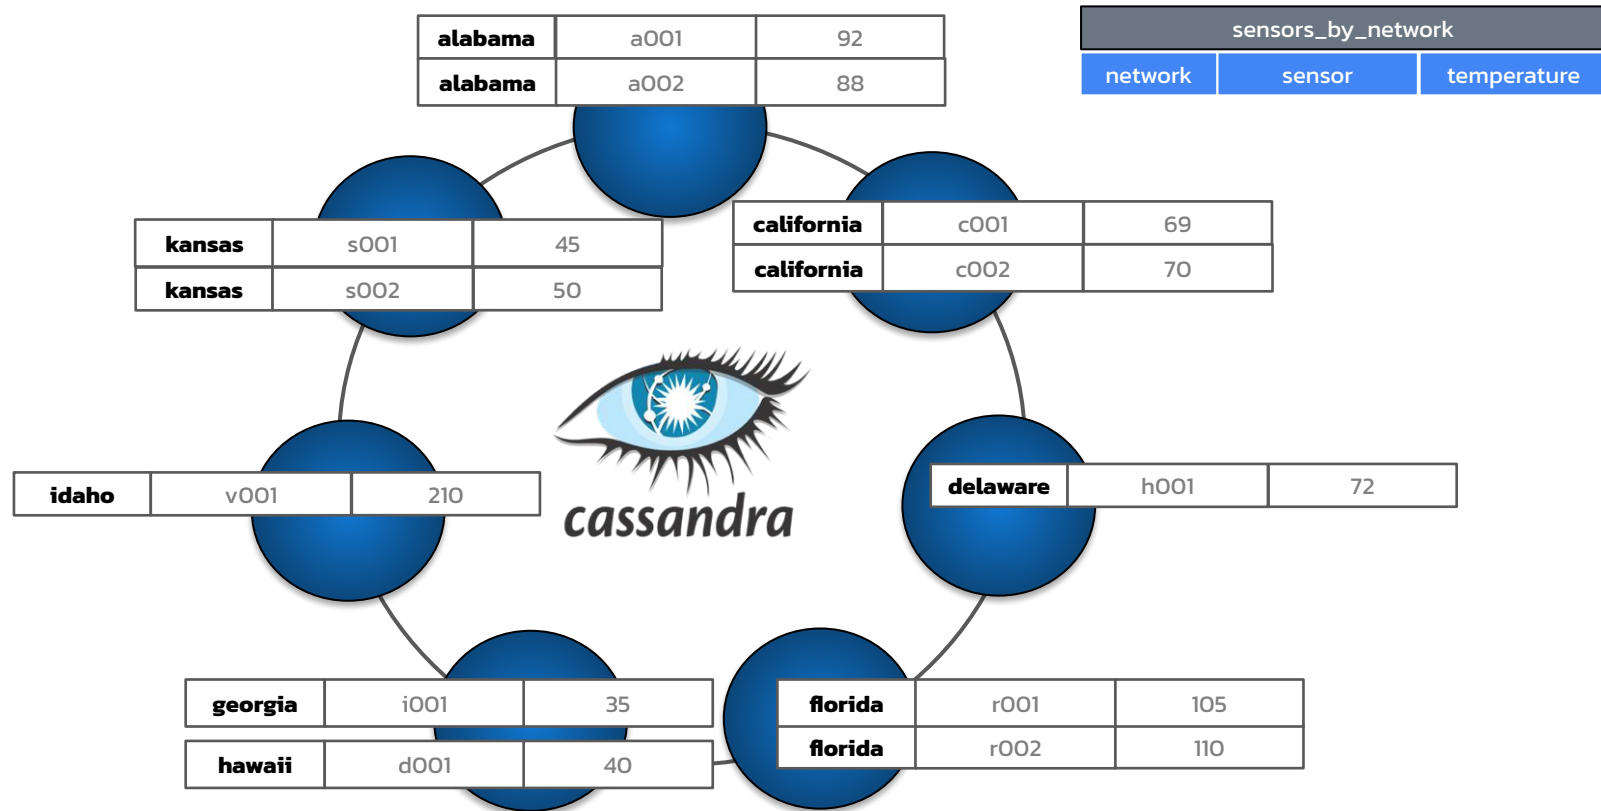
# Partition key and partitions



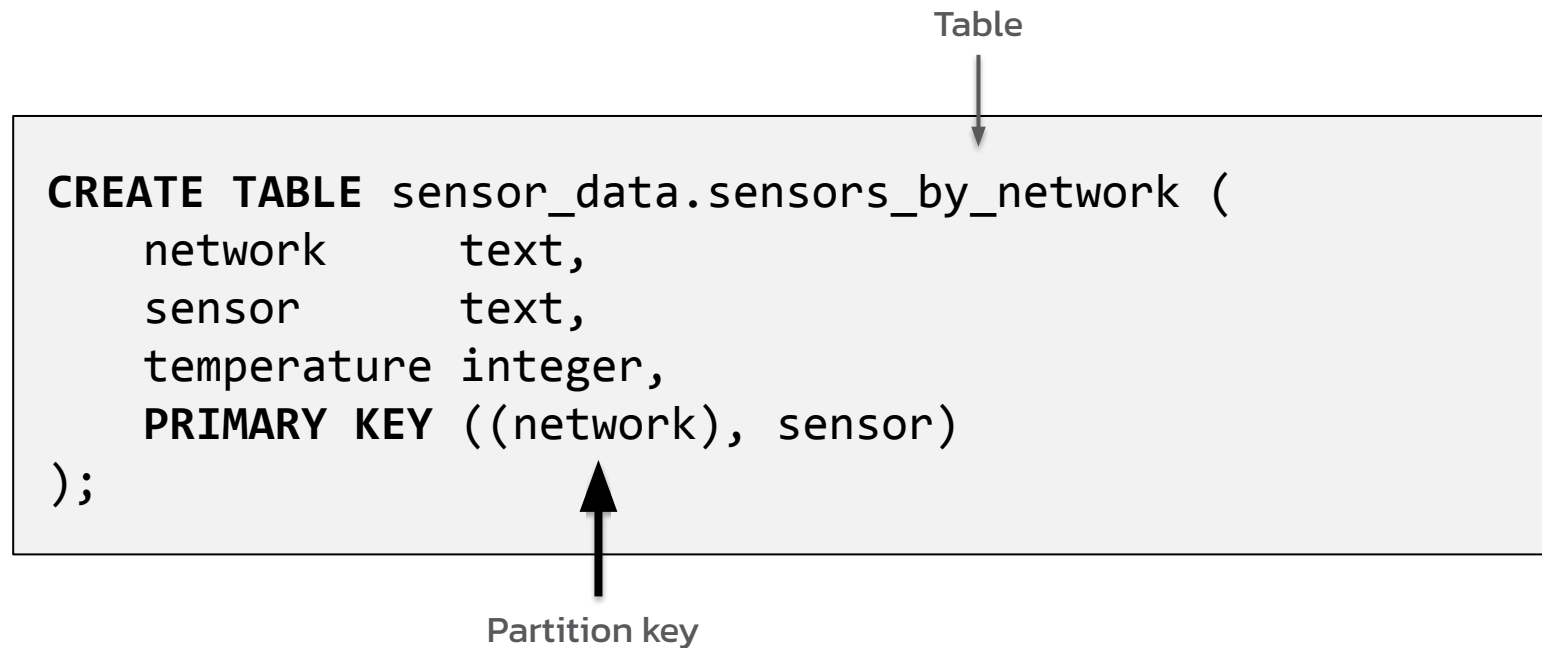
sensors_by_network		
network	sensor	temperature
alabama	a001	92
alabama	a002	88
california	c001	210
delaware	d001	45
delaware	d002	50
florida	f001	72
georgia	g001	69
georgia	g002	70
hawaii	h001	40
idaho	i001	105
idaho	i002	110
kansas	k001	35

  
**Partition Key**

# Data distribution



# Partition key definition in CQL



# Internals: Partitioning and Token Ranges

network	sensor	Value
alabama	f001	92
alabama	f002	88
idaho	s001	45
idaho	s002	50
hawaii	r001	105
hawaii	r002	110

*Partition Keys*



Network	Sensor	Value
<b>59</b>	f001	92
<b>59</b>	f002	88
<b>12</b>	s001	45
<b>12</b>	s002	50
<b>45</b>	r001	105
<b>45</b>	r002	110

*Tokens*

Cassandra Nodes






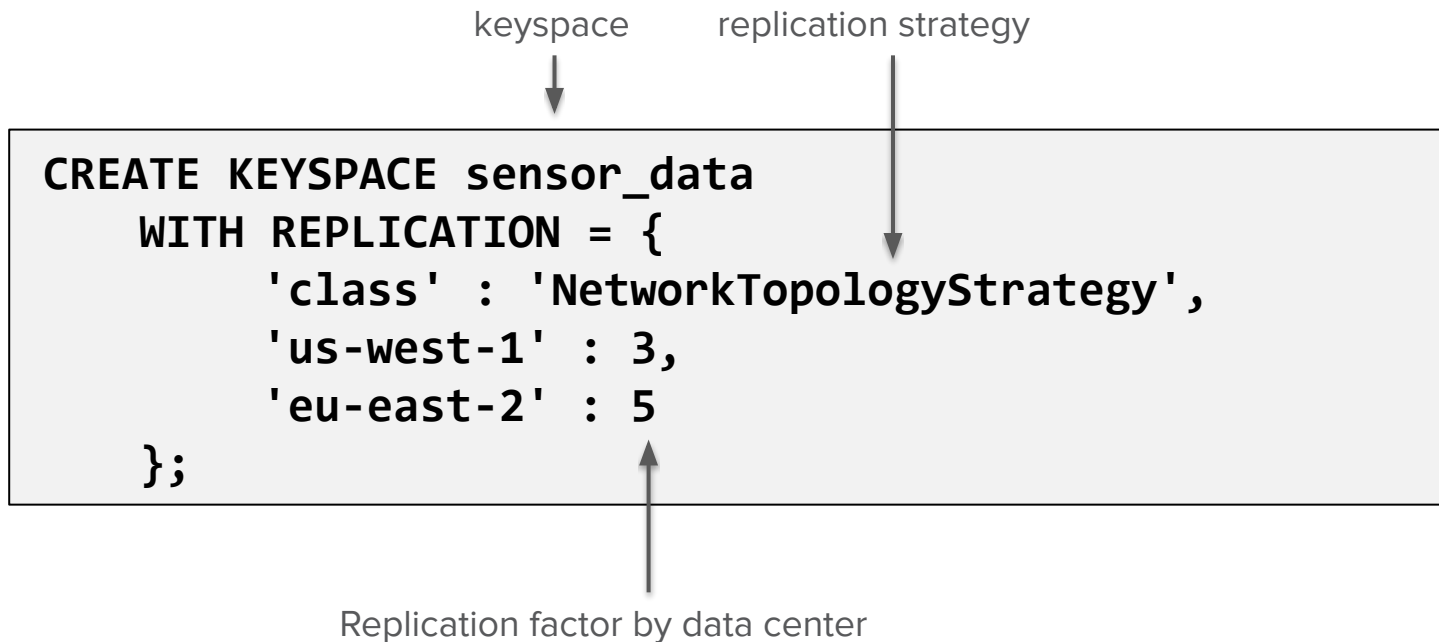
# Replication Factor

RF = ?

Replication Factor  
means the number  
of nodes used to  
store each partition




# Replication is defined per keyspace



# RF = 1

## RF = 1


Replication Factor 1  
means that every  
partition is stored  
on 1 node



# RF = 2

## RF = 2

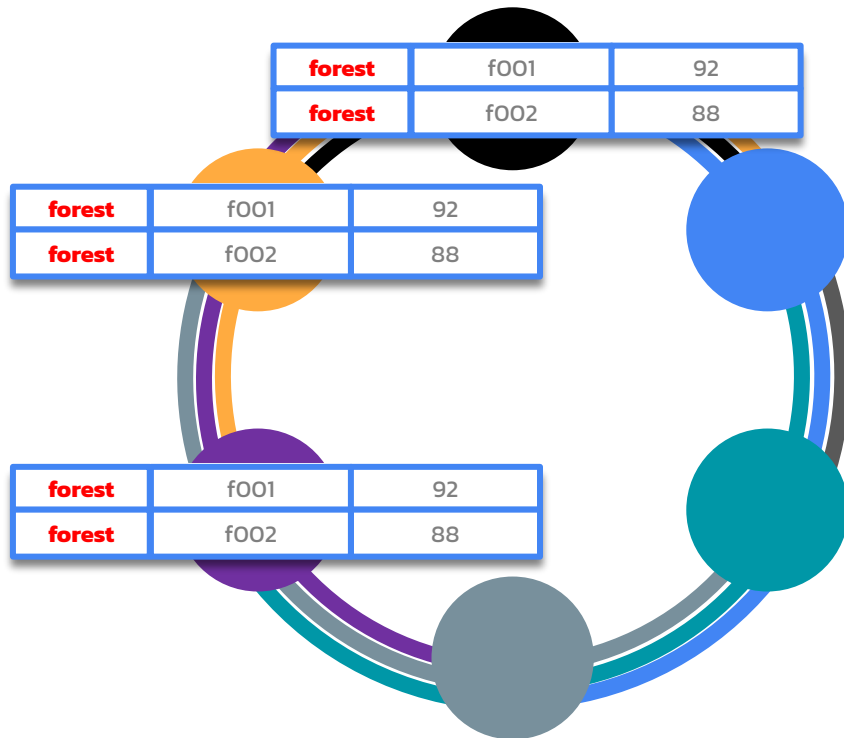
Replication Factor 2  
means that every  
partition is stored  
on 2 nodes



# RF = 3

## RF = 3


Replication Factor 3  
means that every  
partition is stored  
on 3 nodes



# Replication, consistency and availability

**But what if ...**


RF = 3



forest	f001	92
forest	f002	88


# CAP Theorem

Any distributed data store can provide only **two of the three** guarantees




# Tunable Consistency and Consistency Levels

- ANY
- ONE
- TWO
- THREE
- QUORUM
- LOCAL\_ONE
- LOCAL\_QUORUM
- EACH\_QUORUM
- ALL






# Consistency level ONE




# Consistency level QUORUM



# Immediate Consistency

CL Write + CL Read > RF → Immediate Consistency



**01**



**Getting started with  
Apache Cassandra®**

**02**

**Hands-on lab 1:  
Create a database**

**03**

**Data definition and  
manipulation with CQL**

**04**

**Hands-on lab 2:  
Create, populate and query**

**05**

**Application development  
in Java and Python**

**06**

**Hands-on lab 3:  
Create an app with C\* driver**

# Lab 1

Create a Cassandra database instance in the cloud  
Start here: [astra.datastax.com](https://astra.datastax.com)

- ✓ Create database **workshops** and keyspace **sensor\_data**
- ✓ Generate and save an application token

*Follow a demo by the instructor to complete the steps*

# Astra = Cassandra As a Service++



## Free to Use

Up to 80GB storage and/or 20 million operations monthly.



## Serverless

Lower your costs by running Cassandra clusters only when needed.



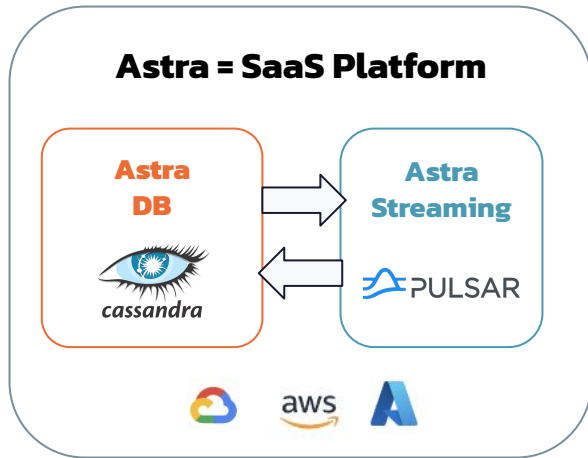
## No Operations

Eliminate the overhead to install, operate, and scale Cassandra.



## Data APIs

Work natively with Document (JSON), REST, GraphQL and gRPC APIs.



## Global Scale

Put your data where you need it without compromising performance, availability or accessibility.



## End-to-End Security

Secure connect with VPC peering and Private Link. Bring your own encryption key management. SAML SSO secure account access.



## Zero Lock-in

Deploy on AWS, GCP or Azure and keep compatibility with open-source Cassandra.



## Relational Indexes

Storage Attached Indexing (SAI) lets you query tables using any columns.

**01**



**Getting started with  
Apache Cassandra®**

**02**

**Hands-on lab 1:  
Create a database**

**03**

**Data definition and  
manipulation with CQL**

**04**

**Hands-on lab 2:  
Create, populate and query**

**05**

**Application development  
in Java and Python**

**06**

**Hands-on lab 3:  
Create an app with C\* driver**

# Table definition and partitioning



# Primary Key, Partition Key, Clustering Key

```
CREATE TABLE sensor_data.temperatures_by_sensor (  
  sensor      TEXT,  
  date        DATE,  
  timestamp   TIMESTAMP,  
  value       FLOAT,  
  PRIMARY KEY ( ( sensor , date ) ,      timestamp )  
 ) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Partition key

Clustering key

**DEFINES DATA DISTRIBUTION  
UNIQUELY IDENTIFIES  
A PARTITION IN A TABLE**

**MULTI-ROW PARTITIONS  
UNIQUELY IDENTIFIES A  
ROW IN A PARTITION**

Primary key

**UNIQUELY IDENTIFIES A ROW IN A TABLE**

# Table Structure $\Rightarrow$ Valid Queries

```
PRIMARY KEY ((sensor, date), timestamp);
```

```
SELECT * FROM temperatures_by_sensor ...
```

```
WHERE sensor = ?;
```

```
WHERE sensor > ?;
```

```
WHERE sensor = ? AND date > ?;
```

```
WHERE sensor = ? AND date = ?;
```

```
WHERE sensor = ? AND date = ? AND timestamp > ?;
```

# Good Partition Rules

- ❖ **Store together what you retrieve together**
- ❖ Avoid big partitions
- ❖ Avoid hot partitions

Q: Show temperature evolution over time for **sensor X** On **Nov 10, 2022**

```
PRIMARY KEY ((sensor, timestamp));
```



```
PRIMARY KEY ((sensor), timestamp);
```



# Good Partition Rules

- ❖ Store together what you retrieve together
- ❖ **Avoid big partitions**
- ❖ Avoid hot partitions

## BUCKETING

```
PRIMARY KEY ((sensor), timestamp);
```



```
PRIMARY KEY ((sensor, month_year), timestamp);
```



- Up to 2 billion cells per partition
- Up to ~100k values in a partition
- Up to ~100MB in a Partition

# Good Partition Rules

- ❖ Store together what you retrieve together
- ❖ Avoid big partitions
- ❖ **Avoid hot partitions**

```
PRIMARY KEY ((date), sensor, timestamp);
```



```
PRIMARY KEY ((date, sensor), timestamp);
```



```
PRIMARY KEY ((sensor, date), timestamp);
```



**Inserts, bulk loading,  
querying capabilities**

# Ingesting Data

- CQL statement `INSERT`
- CQL shell command `COPY FROM`
- Command-line utility `dsbulk`
- Apache Spark with Spark-Cassandra Connector

# Querying Data

```
SELECT [DISTINCT] * |
        select_expression [AS column_name][ , ... ]
FROM    [keyspace_name.] table_name
[WHERE  partition_key_predicate
        [AND clustering_key_predicate]]
[GROUP BY primary_key_column_name][ , ... ]
[ORDER BY clustering_key_column_name ASC|DESC][ , ... ]
[PER PARTITION LIMIT number]
[LIMIT number]
[ALLOW FILTERING]
```



**01**



**Getting started with  
Apache Cassandra®**

**02**

**Hands-on lab 1:  
Create a database**

**03**

**Data definition and  
manipulation with CQL**

**04**

**Hands-on lab 2:  
Create, populate and query**

**05**

**Application development  
in Java and Python**

**06**

**Hands-on lab 3:  
Create an app with C\* driver**

# Lab 2

Create, populate and query Cassandra tables

Start here: [tinyurl.com/cassandra-cql](https://tinyurl.com/cassandra-cql)

- ✓ Understand the data model and queries
- ✓ Create tables and run queries in Cassandra

*Follow a demo by the instructor to complete the steps*

**01**



**Getting started with  
Apache Cassandra®**

**02**

**Hands-on lab 1:  
Create a database**

**03**

**Data definition and  
manipulation with CQL**

**04**

**Hands-on lab 2:  
Create, populate and query**

**05**

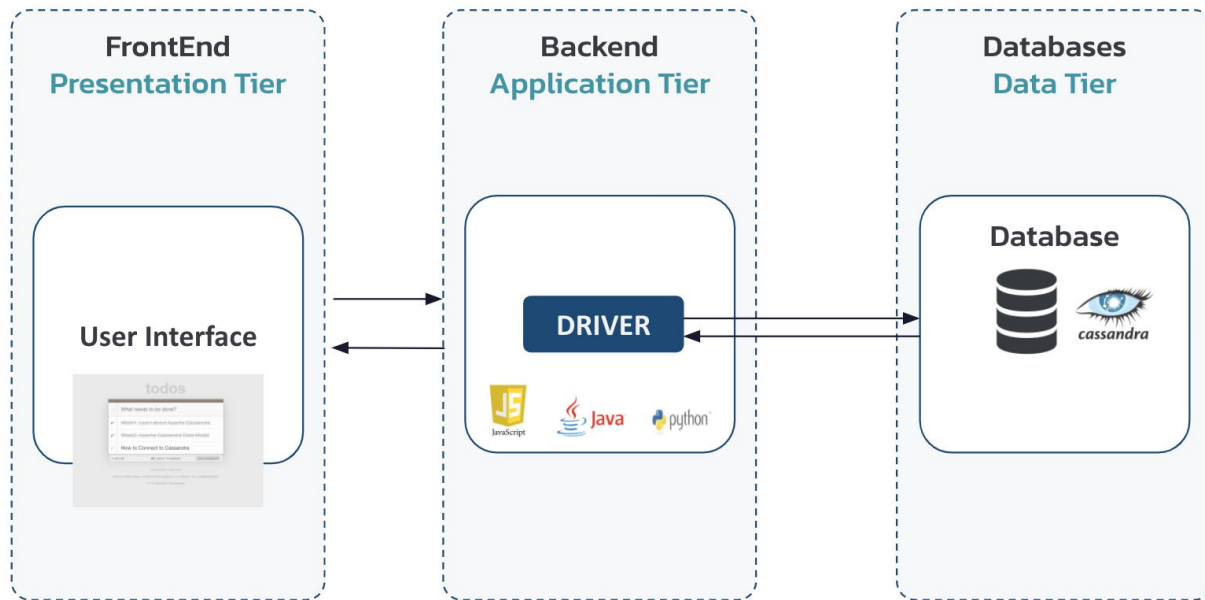
**Application development  
in Java and Python**

**06**

**Hands-on lab 3:  
Create an app with C\* driver**

# Cassandra Drivers

# Application Development with Apache Cassandra®



# Drivers



## Connectivity

- ★ Token & Datacenter Aware
- ★ Load Balancing Policies
- ★ Retry Policies
- ★ Reconnection Policies
- ★ Connection Pooling
- ★ Health Checks
- ★ Authentication | Authorization
- ★ SSL

## Query

- ★ CQL Support
- ★ Schema Management
- ★ Sync/Async/Reactive API
- ★ Query Builder
- ★ Compression
- ★ Paging

## Parsing Results

- ★ Lazy Load
- ★ Object Mapper
- ★ Spring Support
- ★ Paging

# Installing the Drivers

```
<dependency>
```

```
<groupId>com.datastax.oss</groupId>
```

```
<artifactId>java-driver-core</artifactId>
```

```
<version>4.13.1</version>
```

```
</dependency>
```

**Maven™**



```
npm install cassandra-driver
```

4.6.3 **npm**

```
{  
  "dependencies": {  
    "cassandra-driver": "^4.6.3"  
  }  
}
```



```
pip install cassandra-driver==3.25.0
```




**nuget** v3.15.0

```
Install-Package CassandraCSharpDriver -Version 3.15.0
```



# Contact Points (Cassandra)

- One contact point *would be enough*  
... unless that node is down
- ~3 nodes per DC for resilience
- From there, drivers discover whole cluster
- Local Datacenter





# Create "Session" Client (with contact points)

```
CqlSession cqlSession = CqlSession.builder()
    .addContactPoint(new InetSocketAddress("127.0.0.1", 9042))
    .withKeyspace("sensor_data")
    .withLocalDatacenter("dc1")
    .withAuthCredentials("U","P")
    .build();
```



```
const client = new cassandra.Client({
  contactPoints: ['127.0.0.1'],
  localDataCenter: 'dc1',
  keyspace: 'sensor_data',
  credentials: { username: 'U', password: 'P' }
});
await client.connect();
```



```
auth_provider = PlainTextAuthProvider(
    username='U', password='P')
cluster = Cluster(['127.0.0.1'],
    auth_provider=auth_provider, protocol_version=5)
session = cluster.connect('sensor_data')
```



```
Cluster cluster = Cluster.Builder()
    .AddContactPoint("127.0.0.1")
    .WithCredentials("U", "P")
    .Build();
session = cluster.Connect("sensor_data");
```



# Create "Session" Client (with Astra DB)

```
CqlSession cqlSession = CqlSession.builder()  
    .withCloudSecureConnectBundle(Paths.get("secure.zip"))  
    .withAuthCredentials("U","P")  
    .withKeyspace("sensor_data")  
    .build();
```



```
const client = new cassandra.Client({  
    cloud: { secureConnectBundle: 'secure.zip' },  
    credentials: { username: 'u', password: 'p' },  
    keyspace: 'sensor_data'  
});  
await client.connect();
```



```
auth_provider = PlainTextAuthProvider(  
    username='U', password='P')  
cluster = Cluster(  
    cloud ={'secure_connect_bundle': 'secure.zip'},  
    auth_provider=auth_provider, protocol_version=4)  
session= cluster.connect('sensor_data')
```



```
var cluster = Cluster.Builder()  
    .WithCloudSecureConnectionBundle("secure.zip")  
    .WithCredentials("u", "p")  
    .Build();  
var session = cluster.Connect("sensor_data");
```



# There Should Only Be One Session !

- Stateful object handling communications with each node
- Should be unique in the Application (*Singleton*)
- Should be closed at application shutdown (*shutdown hook*) in order to free opened TCP sockets (*stateful*)

<b>Java:</b>	<code>cqlSession.close();</code>
<b>Python:</b>	<code>session.shutdown();</code>
<b>Node:</b>	<code>client.shutdown();</code>
<b>CSharp:</b>	<code>IDisposable</code>

# Executing CQL Queries



```
session.execute(  
    "SELECT * FROM sensors_by_network WHERE network = %s;",  
    (network,) )
```



```
cqlSession.execute(  
    "SELECT * FROM sensors_by_network WHERE network = '" + network + "'" ) ;
```

# Prepared CQL Statements



```
q3_statement = session.prepare(  
    "SELECT * FROM sensors_by_network WHERE network = ?;"  
)  
rows = session.execute(q3_statement, (network,))
```



```
PreparedStatement q3Prepared = session.prepare(  
    "SELECT * FROM sensors_by_network WHERE network = ?");  
BoundStatement q3Bound = q3Prepared.bind(network);  
ResultSet rs = session.execute(q3Bound);
```

# Data APIs



Drivers

Open API



Use native drivers and the Cassandra Query Language to access your data in Cassandra



Go driverless with a high performance RPC (gRPC) API for every Cassandra database



Serve a GraphQL API from any Cassandra database, in schema first or cql first modes

{ REST : API }

Serve a RESTful API from any Cassandra database



Save and search schemaless JSON documents



More Performant



More Flexible

**01**



**Getting started with  
Apache Cassandra®**

**02**

**Hands-on lab 1:  
Create a database**

**03**

**Data definition and  
manipulation with CQL**

**04**

**Hands-on lab 2:  
Create, populate and query**

**05**

**Application development  
in Java and Python**

**06**

**Hands-on lab 3:  
Create an app with C\* driver**



# Lab 3

Create a Java or Python app  
to query a Cassandra database

Start here: [tinyurl.com/cassandra-dev](https://tinyurl.com/cassandra-dev)

- ✓ Explore and deploy a Java app (+ Spring Boot)
- ✓ Explore and deploy a Python app (+ Fast API)

*Follow the steps at the link*

# Stay in touch!

**Discord:** [dtsx.io/discord](https://dtsx.io/discord)

**Academy:** [academy.datastax.com](https://academy.datastax.com)

**Workshops:** [datastax.com/workshops](https://datastax.com/workshops)

**YouTube:** [@DataStaxDevelopers](https://www.youtube.com/DataStaxDevelopers)

# Get Cassandra help



**stackoverflow (\*) :**

[stackoverflow.com/questions/tagged/cassandra](https://stackoverflow.com/questions/tagged/cassandra)



**DBA Stack Exchange (\*) :**

[dba.stackexchange.com/questions/tagged/cassandra](https://dba.stackexchange.com/questions/tagged/cassandra)



**Discord:**

[dtsx.io/discord](https://dtsx.io/discord)

(\*) for best results, follow the "cassandra" tag

**Sponsored by**

**DataStax**

**Thank You**