



Sponsored by
DataStax

Hands-on Workshop

Cassandra Data Modeling & Application Development

You can't build a successful app without data modeling

Some words about data modeling

- * Not a relational database!
- * Denormalization
- * Keys and Partitions
- * Querying tables
- * Data Modeling methodology
- * **HANDS-ON:** explore a full data model

Build a successful app

- * Apps & Drivers (Python, Java)
- * **HANDS-ON:** run your first Cassandra-backed API

Resources, homework, badges

Not a relational database

Definition

Data modeling =

the act of laying out the structure of your database to fit your application's purposes

Begins with an abstract idea of your application's needs

Ends with the commands that create the physical tables and the read/write operations that they support

Data Modeling: Relational vs. Cassandra



Data -> Model -> Application	Application -> Model -> Data
Joins, indexes	Denormalization
Referential integrity	No (would not scale well anyway)
Primary key = uniqueness	Primary key = uniqueness + partitioning + access
Entity-driven	Query-driven

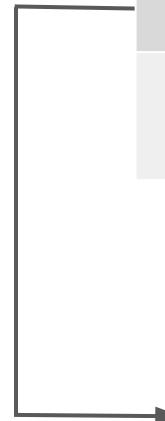
Normalization vs Denormalization

Normalization

“Database normalization is the process of structuring a relational database in accordance with a series of so-called normal forms in order to reduce data redundancy and improve data integrity. It was first proposed by Edgar F. Codd as part of his relational model.”

PROS: Simple write, Data Integrity

CONS: Slow read, Complex Queries



The diagram illustrates a one-to-many relationship between the Employees and Departments tables. A line with an arrow originates from the deptId column in the Employees table and points to the departmentId column in the Departments table.

Employees			
userId	deptId	firstName	lastName
1	1	Edgar	Codd
2	1	Raymond	Boyce

Departments	
departmentId	department
1	Engineering
2	Math

Denormalization

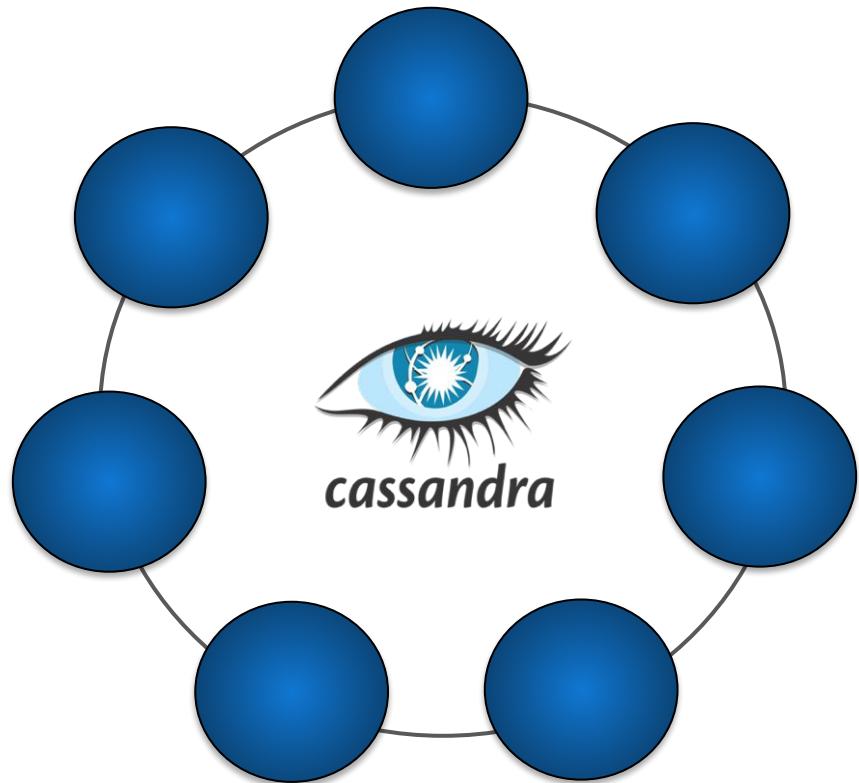
“Denormalization is a strategy used on a database to increase performance. In computing, denormalization is the process of trying to improve the read performance of a database, at the expense of losing some write performance, by adding redundant copies of data”

PROS: Quick Read, Simple Queries

CONS: Multiple Writes, Manual Integrity

Employees			
userId	firstName	lastName	department
1	Edgar	Codd	Engineering
2	Raymond	Boyce	Engineering
3	Sage	Lahja	Math
4	Juniper	Jones	Botany

departmentId	department
1	Engineering
2	Math



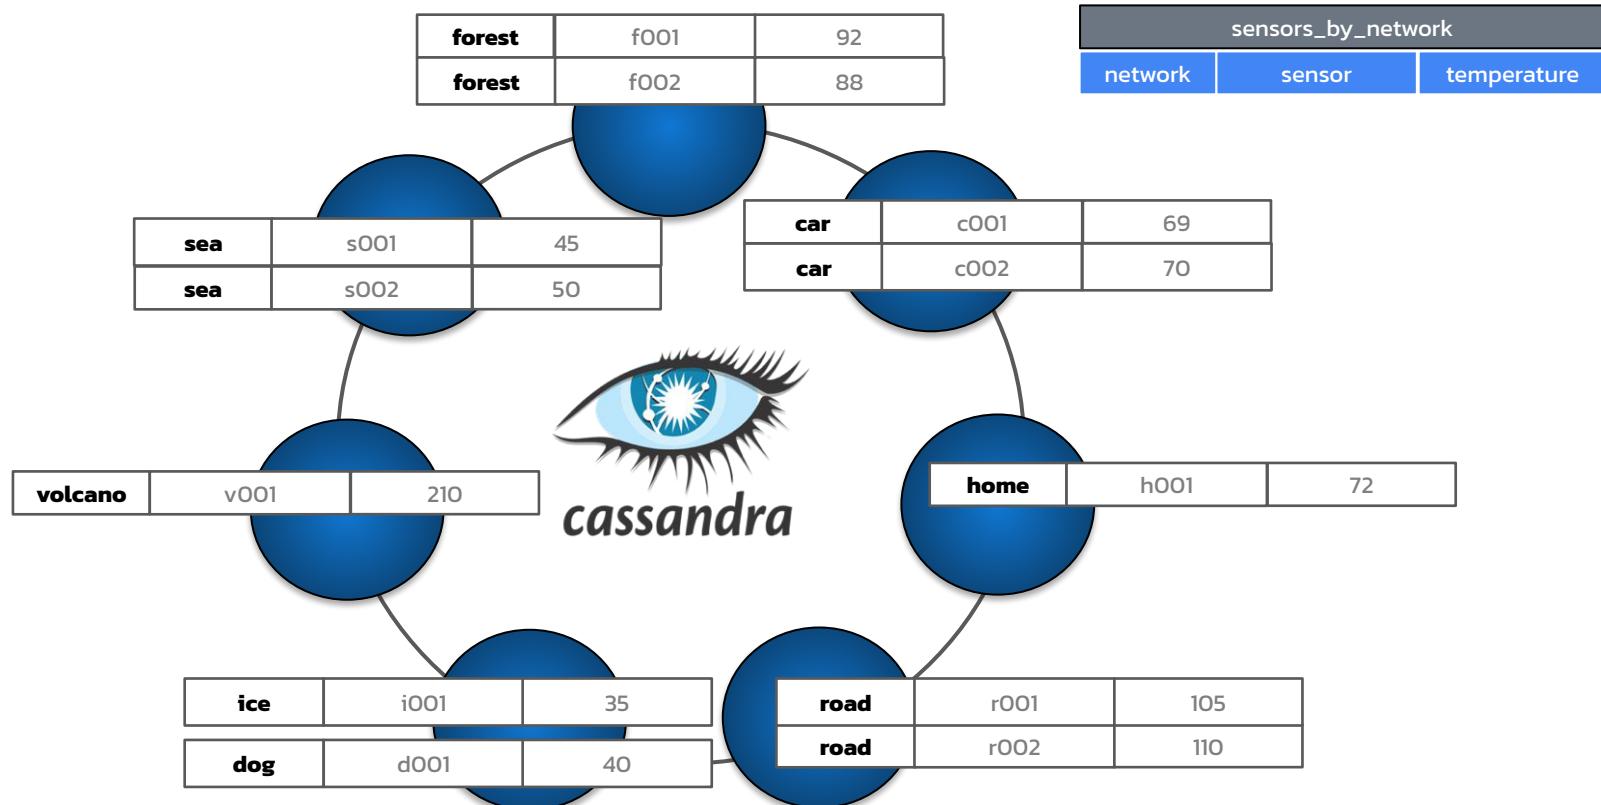
sensors_by_network		
network	sensor	temperature
forest	f001	92
	f002	88
volcano	v001	210
	s001	45
sea	s002	50
	h001	72
car	c001	69
car	c002	70
dog	d001	40
road	r001	105
road	r002	110
ice	i001	35

Partition Key

Primary Key

Keys and Partitions

Data is Organized into Distributed Tables



Partition Key

An identifier for a partition.
Consists of at least one column,
may have more if needed.

PARTITIONS TABLE.

Defines data distribution strategy

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
);
```

Partition key

Examples:

```
PRIMARY KEY ((sensor), timestamp);
```

```
PRIMARY KEY ((sensor), date, timestamp);
```

```
PRIMARY KEY ((sensor, timestamp));
```

Clustering Columns

Used to **ensure uniqueness** and **sorting order**. Optional.

Defines physical ordering on disk

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Clustering columns

PRIMARY KEY ((**sensor**));

Not Unique 

PRIMARY KEY ((**sensor**), timestamp);

Not filter on dates 

PRIMARY KEY ((**sensor**), value, timestamp);

Not sorted 

PRIMARY KEY ((**sensor**), date, timestamp);



Primary Key

An identifier for a row. Consists of at least one Partition Key and zero or more Clustering Columns.

UNIQUELY IDENTIFIES A ROW

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Primary key

Examples:

```
PRIMARY KEY ((sensor), timestamp);
```

```
PRIMARY KEY ((sensor, date), timestamp);
```

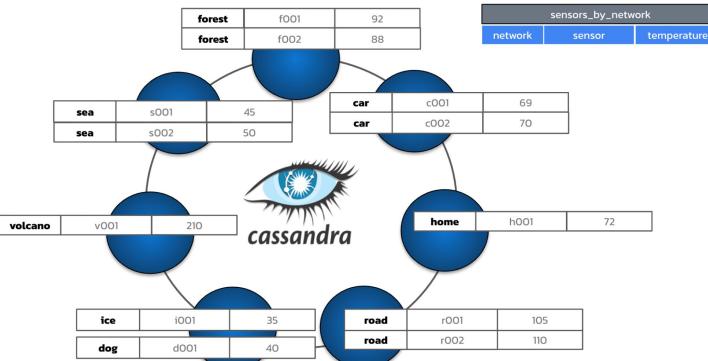
Primary Key = Ensure Uniqueness = PK + CC

Partition Keys:

Defines data distribution over the cluster

Clustering Columns

Defines how data is physically stored (written on disk)



```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Corollary: Schema Immutability

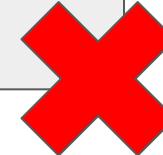
Important:

Once created, the data model cannot be changed! You will need new tables and migration. Stay lazy, design it right in advance!

```
ALTER TABLE temperatures_by_sensor  
ADD season TEXT;
```



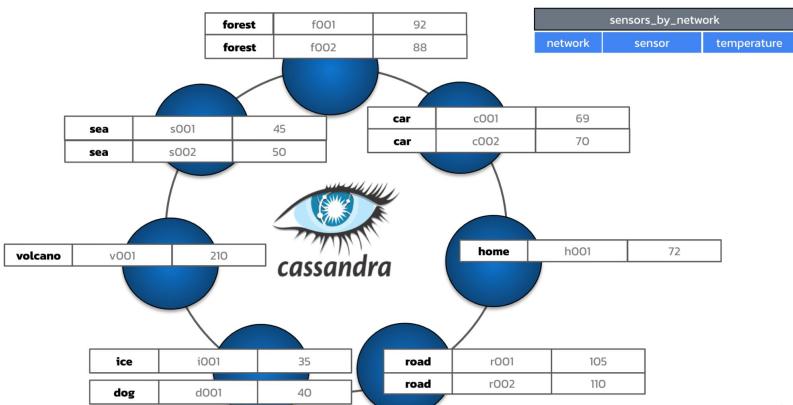
```
ALTER TABLE temperatures_by_sensor  
DROP PRIMARY KEY  
ADD PRIMARY KEY (sensor, timestamp)
```



Keys and Partitions (Recap)

Primary Key, Partition Key, Clustering Columns, Oh My!

```
CREATE TABLE sensor_data.temperatures_by_sensor (  
    sensor      TEXT,  
    date        DATE,  
    timestamp   TIMESTAMP,  
    value       FLOAT,  
    PRIMARY KEY ( ( sensor , date ) , timestamp )  
) WITH CLUSTERING ORDER BY (timestamp DESC);
```



Partition Key
DEFINES DATA DISTRIB.

Clustering column(s)
ENSURE UNIQUENESS

Primary key
UNIQUELY IDENTIFIES A ROW

Querying Tables

"Rule Zero" of Querying

"I want to read ALL rows of my table"

Sorry, not today.

If you really need to,

⇒ OLAP workload: use Apache Spark & co.

but in a typical OLTP app,
you actually **don't need** to run such a query!

**Reading all rows ("full-table scan") is a performance capital sin,
since it involves a lot of nodes in the cluster!**

Table Structure ⇒ Valid Queries

```
PRIMARY KEY ((sensor, date), timestamp);
```

```
SELECT * FROM temperatures_by_sensor ...  
  
WHERE sensor = ?;  
  
WHERE sensor > ?;  
  
WHERE sensor = ? AND date > ?;  
  
WHERE sensor = ? AND date = ?;  
  
WHERE sensor = ? AND date = ? AND timestamp > ?;
```

Rules of a Good Partition 1/4

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Rules of a Good Partition 2/4

- ❖ **Store together what you retrieve together**
- ❖ Avoid big partitions
- ❖ Avoid hot partitions

Q: Show temperature evolution over time for **sensor X** On **Sept 20th 2022**

PRIMARY KEY ((**sensor**, timestamp));



PRIMARY KEY (**sensor**, timestamp);



Rules of a Good Partition 3/4

- ❖ Store together what you retrieve together
- ❖ **Avoid big partitions**
- ❖ Avoid hot partitions

BUCKETING

PRIMARY KEY ((*sensor*), timestamp);



PRIMARY KEY ((*sensor*, month), timestamp);



- Up to 2 billion cells per partition
- Up to ~100k values in a partition
- Up to ~100MB in a Partition

Rules of a Good Partition 4/4

- ❖ Store together what you retrieve together
- ❖ Avoid big partitions
- ❖ **Avoid hot partitions**

```
PRIMARY KEY ((date), sensor, timestamp);
```



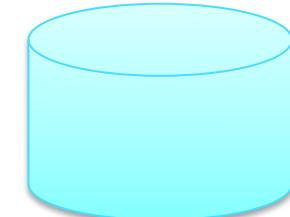
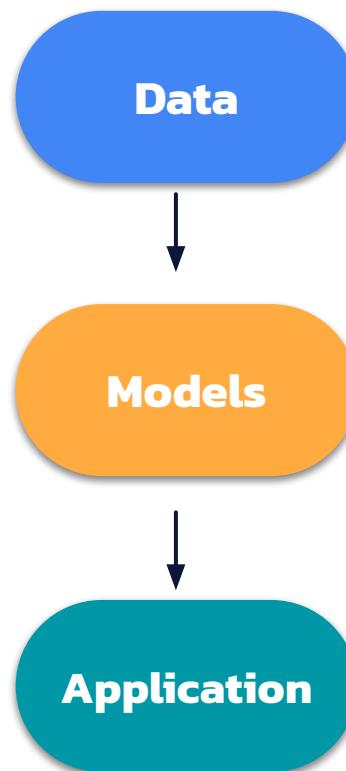
```
PRIMARY KEY ((date, sensor), timestamp);
```



Data Modeling Methodology

Relational Data Modeling

1. Analyze raw data
2. Identify entities, their properties and relations
3. Design tables, using **normalization** and foreign keys.
4. Use JOIN when doing queries to join normalized data from multiple tables



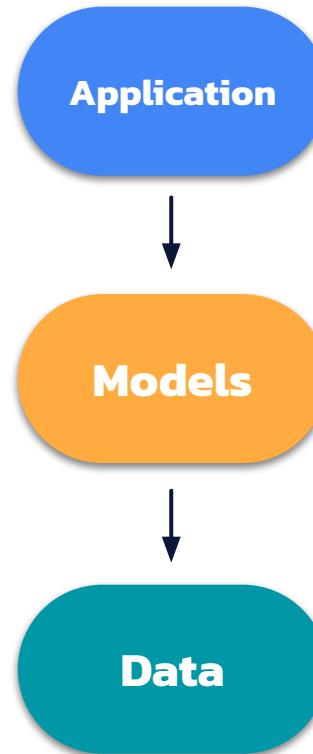
Employees		
userId	firstName	lastName
1	Edgar	Codd
2	Raymond	Boyce

Departments	
departmentId	department
1	Engineering
2	Math

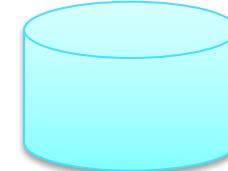


NoSQL Data Modeling

1. Analyze user behaviour
(customer first!)
2. Identify workflows, their dependencies and needs
3. Define Queries to fulfill these workflows
4. Knowing the queries, design tables, using **denormalization**.
5. Use BATCH when inserting or updating denormalized data of multiple tables



Employees			
userId	firstName	lastName	department
1	Edgar	Codd	Engineering
2	Raymond	Boyce	Math
3	Sage	Lahja	Math
4	Juniper	Jones	Botany



What is Data Modeling?

- Collection and analysis of **data requirements**
- Identification of participating **entities** and relationships
- Identification of data **access patterns**
- A particular way of **organizing** and structuring data
- Design and specification of a **database schema**
- Schema **optimization** and data **indexing** techniques



Data Quality: completeness consistency accuracy
Data Access: queryability efficiency scalability

Cassandra Data Modeling Principles

Modeling principle 1: “Know your data”

- Key and cardinality constraints are fundamental to schema design

Modeling principle 2: “Know your queries”

- Queries drive schema design

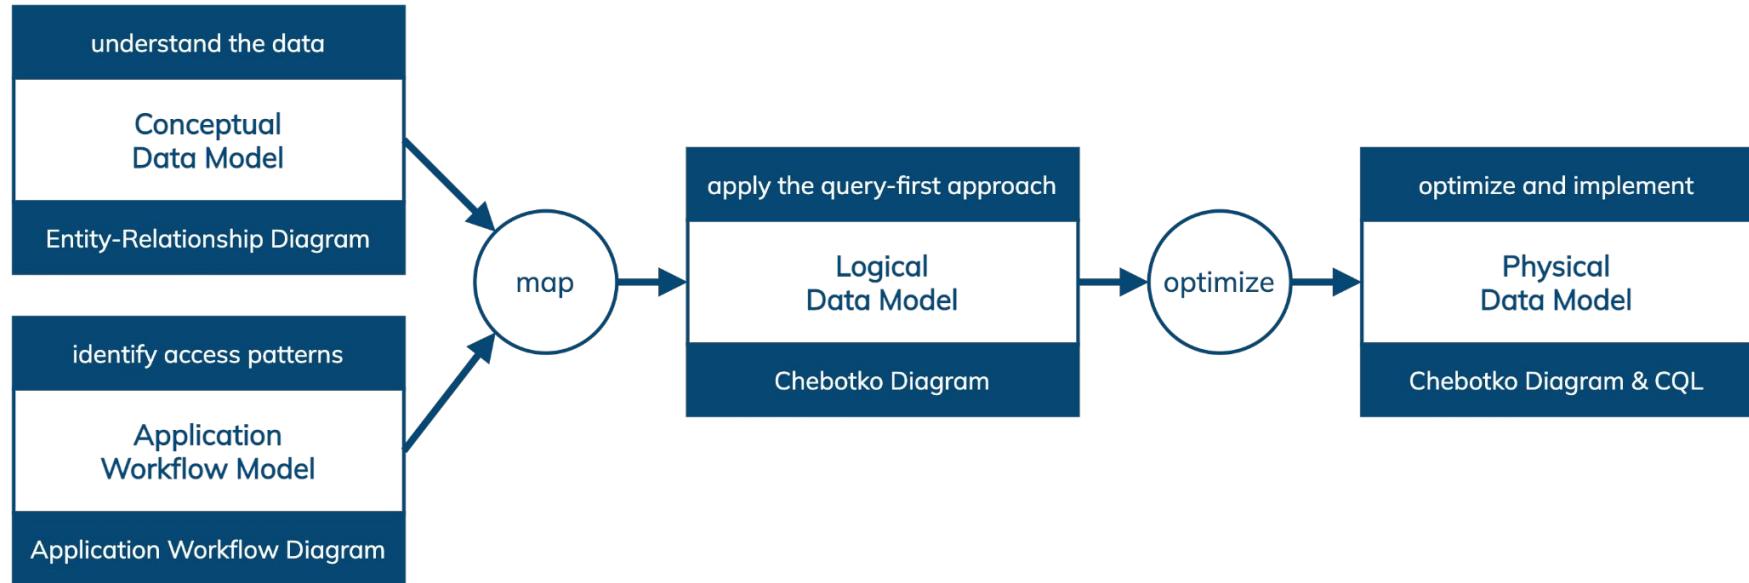
Modeling principle 3: “Nest data”

- Data nesting is the main data modeling technique

Modeling principle 4: “Duplicate data”

- Better to duplicate than to join

Modeling Workflow



Data Modeling Methodology: Step I

understand the data

Conceptual
Data Model

Entity-Relationship Diagram

Analyze the Domain

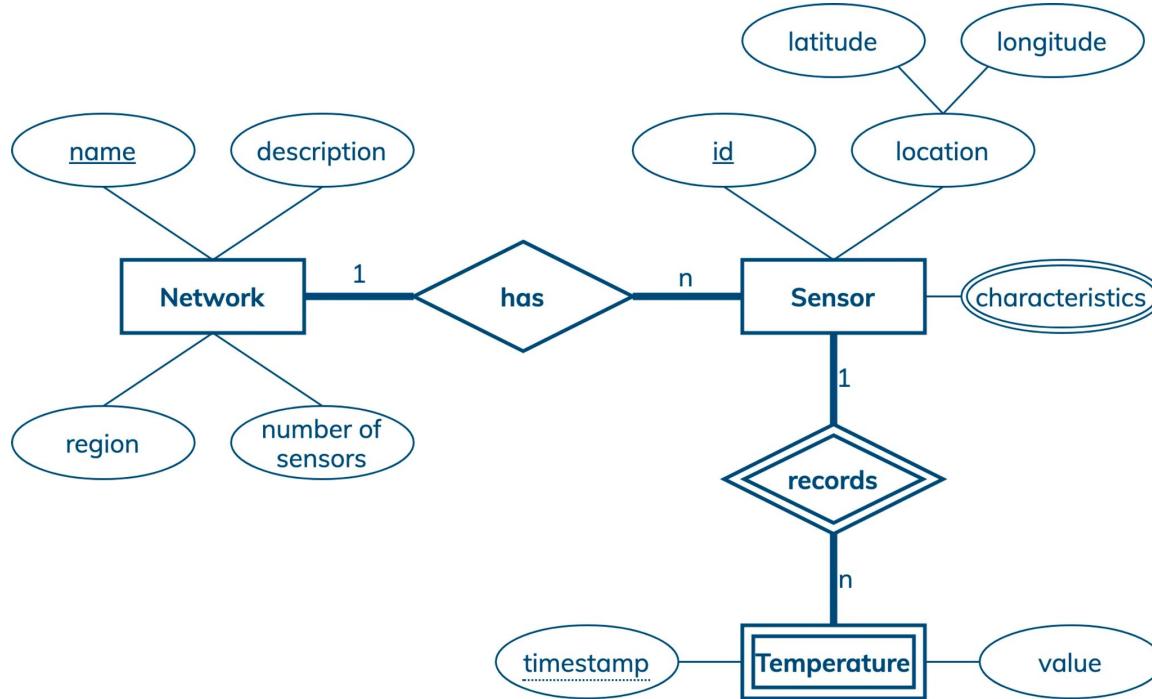
identify access patterns

Application
Workflow Model

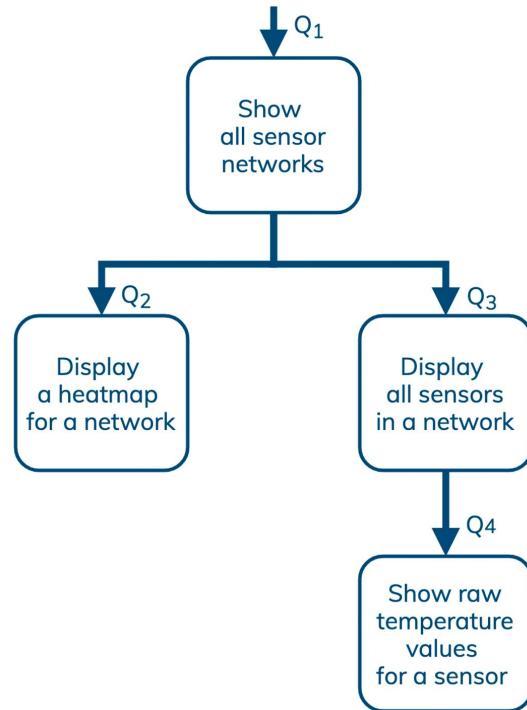
Application Workflow Diagram

Analyze Customer Workflows

Entity-Relationship Diagram



Application Workflow Diagram



Data access patterns

Q₁: Find information about all networks; order by name (asc)

Q₂: Find hourly average temperatures for every sensor in a specified network for a given date range; order by date (desc) and hour (desc)

Q₃: Find information about all sensors in a specified network

Q₄: Find raw measurements for a particular sensor on a specified date; order by timestamp (desc)

Data Modeling Methodology: Step II



Design queries and build
tables based on the queries

Mapping Rules

Mapping rule 1: “Entities and relationships”

- Entity and relationship types map to tables

Based on
a conceptual
data model

Mapping rule 2: “Equality search attributes”

- Equality search attributes map to the beginning columns of a primary key

Based on
a query

Mapping rule 3: “Inequality search attributes”

- Inequality search attributes map to clustering columns

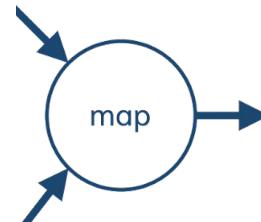
Mapping rule 4: “Ordering attributes”

- Ordering attributes map to clustering columns

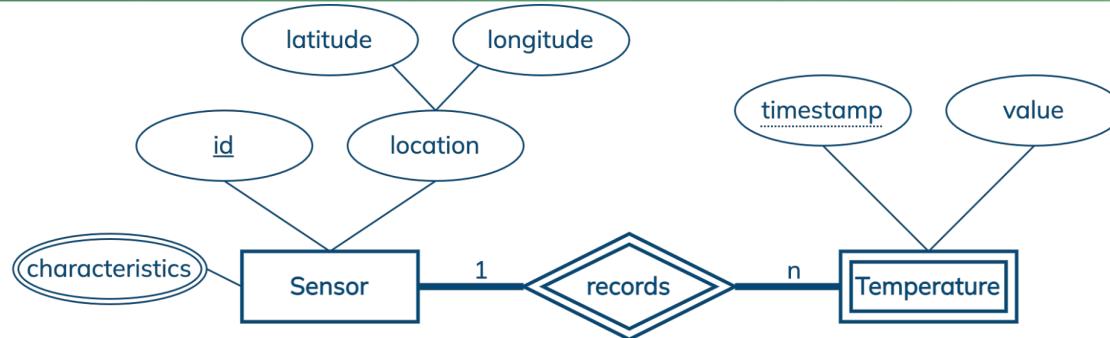
Based on
a conceptual
data model

Mapping rule 5: “Key attributes”

- Key attributes map to primary key columns



Mapping Rules: Example



Q5: Find raw measurements for a given location and a date/time range; order by timestamp (desc)

The diagram shows five intermediate results (MR1 to MR5) for a map-reduce process. Each result is represented as a table with columns for latitude, longitude, timestamp, and value. The tables show the progression of data and key-value pairs (K) through the process:

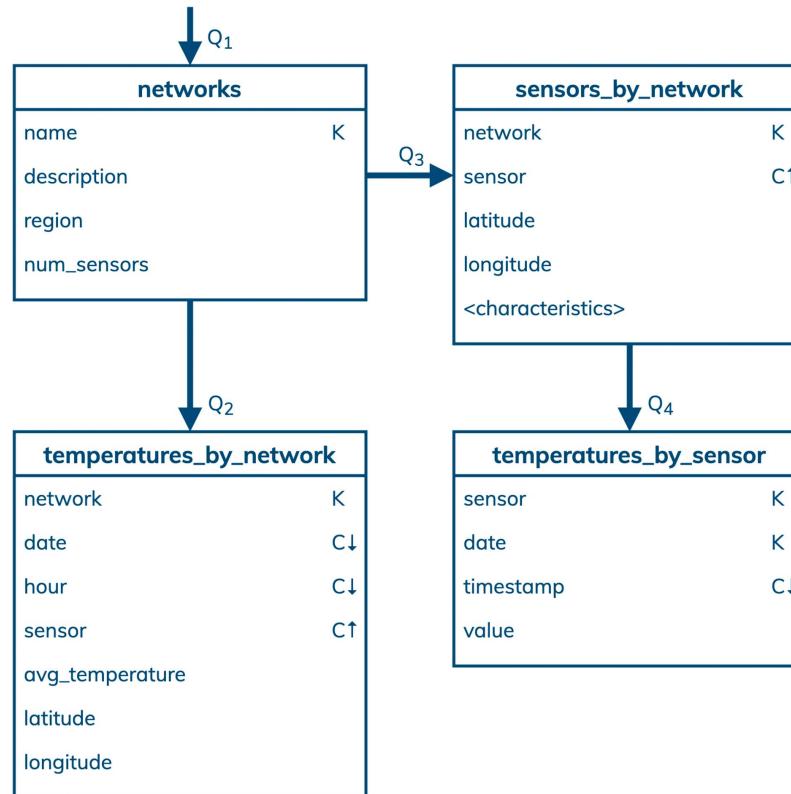
- MR1:** Shows raw data with columns: latitude, longitude, timestamp, value.
- MR2:** Shows data grouped by location (latitude, longitude) with values K. The timestamp column is present.
- MR3:** Shows data grouped by location and timestamp (latitude, longitude, timestamp) with values K. The timestamp column is present.
- MR4:** Shows data grouped by location, timestamp, and value. The timestamp column has a C↑ indicator, and the value column has a C↓ indicator.
- MR5:** Shows the final output with all columns: latitude, longitude, timestamp, and value. The timestamp column has a C↓ indicator, and the value column has a C↑ indicator.

temps_by_sensor			
latitude	longitude	timestamp	value
MR1	MR2	MR3	MR4
latitude	longitude	K	K
longitude		K	
timestamp		C↑	C↓
value	value	value	value

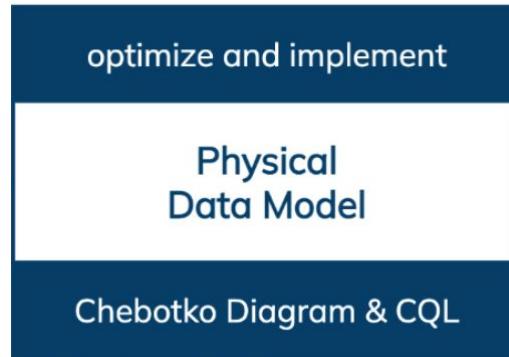
Chebotko Diagram

Data access patterns

- Q1: Find information about all networks; order by name (asc)
- Q2: Find hourly average temperatures for every sensor in a specified network for a given date range; order by date (desc) and hour (desc)
- Q3: Find information about all sensors in a specified network
- Q4: Find raw measurements for a particular sensor on a specified date; order by timestamp (desc)



Data Modeling Methodology: Step III



Optimize and Create

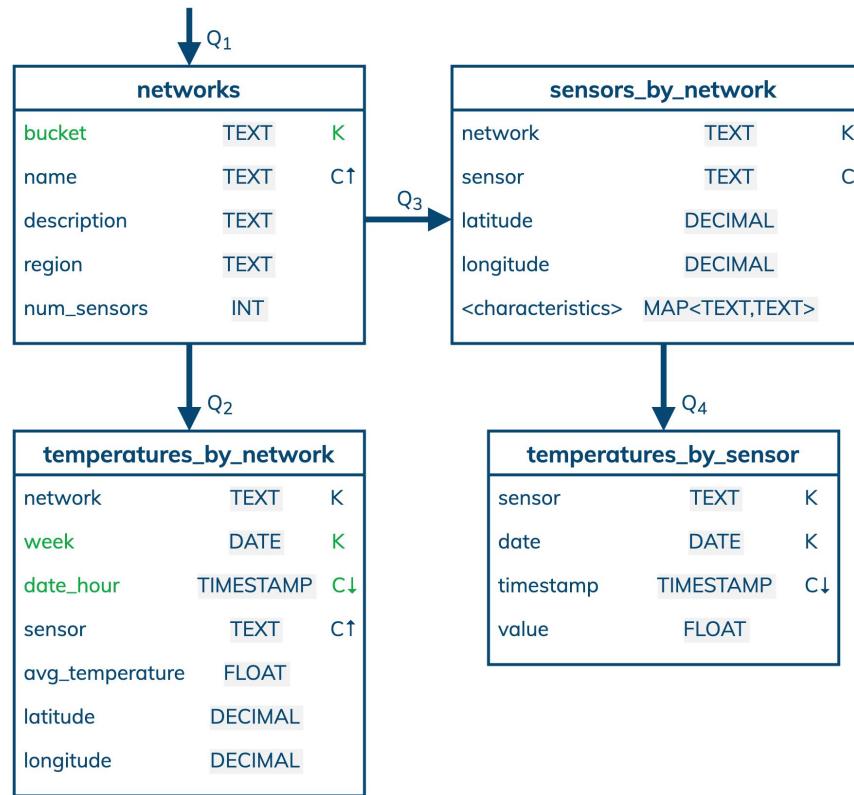
Optimization Techniques

- Partition size limits and splitting large partitions
- Data duplication and batches
- Indexes and materialized views*
- Concurrent data access and lightweight transactions*
- Dealing with tombstones*



* Explained in details at our free DS220 course at academy.datastax.com

Physical Data Model: Chebotko Diagram



Cassandra Query Language (CQL)

Q1

```
CREATE TABLE networks (
    bucket TEXT,
    name TEXT,
    description TEXT,
    region TEXT,
    num_sensors INT,
    PRIMARY KEY ((bucket),name)
);
```

Q3

```
CREATE TABLE sensors_by_network (
    network TEXT,
    sensor TEXT,
    latitude DECIMAL,
    longitude DECIMAL,
    characteristics MAP<TEXT,TEXT>,
    PRIMARY KEY ((network),sensor)
);
```

Q2

```
CREATE TABLE temperatures_by_network (
    network TEXT,
    week DATE,
    date_hour TIMESTAMP,
    sensor TEXT,
    avg_temperature FLOAT,
    latitude DECIMAL,
    longitude DECIMAL,
    PRIMARY KEY ((network,week),date_hour,sensor)
) WITH CLUSTERING ORDER BY (date_hour DESC, sensor ASC);
```

Q4

```
CREATE TABLE temperatures_by_sensor (
    sensor TEXT,
    date DATE,
    timestamp TIMESTAMP,
    value FLOAT,
    PRIMARY KEY ((sensor,date),timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

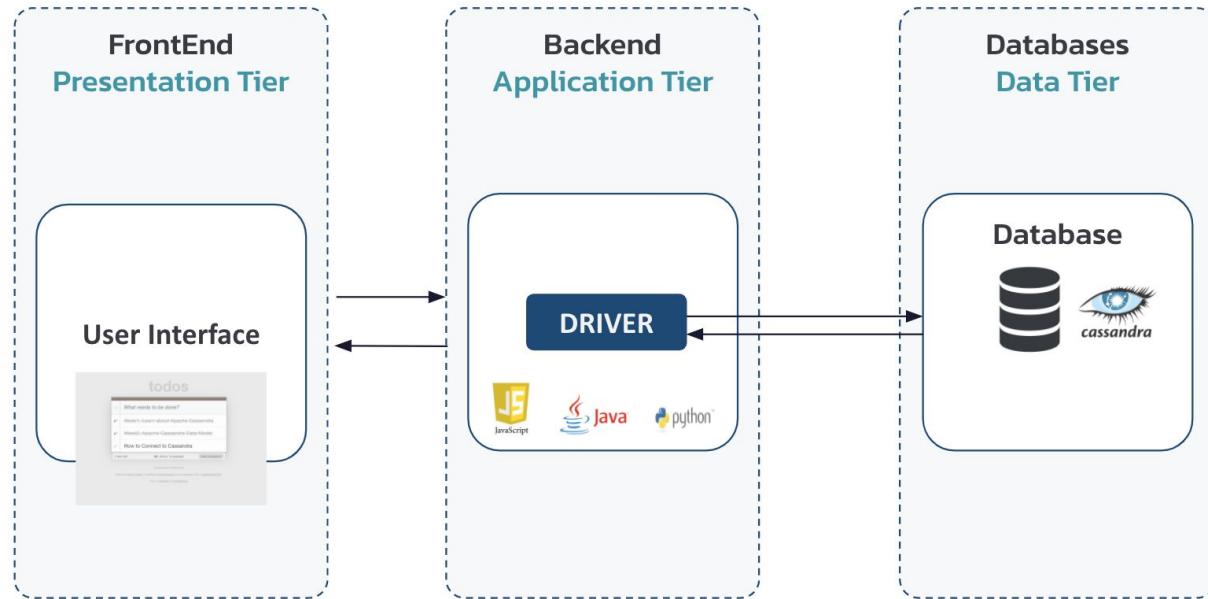
Lab 1

Explore a data model

Start here: dtsx.io/cday-ws2

Apps & Drivers

Application Development with Apache Cassandra™



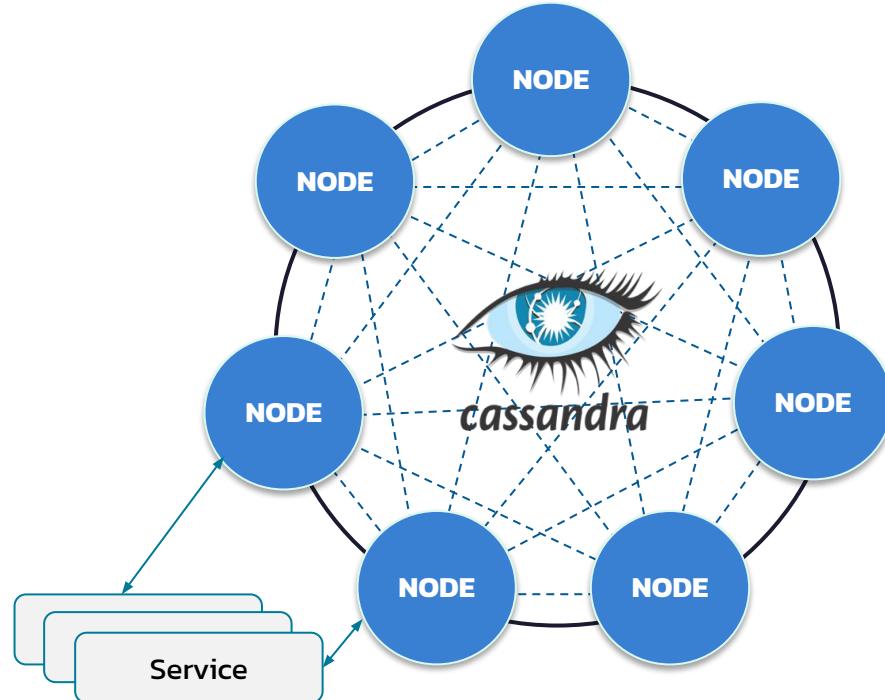
Cassandra ❤️ Microservices

Loose Coupling: Data resiliency

- Data Replicated on multiple Nodes
- Load Balancing at driver side
- Health Check at driver side
- Hinted HandOff

Shared nothing: Data isolation

- Per Keyspace (with replications)
- Per Tables (1 query = 1 table)
- Per profile (RBAC)



Drivers



Connectivity

- ★ Token & Datacenter Aware
- ★ Load Balancing Policies
- ★ Retry Policies
- ★ Reconnection Policies
- ★ Connection Pooling
- ★ Health Checks
- ★ Authentication | Authorization
- ★ SSL

Query

- ★ CQL Support
- ★ Schema Management
- ★ Sync/Async/Reactive API
- ★ Query Builder
- ★ Compression
- ★ Paging

Parsing Results

- ★ Lazy Load
- ★ Object Mapper
- ★ Spring Support
- ★ Paging

Installing the Drivers

```
<dependency>  
  
    <groupId>com.datastax.oss</groupId>  
  
    <artifactId>java-driver-core</artifactId>  
  
    <version>4.13.1</version>  
  
</dependency>
```



```
pip install cassandra-driver==3.25.0
```



```
npm install cassandra-driver
```

```
{  
  "dependencies": {  
    "cassandra-driver": "^4.6.3"  
  }  
}
```

4.6.3



JavaScript

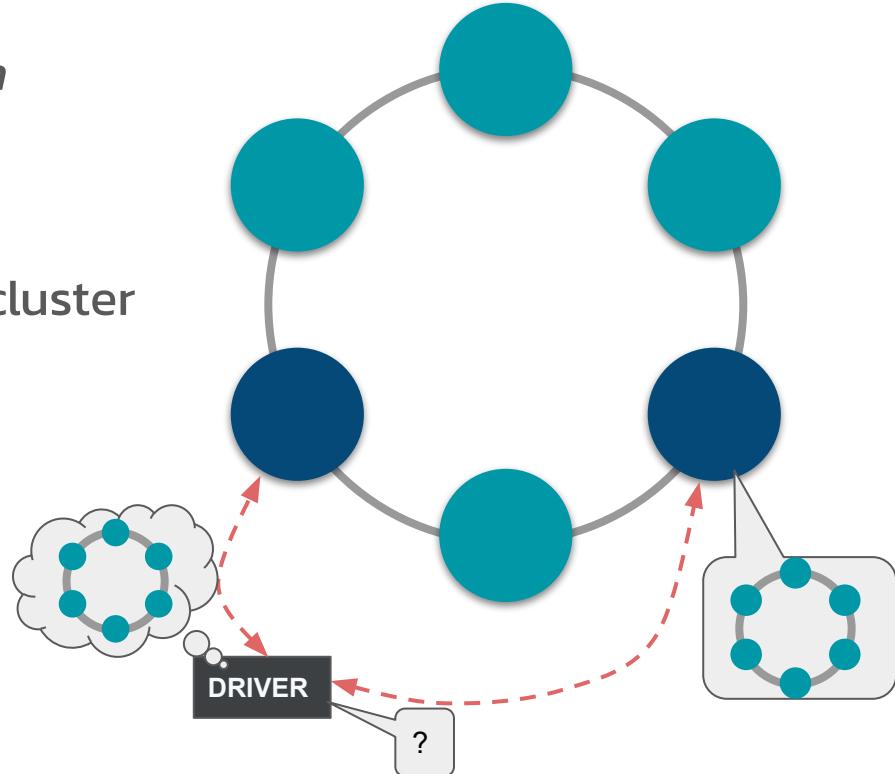
nuget v3.15.0

```
Install-Package CassandraCSharpDriver -Version 3.15.0
```



Contact Points (Cassandra)

- One contact point *would be enough*
... unless that node is down
- ~3 nodes per DC for resilience
- From there, drivers discover whole cluster
- Local Datacenter



Create "Session" Client (with contact points)

```
CqlSession cqlSession = CqlSession.builder()  
    .addContactPoint(new InetSocketAddress("127.0.0.1", 9042))  
    .withKeyspace("sensor_data")  
    .withLocalDatacenter("dc1")  
    .withAuthCredentials("U", "P")  
    .build();
```



```
const client = new cassandra.Client({  
  contactPoints: ['127.0.0.1'],  
  localDataCenter: 'dc1',  
  keyspace: 'sensor_data',  
  credentials: { username: 'U', password: 'P' }  
});  
await client.connect();
```



```
auth_provider = PlainTextAuthProvider(  
    username='U', password='P')  
  
cluster = Cluster(['127.0.0.1'],  
    auth_provider=auth_provider, protocol_version=5)  
  
session = cluster.connect('sensor_data')
```



```
Cluster cluster = Cluster.Builder()  
    .AddContactPoint("127.0.0.1")  
    .WithCredentials("U", "P")  
    .Build();  
  
session = cluster.Connect("sensor_data");
```



Create "Session" Client (with Astra DB)

```
CqlSession cqlSession = CqlSession.builder()  
    .withCloudSecureConnectBundle(Paths.get("secure.zip"))  
    .withAuthCredentials("U","P")  
    .withKeyspace("sensor_data")  
    .build();
```



```
auth_provider = PlainTextAuthProvider(  
    username='U', password='P')  
  
cluster = Cluster(  
    cloud ={  
        'secure_connect_bundle': 'secure.zip'},  
    auth_provider=auth_provider, protocol_version=4)  
  
session= cluster.connect('sensor_data')
```



```
const client = new cassandra.Client({  
    cloud: { secureConnectBundle: 'secure.zip' },  
    credentials: { username: 'u', password: 'p' },  
    keyspace: 'sensor_data'  
});  
await client.connect();
```



```
var cluster = Cluster.Builder()  
    .WithCloudSecureConnectionBundle("secure.zip")  
    .WithCredentials("u", "p")  
    .Build();  
  
var session = cluster.Connect("sensor_data");
```



There Should Only Be One Session !

- Stateful object handling communications with each node
- Should be unique in the Application (*Singleton*)
- Should **be closed** at application shutdown (*shutdown hook*) in order to free opened TCP sockets (*stateful*)

```
Java:      cqlSession.close();  
Python:    session.shutdown();  
Node:      client.shutdown();  
CSharp:    IDisposable
```

Executing CQL Queries

 python™

```
session.execute(  
    "SELECT * FROM sensors_by_network WHERE network = %s;",  
    (network,)  
)
```

 Java

```
cqlSession.execute(  
    "SELECT * FROM sensors_by_network WHERE network = '" + network + "'")
```

Prepare Your Statements

```
q3_statement = session.prepare(  
    "SELECT * FROM sensors_by_network WHERE network = ?;"  
)  
rows = session.execute(q3_statement, (network,) )
```



Prepare Your Statements

```
PreparedStatement q3Prepared = session.prepare(  
    "SELECT * FROM sensors_by_network WHERE network = ?");  
BoundStatement q3Bound = q3Prepared.bind(network);  
ResultSet rs = session.execute(q3Bound);
```



Advantages of Prepared Statements

- Parse once, run many times.
- Saves network trips for result set metadata.
- Client-side type validation.
- Statements binding on partition keys compute their own cluster routing.

Pick Your Language (For Today)



today: +Spring Boot



today: +FastAPI

Lab 2

Create DB and run API

Start here: dtsx.io/cday-ws2

Conclusion & Next Steps

Homework (see AppDev Github Repo!)



A couple of "theory" questions, plus ...

Coding exercise: *Enrich the API with a new GET endpoint for Q1 ("get all networks")*

Get help about Cassandra



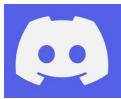
stack**overflow** (*):

stackoverflow.com/questions/tagged/cassandra



DBA Stack Exchange (*):

dba.stackexchange.com/questions/tagged/cassandra



Discord:

dtsx.io/discord

(*) for best results, follow the "cassandra" tag

This is just the beginning of your journey



Roads? Where we're going we don't need roads

Enjoy the rest of today
Do the homework ⇒ get a badge
We have more workshops & hands-on labs!
Get Cassandra Certified (Academy & Exam voucher)

Sponsored by

DataStax

Thank You