



**Cassandra Day**  
2022

Hands-on Workshop

# **Cassandra Data Modeling & Application Development**

**Sponsored by DataStax**

# You can't build a successful app without data modeling

## Some words about data modeling

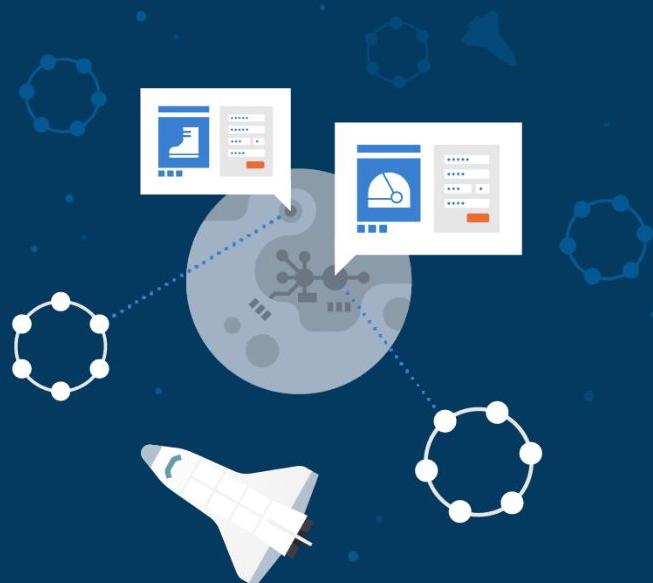
- \* Not a relational database!
- \* Denormalization
- \* Keys and Partitions
- \* Querying tables
- \* Data Modeling methodology
- \* **HANDS-ON:** explore a full data model

## Build a successful app

- \* Apps & Drivers (Python, Java)
- \* **HANDS-ON:** run your first Cassandra-backed API

## Resources, homework, badges

Programme



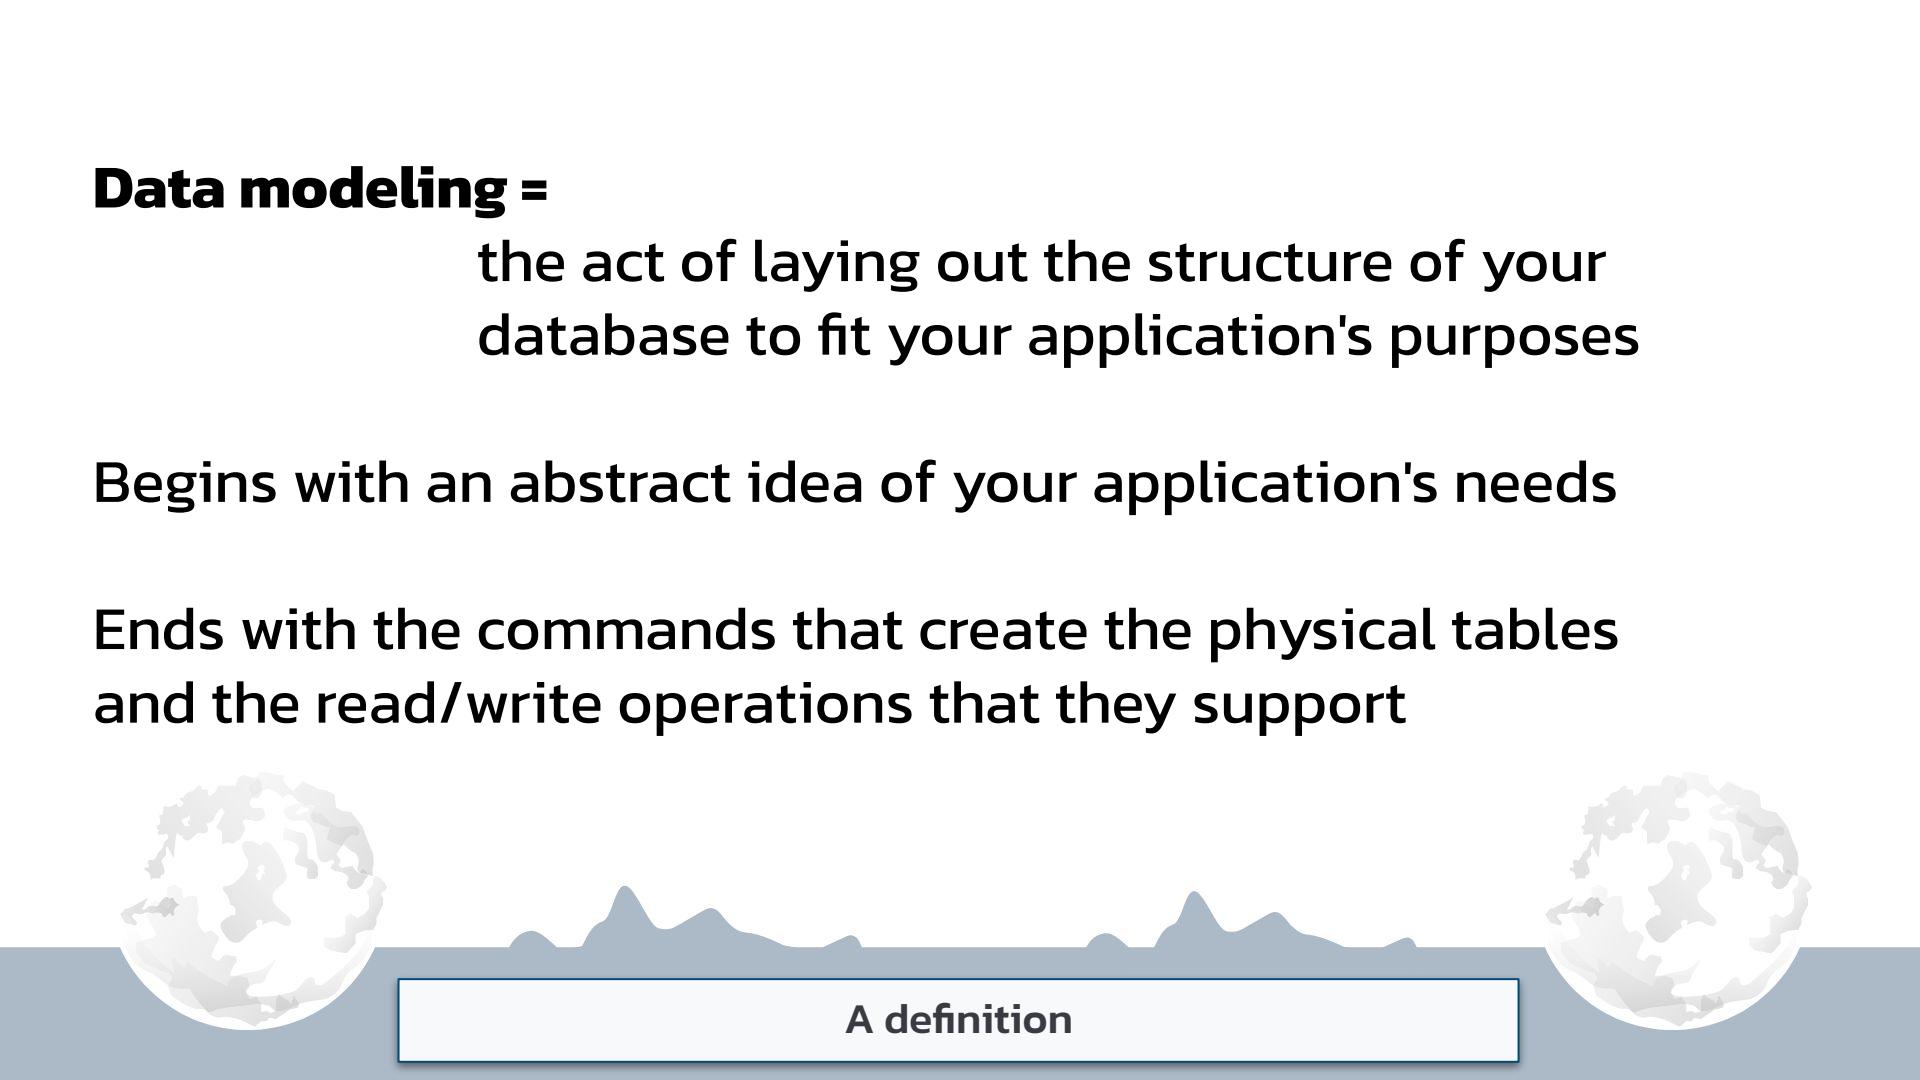
**Not a relational database!**

## **Data modeling =**

the act of laying out the structure of your database to fit your application's purposes

Begins with an abstract idea of your application's needs

Ends with the commands that create the physical tables and the read/write operations that they support



A definition

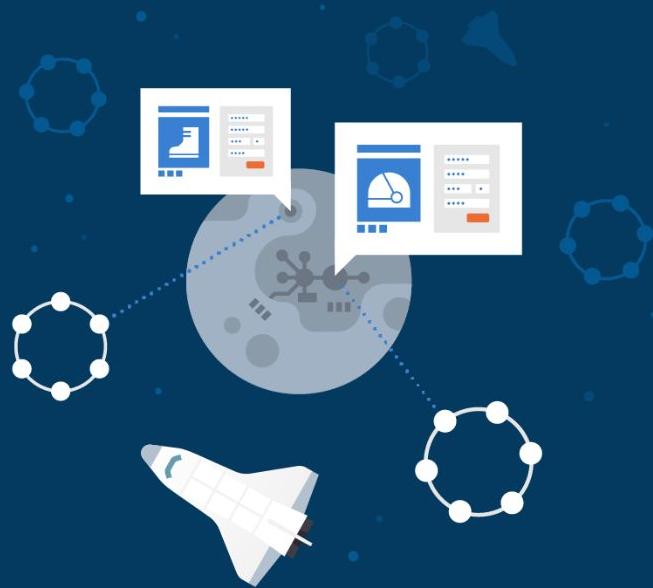


| Data -> Model -> Application | Application -> Model -> Data                            |
|------------------------------|---|
| Joins, indexes               | <b>Denormalization</b>                                  |
| Referential integrity        | <b>No</b> (would not scale well anyway)                 |
| Primary key = uniqueness     | <b>Primary key = uniqueness + partitioning + access</b> |
| Entity-driven                | <b>Query-driven</b>                                     |



**Data modeling: relational VS Cassandra**





## Normalization vs Denormalization

"Database normalization is the process of structuring a relational database in accordance with a series of so-called normal forms in order to reduce data redundancy and improve data integrity. It was first proposed by Edgar F. Codd as part of his relational model."

**PROS:** Simple write, Data Integrity  
**CONS:** Slow read, Complex Queries



## Normalization

Employees

| userId | deptId | firstName | lastName |
|--------|--------|-----------|----------|
| 1      | 1      | Edgar     | Codd     |
| 2      | 1      | Raymond   | Boyce    |

Departments

| departmentId | department  |
|--------------|-------------|
| 1            | Engineering |
| 2            | Math        |



**"Denormalization** is a strategy used on a database to increase performance. In computing, denormalization is the process of trying to improve the read performance of a database, at the expense of losing some write performance, by adding redundant copies of data"

**PROS:** Quick Read, Simple Queries

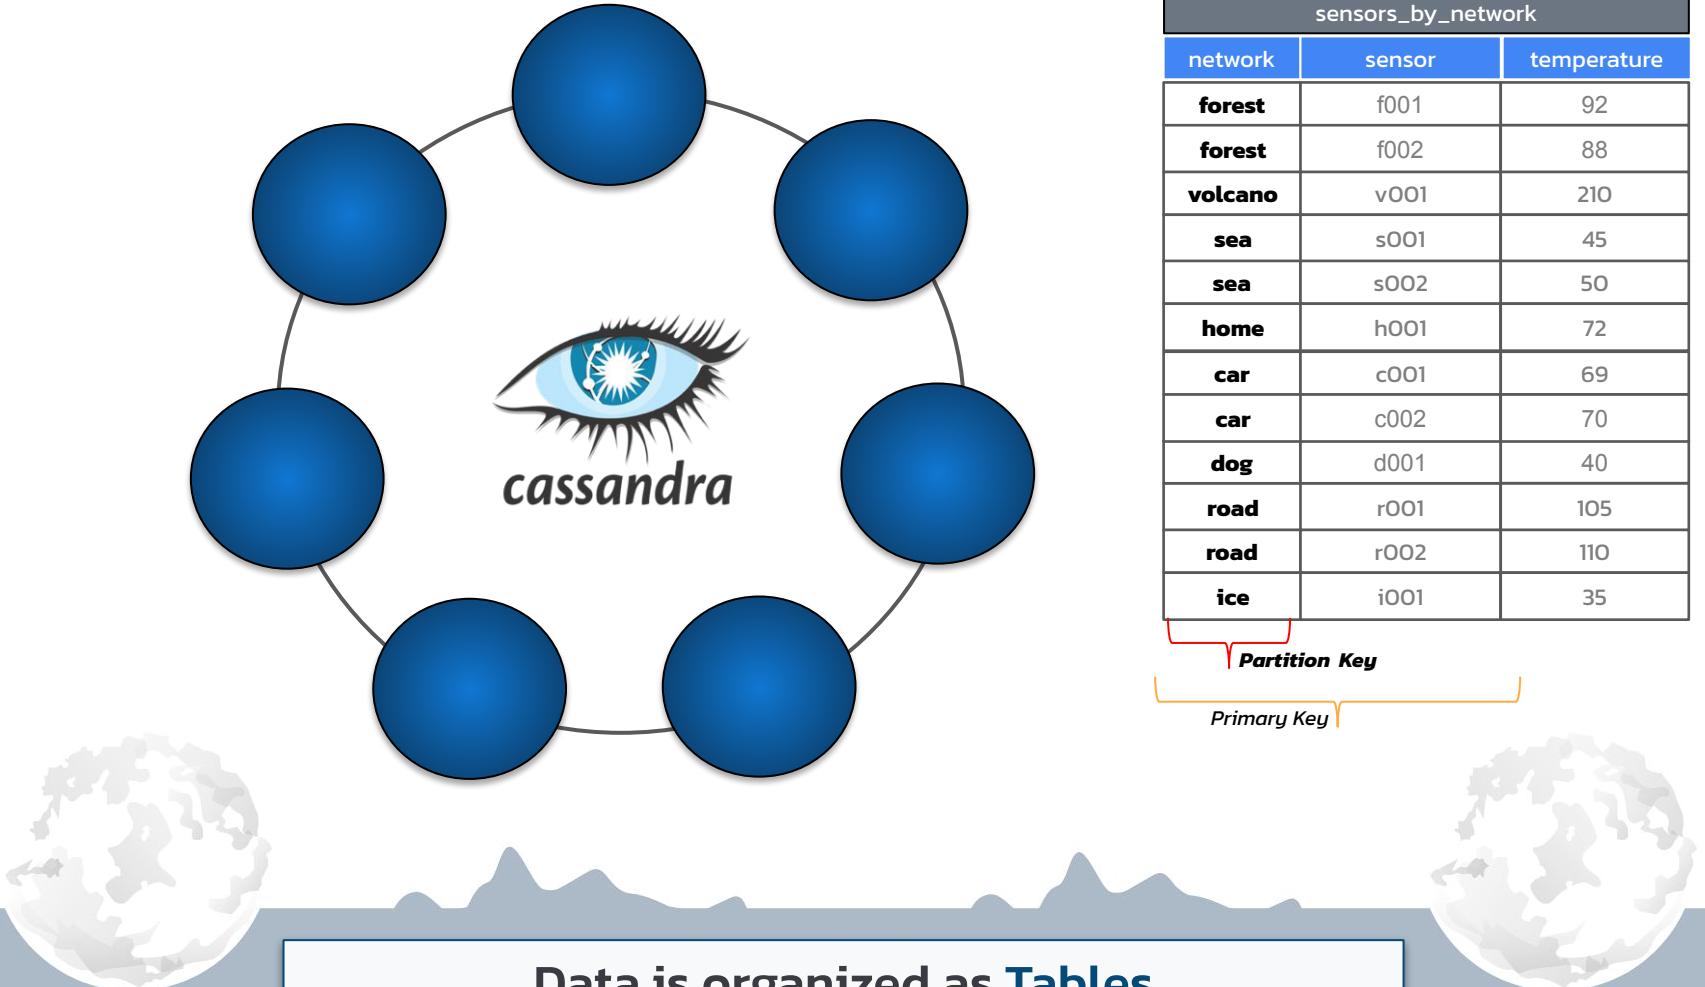
**CONS:** Multiple Writes, Manual Integrity

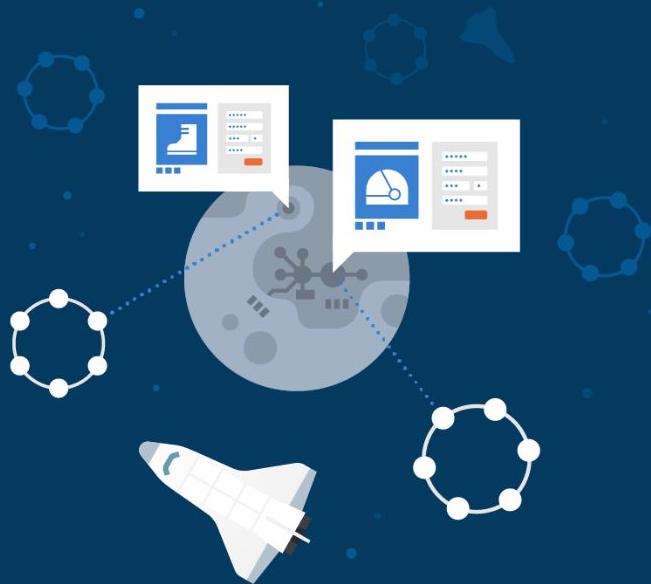
Employees

| userId | firstName | lastName | department  |
|--------|-----------|----------|-------------|
| 1      | Edgar     | Codd     | Engineering |
| 2      | Raymond   | Boyce    | Engineering |
| 3      | Sage      | Lahja    | Math        |
| 4      | Juniper   | Jones    | Botany      |

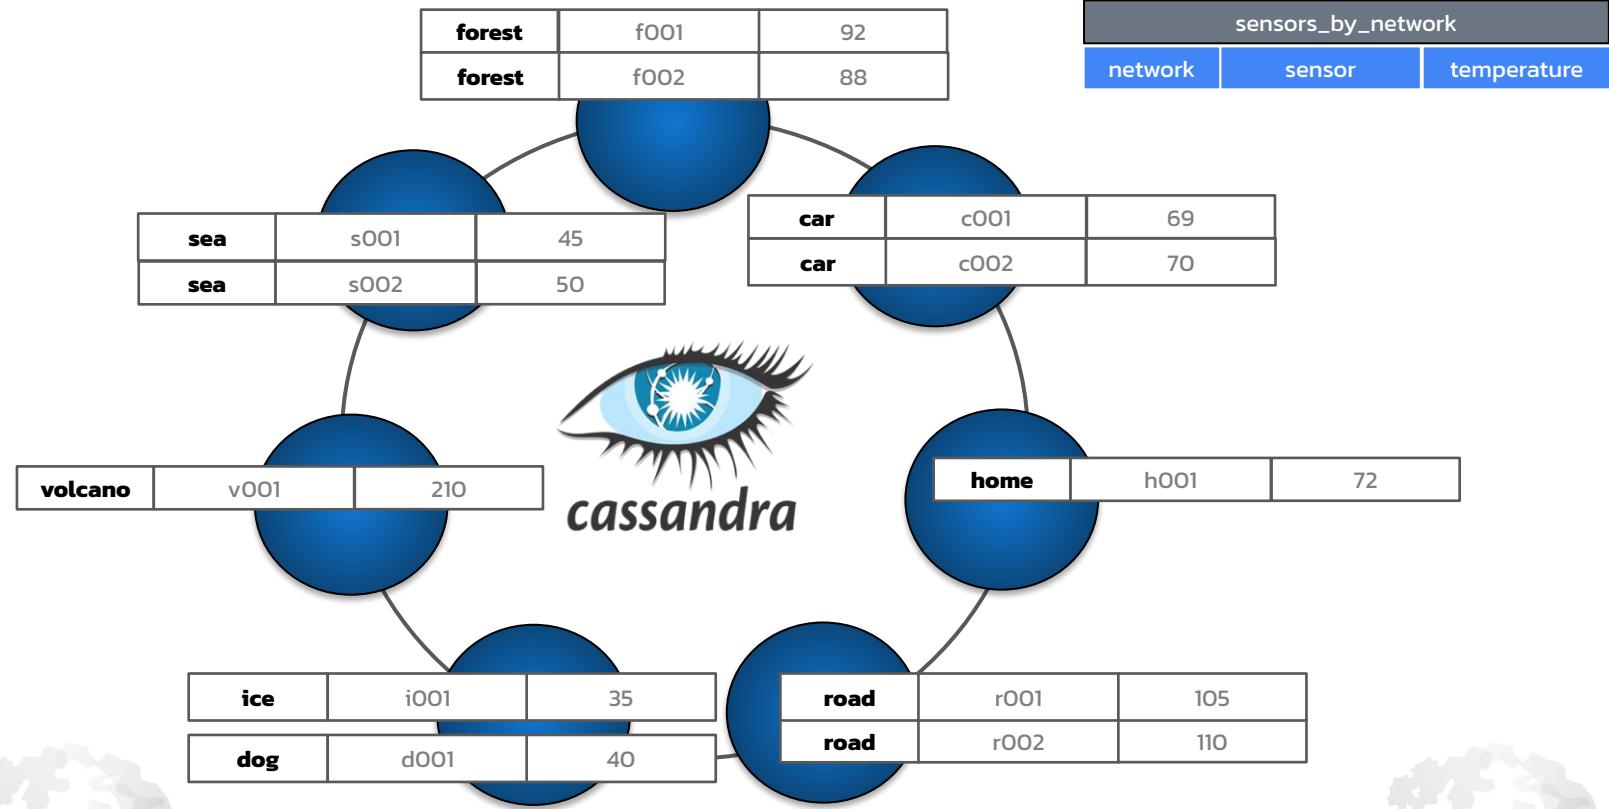
| departmentId | department  |
|--------------|-------------|
| 1            | Engineering |
| 2            | Math        |

Denormalization





## Keys and Partitions



Data is organized as Distributed tables

An identifier for a partition.  
Consists of at least one column,  
may have more if needed.

## PARTITIONS TABLE.

Defines data distribution strategy

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
);
```

Partition key

Examples:

```
PRIMARY KEY ((sensor), timestamp);
```

```
PRIMARY KEY ((sensor), date, timestamp);
```

```
PRIMARY KEY ((sensor, timestamp));
```

Partition Key

Used to **ensure uniqueness** and **sorting order**. Optional.

Defines physical ordering on disk

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Clustering columns

PRIMARY KEY ((**sensor**));

Not Unique

PRIMARY KEY ((**sensor**), timestamp);

Not filter on dates

PRIMARY KEY ((**sensor**), value, timestamp);

Not sorted

PRIMARY KEY ((**sensor**), date, timestamp);



Clustering Columns

An identifier for a row. Consists of at least one Partition Key and zero or more Clustering Columns.

## UNIQUELY IDENTIFIES A ROW

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Primary key

Examples:

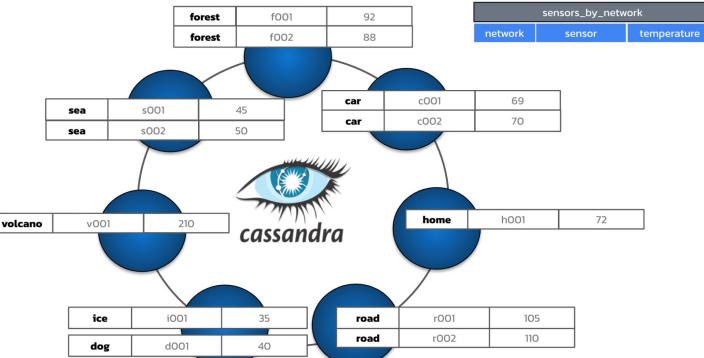
```
PRIMARY KEY ((sensor), timestamp);
```

```
PRIMARY KEY ((sensor, date), timestamp);
```

Primary Key

## Partition Keys:

Defines data distribution over the cluster



## Clustering Columns

Defines how data is physically stored (written on disk)

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

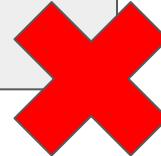
## Important:

Once created, the data model cannot be changed! You will need new tables and migration. Stay lazy, design it right in advance!

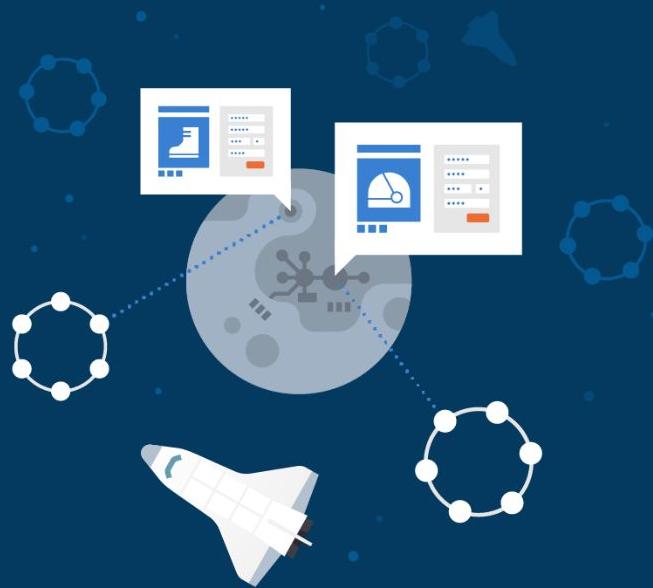
```
ALTER TABLE temperatures_by_sensor  
ADD season TEXT;
```



```
ALTER TABLE temperatures_by_sensor  
DROP PRIMARY KEY  
ADD PRIMARY KEY (sensor, timestamp)
```



Corollary: Schema Immutability

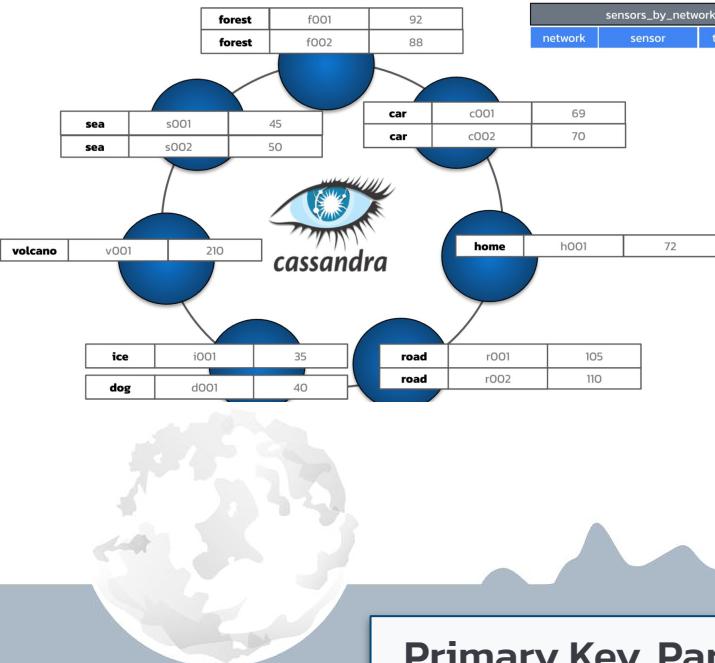


## Keys and Partitions (recap)

```

CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ( ( sensor , date ) , timestamp )
) WITH CLUSTERING ORDER BY (timestamp DESC);

```

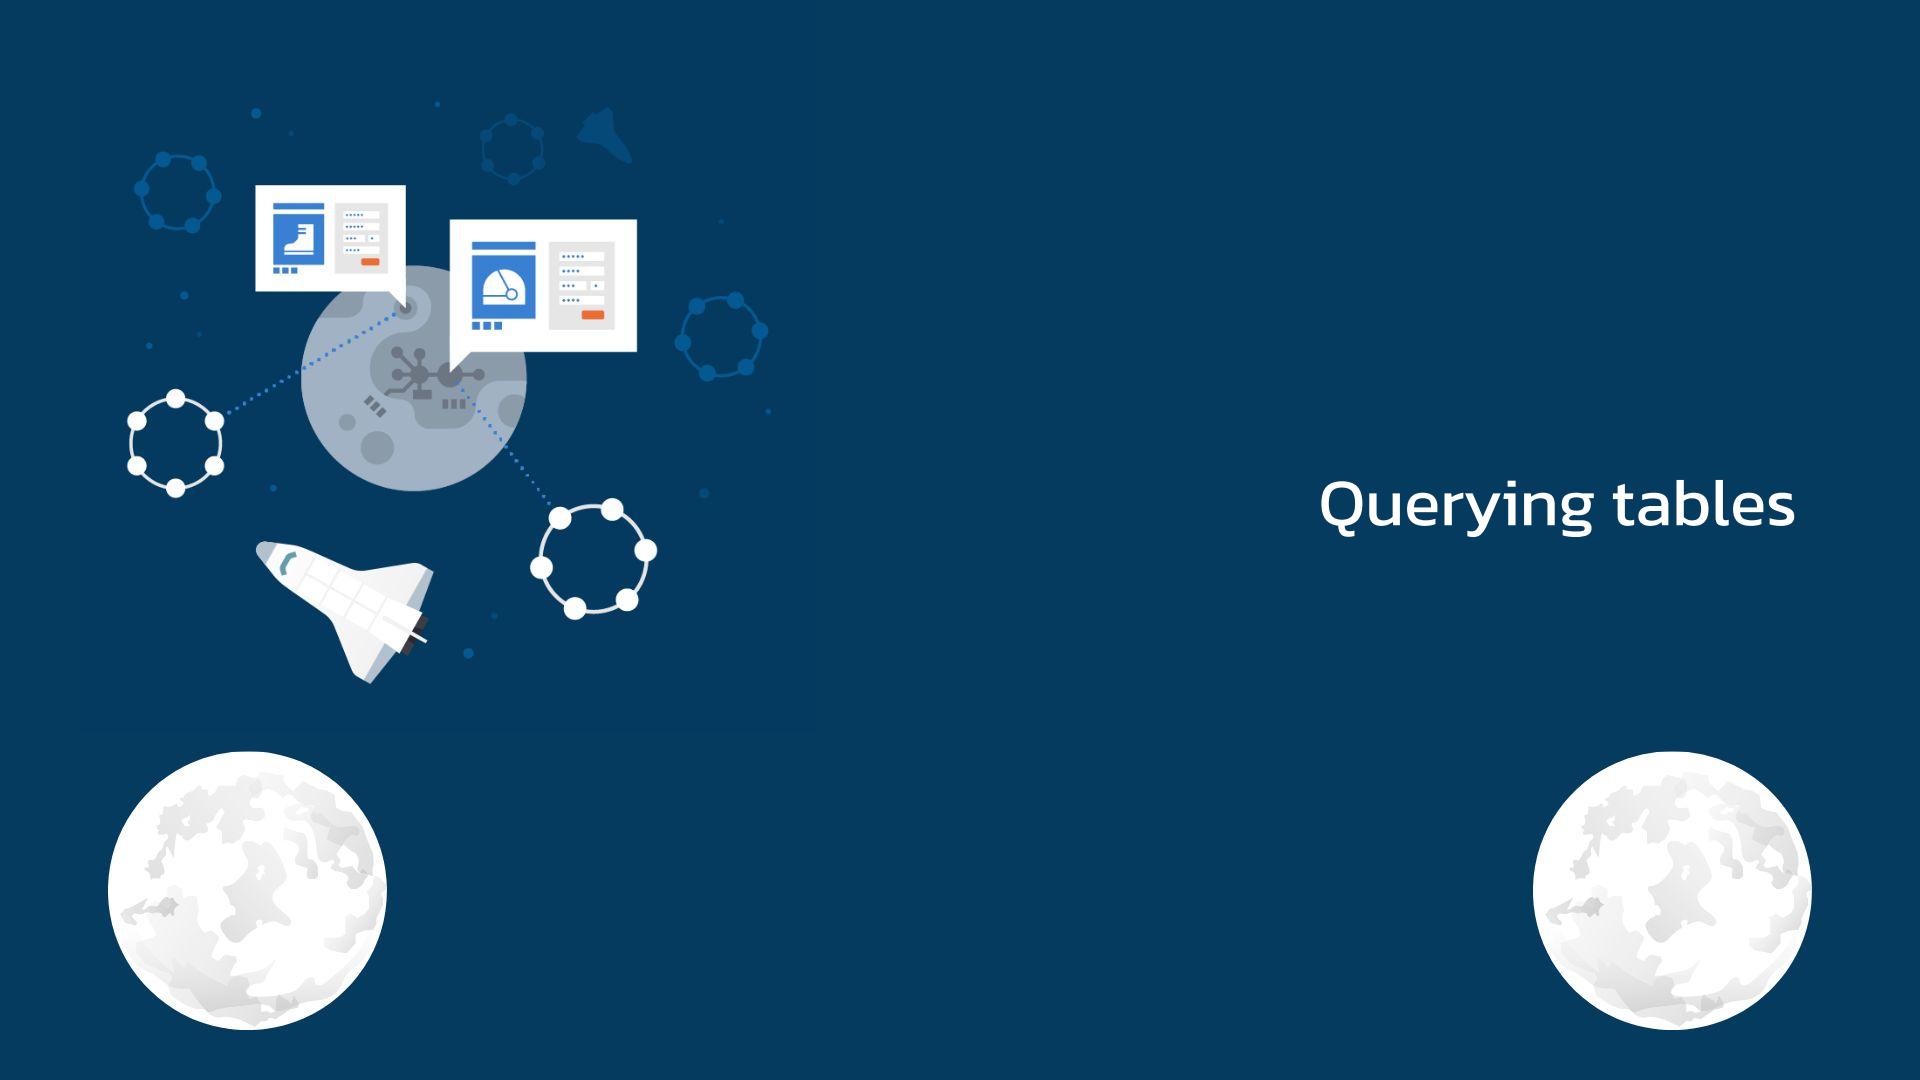


Partition Key  
DEFINES DATA DISTRIB.

Clustering column(s)  
ENSURE UNIQUENESS

Primary key  
UNIQUELY IDENTIFIES A ROW

Primary Key, Partition Key, Clustering Columns, Oh My!



# Querying tables

# "I want to read ALL rows of my table"

Sorry, **not today**.

If you really need to,

⇒ OLAP workload: use Apache Spark & co.

but in a typical OLTP app,  
you actually **don't need** to run such a query!

Reading all rows (*full-table scan*) is a performance capital sin,  
since it involves a lot of nodes in the cluster!

"Rule zero" of querying

```
PRIMARY KEY ((sensor, date), timestamp);
```

```
SELECT * FROM temperatures_by_sensor ...  
  
WHERE sensor = ?;  
  
WHERE sensor > ?;  
  
WHERE sensor = ? AND date > ?;  
  
WHERE sensor = ? AND date = ?;  
  
WHERE sensor = ? AND date = ? AND timestamp > ?;
```

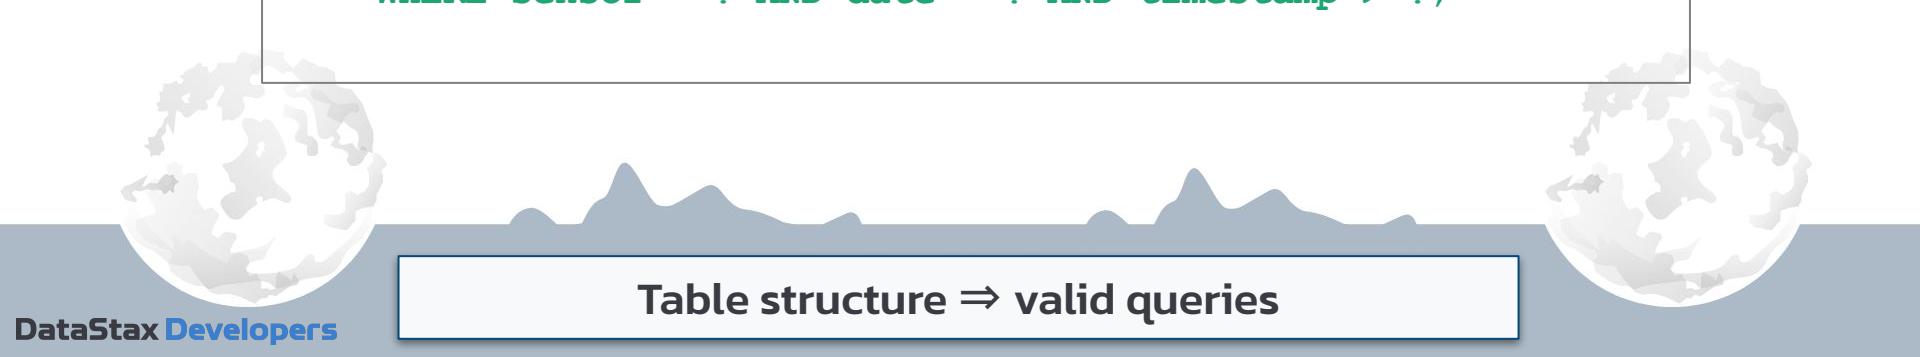
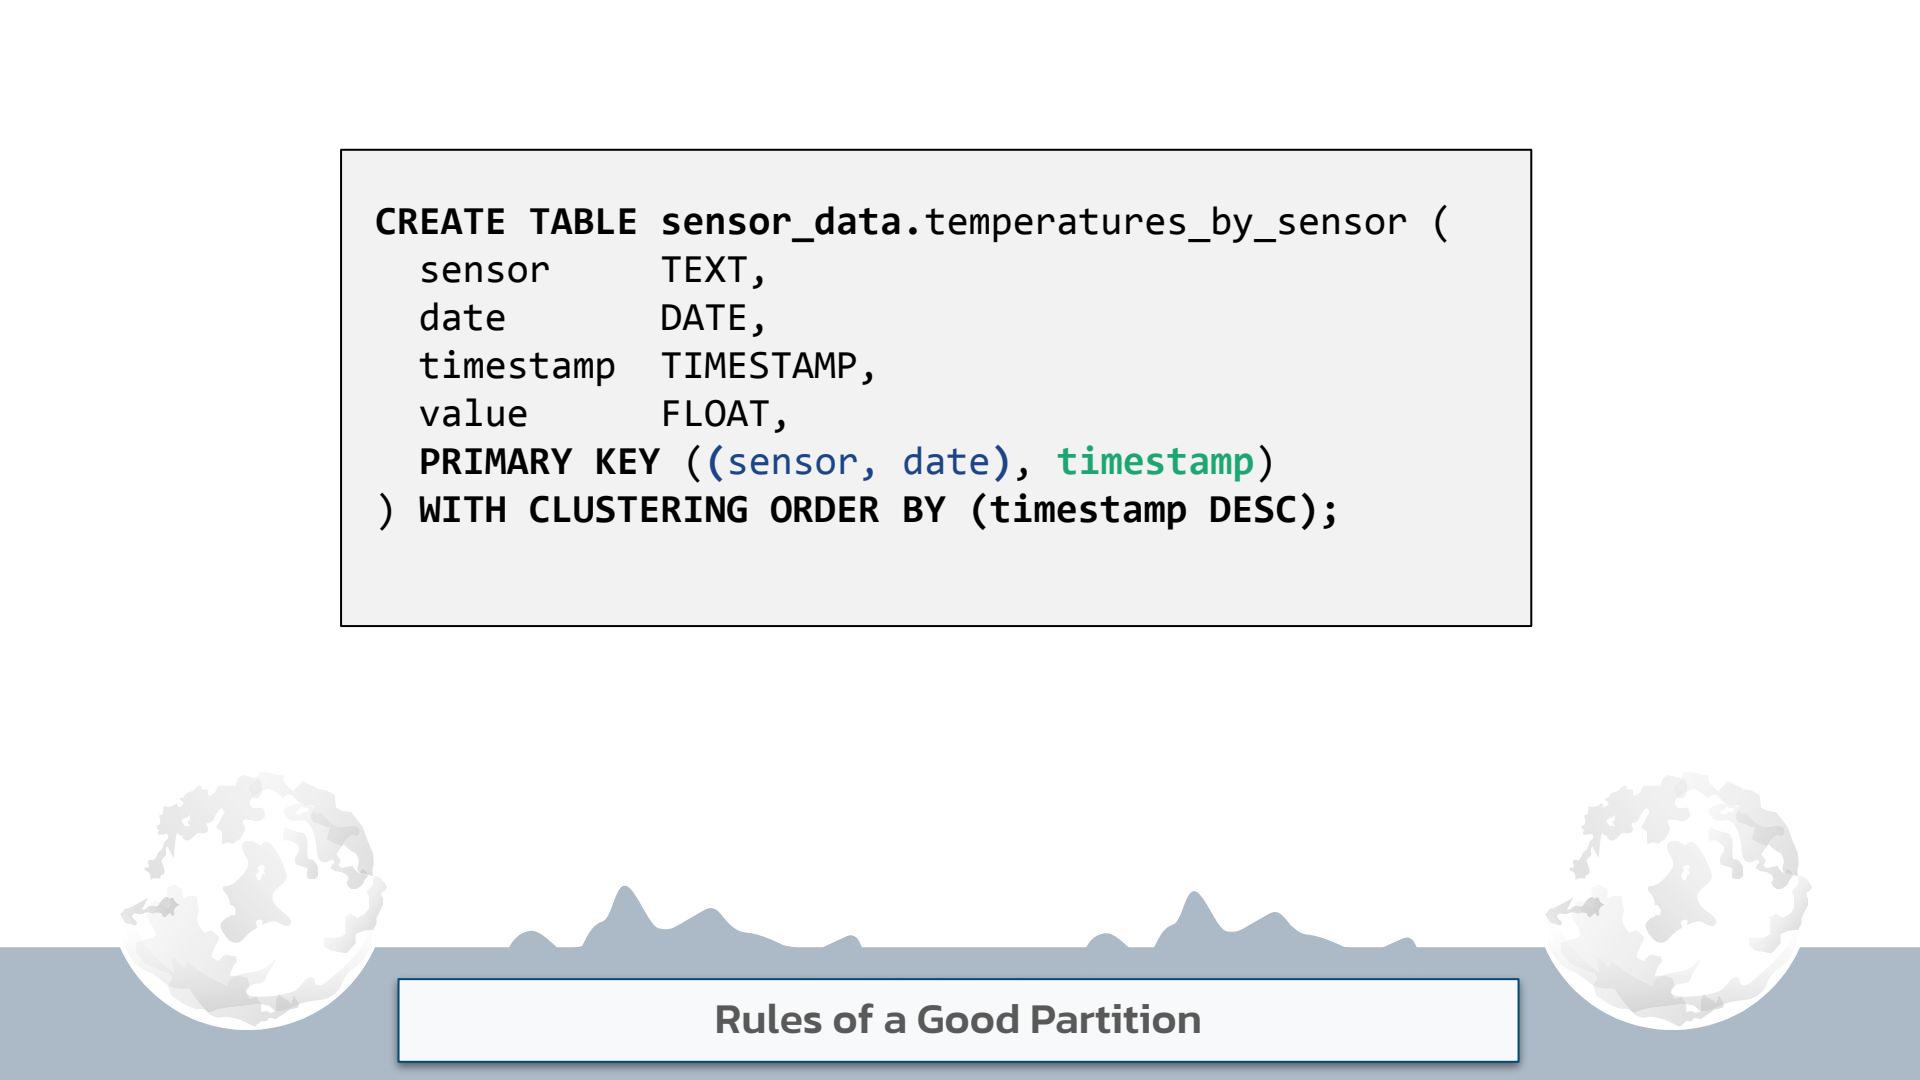


Table structure ⇒ valid queries

```
CREATE TABLE sensor_data.temperatures_by_sensor (
    sensor      TEXT,
    date        DATE,
    timestamp   TIMESTAMP,
    value       FLOAT,
    PRIMARY KEY ((sensor, date), timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```



Rules of a Good Partition

- ❖ **Store together what you retrieve together**
- ❖ Avoid big partitions
- ❖ Avoid hot partitions

**Q:** Show temperature evolution over time for **sensor X** On Sept 20th 2022

```
PRIMARY KEY ((sensor, timestamp));
```



```
PRIMARY KEY ((sensor), timestamp);
```



Rules of a Good Partition

- ❖ Store together what you retrieve together
- ❖ **Avoid big partitions**
- ❖ Avoid hot partitions

# BUCKETING

PRIMARY KEY ((sensor), timestamp);



PRIMARY KEY ((sensor, month), timestamp);



- Up to 2 billion cells per partition
- Up to ~100k values in a partition
- Up to ~100MB in a Partition

Rules of a Good Partition

- ❖ Store together what you retrieve together
- ❖ Avoid big partitions
- ❖ **Avoid hot partitions**

```
PRIMARY KEY ((date), sensor, timestamp);
```



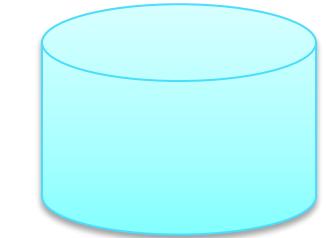
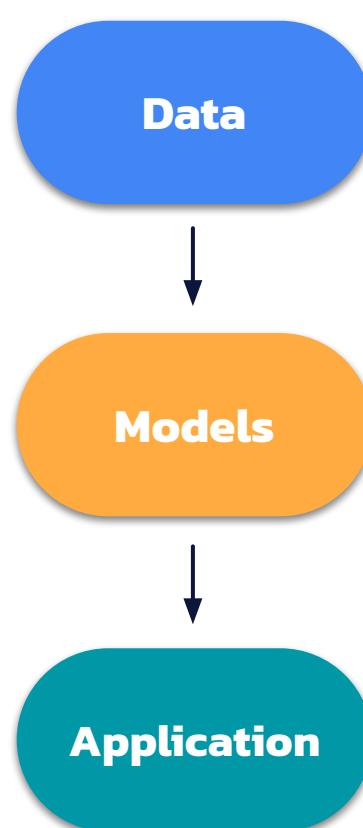
```
PRIMARY KEY ((date, sensor), timestamp);
```





## Data modeling methodology

1. Analyze raw data
2. Identify entities, their properties and relations
3. Design tables, using **normalization** and foreign keys.
4. Use JOIN when doing queries to join normalized data from multiple tables



| Employees |           |          |
|-----------|-----------|----------|
| userId    | firstName | lastName |
| 1         | Edgar     | Codd     |
| 2         | Raymond   | Boyce    |

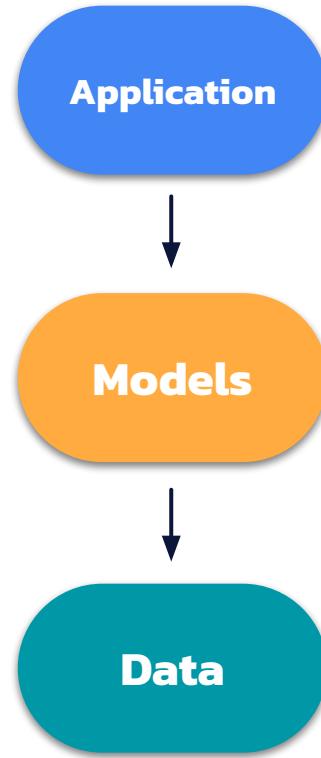
  

| Departments  |             |
|--------------|-------------|
| departmentId | department  |
| 1            | Engineering |
| 2            | Math        |

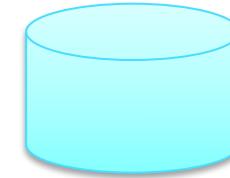


Relational Data Modelling

1. Analyze user behaviour  
(customer first!)
2. Identify workflows, their dependencies  
and needs
3. Define Queries to fulfill these workflows
4. Knowing the queries, design tables,  
using **denormalization**.
5. Use BATCH when inserting or updating  
denormalized data of multiple tables



| Employees |           |          |             |
|-----------|-----------|----------|-------------|
| userId    | firstName | lastName | department  |
| 1         | Edgar     | Codd     | Engineering |
| 2         | Raymond   | Boyce    | Math        |
| 3         | Sage      | Lahja    | Math        |
| 4         | Juniper   | Jones    | Botany      |



- Collection and analysis of **data requirements**
- Identification of participating **entities and relationships**
- Identification of **data access patterns**
- A particular way of **organizing and structuring data**
- Design and specification of a **database schema**
- Schema **optimization** and data **indexing** techniques



Data Quality: completeness consistency accuracy  
Data Access: queryability efficiency scalability



What is Data Modelling ?

## **Modeling principle 1: "Know your data"**

- Key and cardinality constraints are fundamental to schema design

## **Modeling principle 2: "Know your queries"**

- Queries drive schema design

## **Modeling principle 3: "Nest data"**

- Data nesting is the main data modeling technique

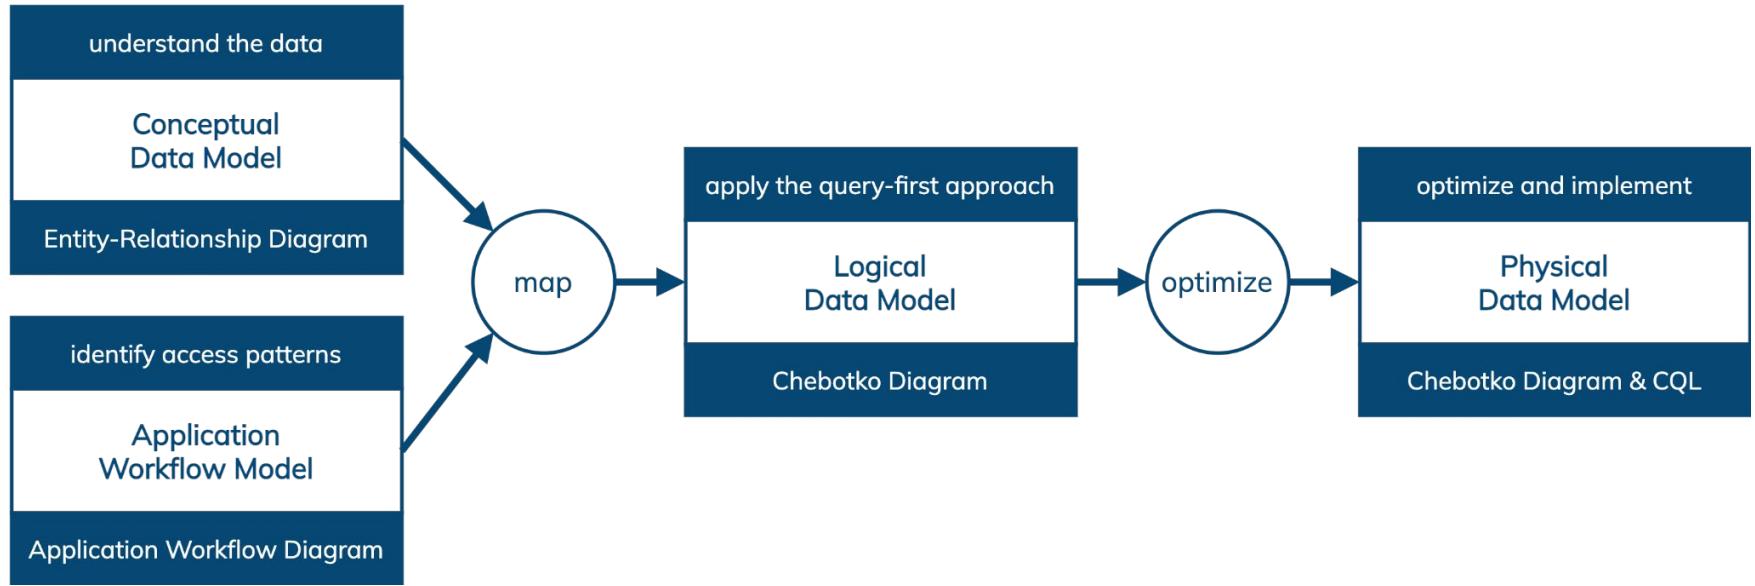
## **Modeling principle 4: "Duplicate data"**

- Better to duplicate than to join



**Cassandra Data Modelling Principles**





Modeling workflow

understand the data

### Conceptual Data Model

Entity-Relationship Diagram

Analyze the Domain

identify access patterns

### Application Workflow Model

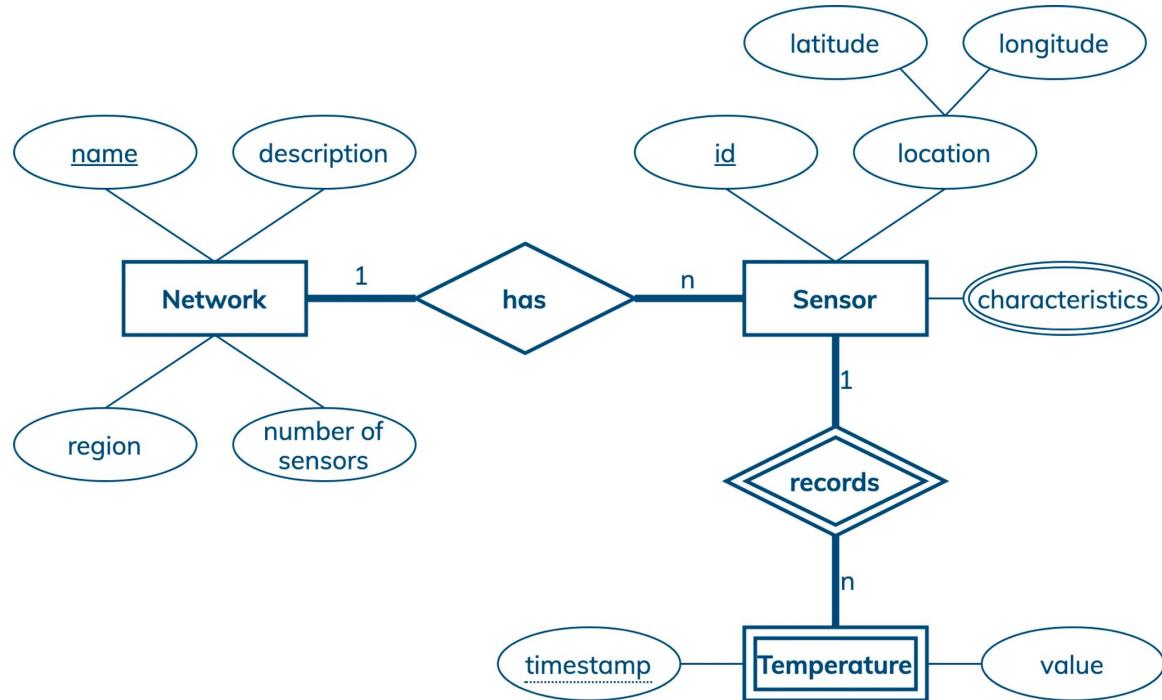
Application Workflow Diagram

Analyze Customer Workflows

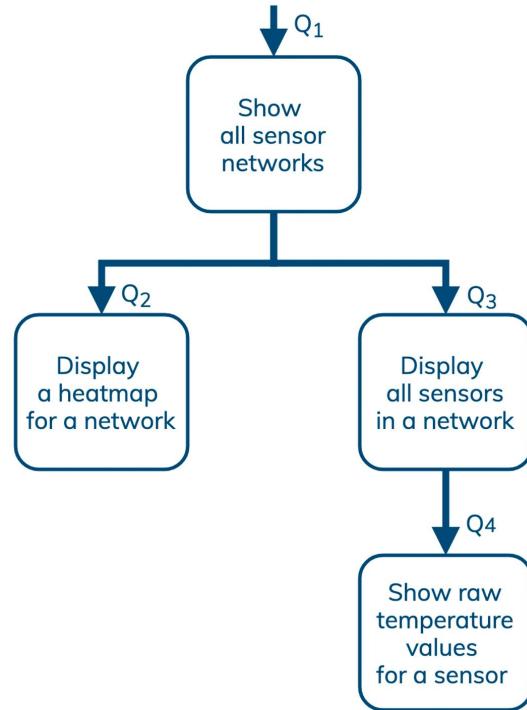


## Data Modelling Methodology: Step I





Entity-Relationship Diagram



#### Data access patterns

- Q1: Find information about all networks; order by name (asc)  
Q2: Find hourly average temperatures for every sensor in a specified network for a given date range; order by date (desc) and hour (desc)  
Q3: Find information about all sensors in a specified network  
Q4: Find raw measurements for a particular sensor on a specified date; order by timestamp (desc)

apply the query-first approach

Logical  
Data Model

Chebotko Diagram

Design queries and build  
tables based on the queries



Data Modelling Methodology: Step II

#### Mapping rule 1: "Entities and relationships"

- Entity and relationship types map to tables

Based on  
a conceptual  
data model

#### Mapping rule 2: "Equality search attributes"

- Equality search attributes map to the beginning columns of a primary key

Based on  
a query

#### Mapping rule 3: "Inequality search attributes"

- Inequality search attributes map to clustering columns

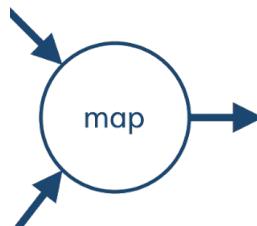
#### Mapping rule 4: "Ordering attributes"

- Ordering attributes map to clustering columns

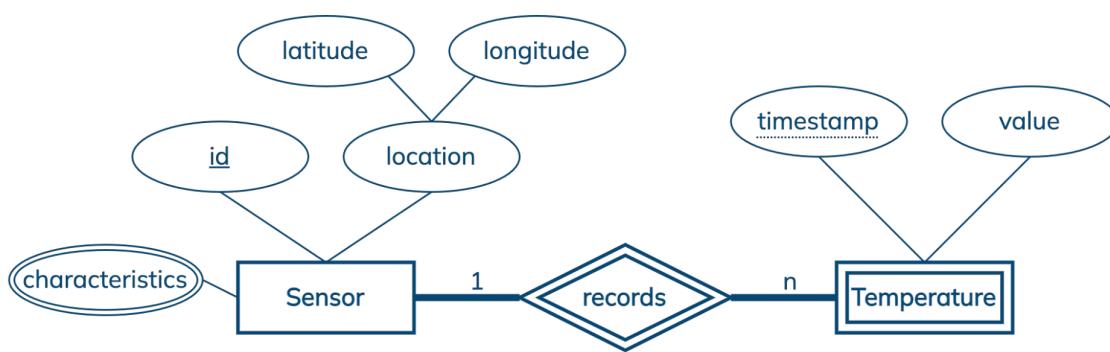
Based on  
a conceptual  
data model

#### Mapping rule 5: "Key attributes"

- Key attributes map to primary key columns



## Mapping Rules

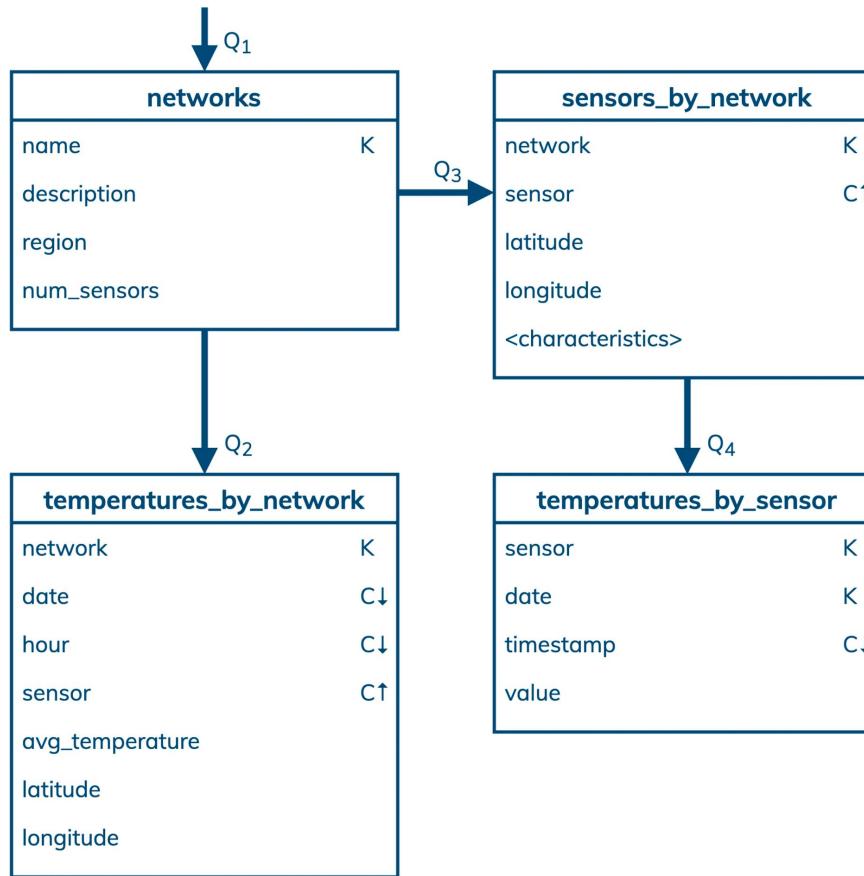


Q5: Find raw measurements for a given location and a date/time range; order by timestamp (desc)

The table shows the results of applying five mapping rules (MR1 to MR5) to the data. The columns represent different mapping configurations, and the rows represent the fields being mapped.

|                 | MR <sub>1</sub>                             | MR <sub>2</sub>                             | MR <sub>3</sub>                             | MR <sub>4</sub>                             | MR <sub>5</sub>                             |
|-----------------|---|---|---|---|---|
| temps_by_sensor | latitude<br>longitude<br>timestamp<br>value | latitude<br>longitude<br>timestamp<br>value | latitude<br>longitude<br>timestamp<br>value | latitude<br>longitude<br>timestamp<br>value | latitude<br>longitude<br>timestamp<br>value |
|                 |   | K   | K   | C↑  | K   |
|                 |   | K   | K   | C↓  | K   |
|                 |   |   |   |   | C↑  |

## Example: Applying Mapping Rules



### Data access patterns

- Q1: Find information about all networks; order by name (asc)
- Q2: Find hourly average temperatures for every sensor in a specified network for a given date range; order by date (desc) and hour (desc)
- Q3: Find information about all sensors in a specified network
- Q4: Find raw measurements for a particular sensor on a specified date; order by timestamp (desc)



Chebotko Diagram

optimize and implement

## Physical Data Model

Chebotko Diagram & CQL

Optimize and Create



Data Modelling Methodology: Step III

- Partition size limits and splitting large partitions
- Data duplication and batches
- Indexes and materialized views\*
- Concurrent data access and lightweight transactions\*
- Dealing with tombstones\*

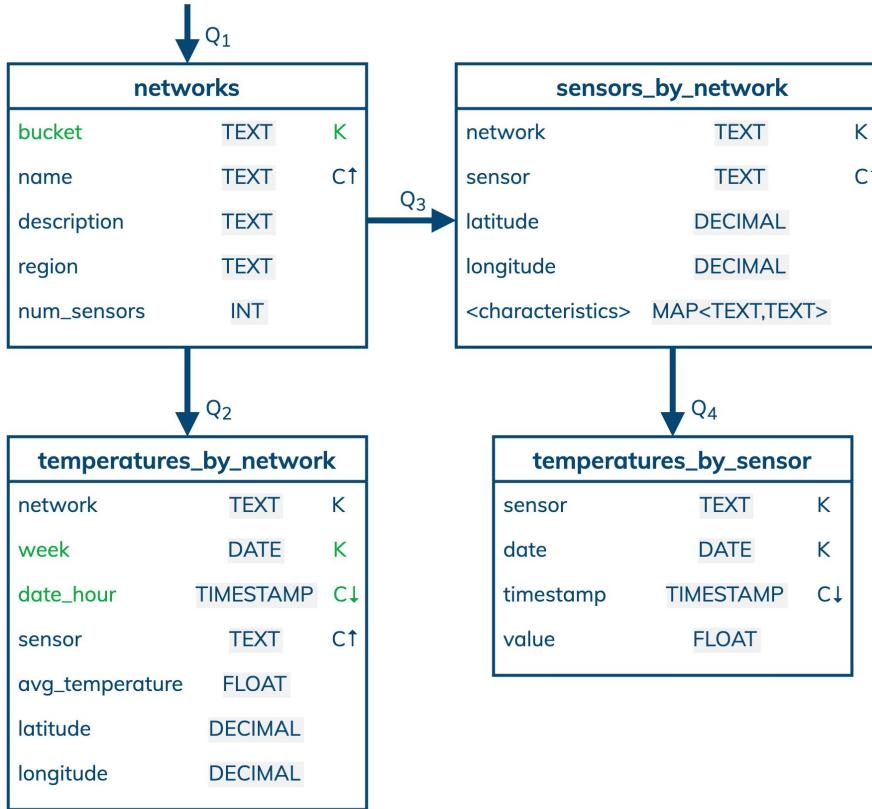


\* Explained in details at our free DS220 course at [academy.datastax.com](http://academy.datastax.com)



## Optimization Techniques





## Physical Data Model: Chebotko Diagram

Q1

```
CREATE TABLE networks (
    bucket TEXT,
    name TEXT,
    description TEXT,
    region TEXT,
    num_sensors INT,
    PRIMARY KEY ((bucket),name)
);
```

Q3

```
CREATE TABLE sensors_by_network (
    network TEXT,
    sensor TEXT,
    latitude DECIMAL,
    longitude DECIMAL,
    characteristics MAP<TEXT,TEXT>,
    PRIMARY KEY ((network),sensor)
);
```

```
CREATE TABLE temperatures_by_network (
    network TEXT,
    week DATE,
    date_hour TIMESTAMP,
    sensor TEXT,
    avg_temperature FLOAT,
    latitude DECIMAL,
    longitude DECIMAL,
    PRIMARY KEY ((network,week),date_hour,sensor)
) WITH CLUSTERING ORDER BY (date_hour DESC, sensor ASC);
```

Q2

Q4

```
CREATE TABLE temperatures_by_sensor (
    sensor TEXT,
    date DATE,
    timestamp TIMESTAMP,
    value FLOAT,
    PRIMARY KEY ((sensor,date),timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

Cassandra Query Language



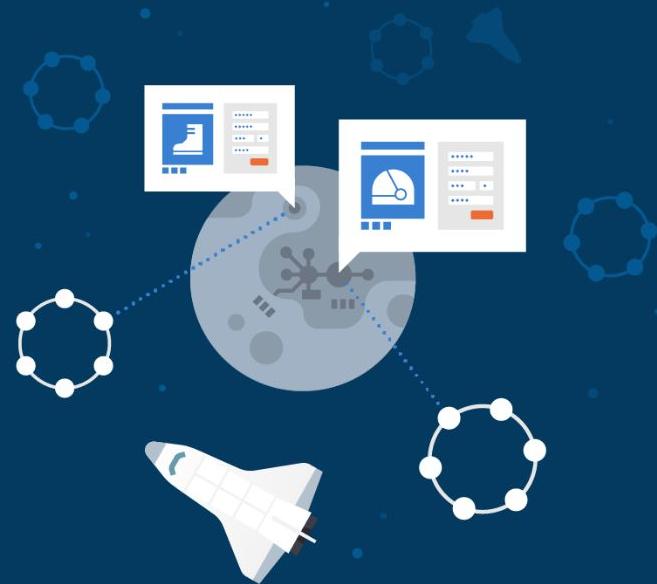
# Lab 1

## Explore a data model

Start here: [dtsx.io/cday-ws2](https://dtsx.io/cday-ws2)



# Apps & Drivers

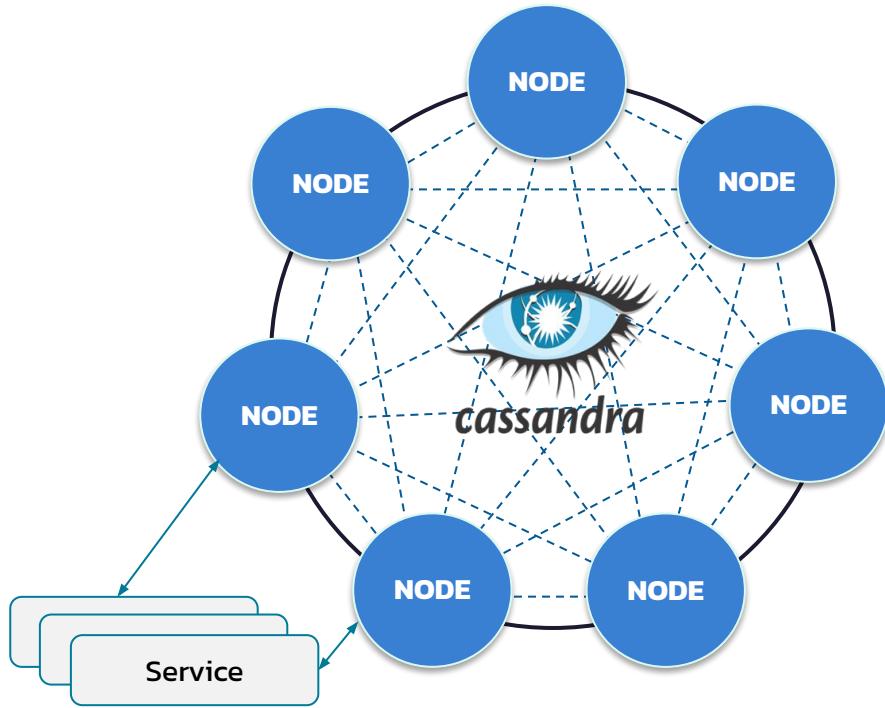


## Loose Coupling: Data resiliency

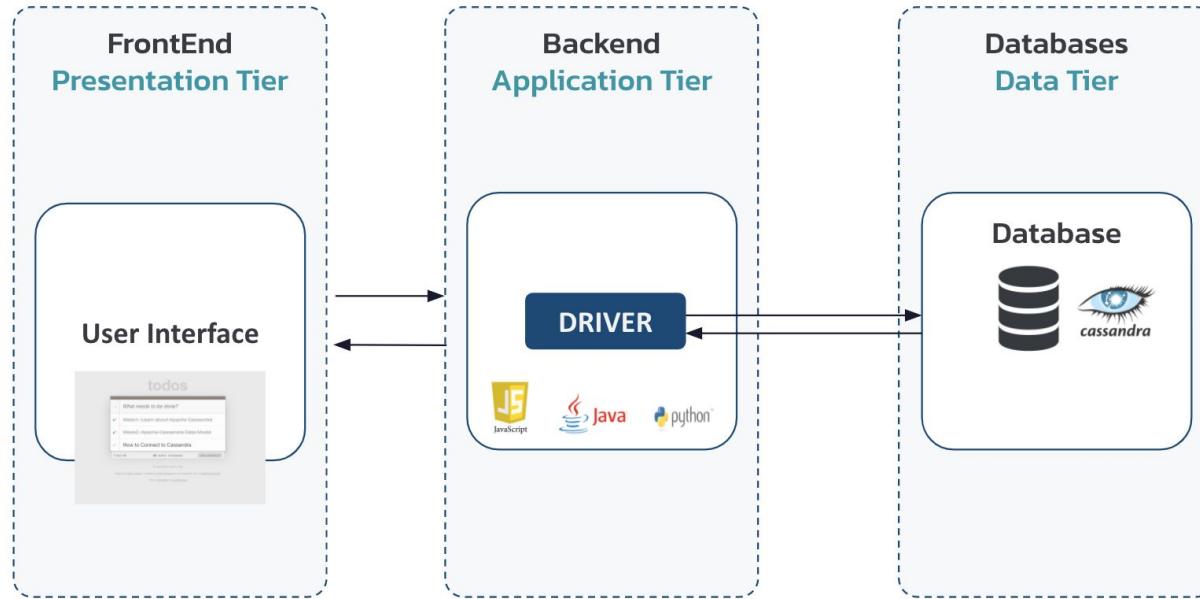
- Data Replicated on multiple Nodes
- Load Balancing at driver side
- Health Check at driver side
- Hinted HandOff

## Shared nothing: Data isolation

- Per Keyspace (with replications)
- Per Tables (1 query = 1 table)
- Per profile (RBAC)



# Application Development with Apache Cassandra



High-level Application Development w/ Cassandra



*today: +Spring Boot*



*today: +FastAPI*



**Choose your language for today**





### Connectivity

- ★ Token & Datacenter Aware
- ★ Load Balancing Policies
- ★ Retry Policies
- ★ Reconnection Policies
- ★ Connection Pooling
- ★ Health Checks
- ★ Authentication | Authorization
- ★ SSL

### Query

- ★ CQL Support
- ★ Schema Management
- ★ Sync/Async/Reactive API
- ★ Query Builder
- ★ Compression
- ★ Paging

### Parsing Results

- ★ Lazy Load
- ★ Object Mapper
- ★ Spring Support
- ★ Paging



Drivers



```
<dependency>
```

```
  <groupId>com.datastax.oss</groupId>  
  <artifactId>java-driver-core</artifactId>  
  <version>4.13.1</version>  
</dependency>
```



```
pip install cassandra-driver
```



```
npm install cassandra-driver
```

4.6.3



```
{  
  "dependencies": {  
    "cassandra-driver": "^4.6.3"  
  }  
}
```

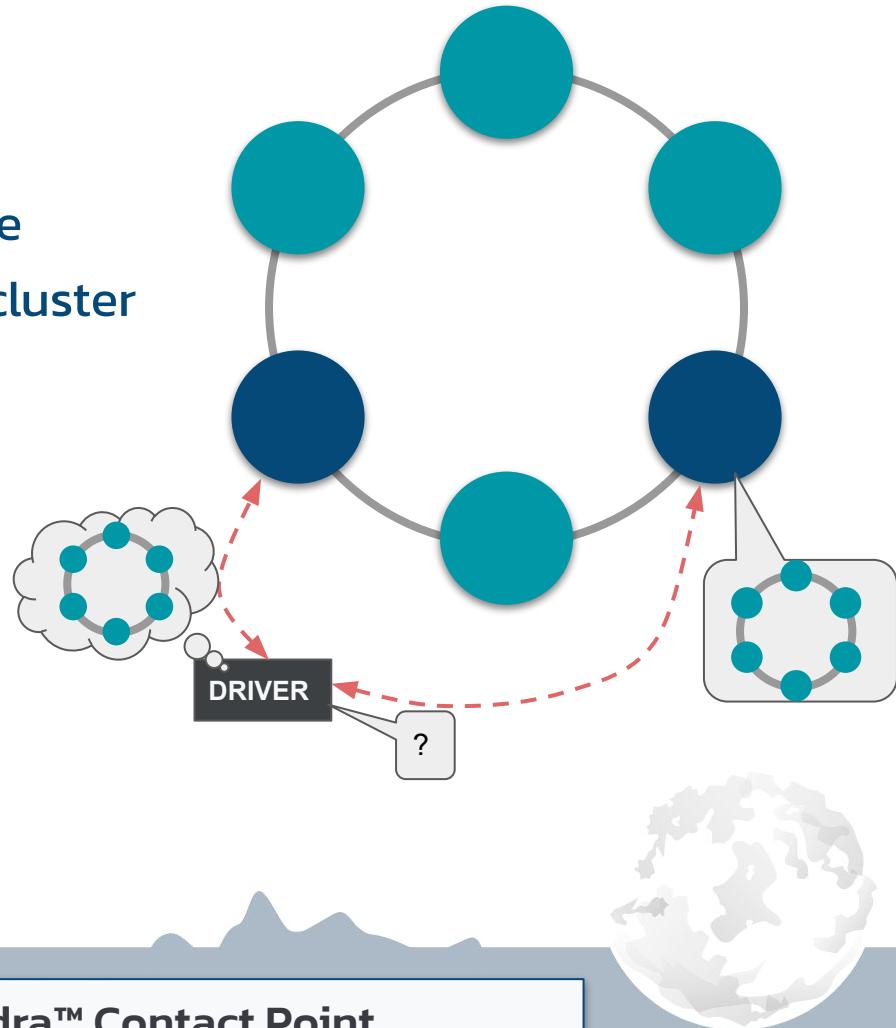


nuget v3.15.0

```
Install-Package CassandraCSharpDriver -Version 3.15.0
```



- Only one necessary
- Unless that node is down
- Better ~3 nodes per DC for resilience
- From there, drivers discover whole cluster
- Local Datacenter



Apache Cassandra™ Contact Point

```
CqlSession cqlSession = CqlSession.builder()  
    .addContactPoint(new InetSocketAddress("127.0.0.1", 9042))  
    .withKeyspace("killrvideo")  
    .withLocalDatacenter("dc1")  
    .withAuthCredentials("U", "P")  
    .build();
```



```
auth_provider = PlainTextAuthProvider(  
    username='U', password='P')  
  
cluster = Cluster(['127.0.0.1'],  
    auth_provider=auth_provider, protocol_version=2)  
  
session = cluster.connect('todos')
```



```
const client = new cassandra.Client({  
    contactPoints: ['127.0.0.1'],  
    localDataCenter: 'dc1',  
    keyspace: 'killrvideo',  
    credentials: { username: 'U', password: 'P' }  
});
```



```
Cluster cluster = Cluster.Builder()  
    .AddContactPoint("127.0.0.1")  
    .WithCredentials("U", "P")  
    .Build();  
  
session = cluster.Connect("todos");
```



```
CqlSession cqlSession = CqlSession.builder()  
    .withCloudSecureConnectBundle(Paths.get("secure.zip"))  
    .withAuthCredentials("U", "P")  
    .withKeyspace("todos")  
    .build();
```



```
auth_provider = PlainTextAuthProvider(  
    username='U', password='P')  
  
cluster = Cluster(  
    cloud ={ 'secure_connect_bundle': 'secure.zip'},  
    auth_provider=auth_provider)  
  
session= cluster.connect('todos')
```



```
const client = new cassandra.Client({  
  cloud: { secureConnectBundle: 'secure.zip' },  
  credentials: { username: 'u', password: 'p' }  
});
```



```
var cluster = Cluster.Builder()  
    .WithCloudSecureConnectionBundle("secure.zip")  
    .WithCredentials("u", "p")  
    .Build();  
  
var session = cluster.Connect("todos");
```



- **Stateful** object handling communications with each node
- Should be unique in the Application (*Singleton*)
- Should be closed at application shutdown (*shutdown hook*) in order to free opened TCP sockets (*stateful*)

```
Java:      cqlSession.close();
```

```
Python:     session.shutdown();
```

```
Node:      client.shutdown();
```

```
CSharp:    IDisposable
```



# Executing CQL Queries

```
session.execute(     python™  
    "SELECT * FROM sensors_by_network WHERE network = %s;",  
    (network,)  
)
```

```
cqlSession.execute(     Java  
    "SELECT * FROM sensors_by_network WHERE network = '" + network + "'"  
);
```

# Prepared Statements



```
q3_statement = session.prepare(  
    "SELECT * FROM sensors_by_network WHERE network = ?;"  
)  
rows = session.execute(q3_statement, (network,) )
```



Prepared Statements in Python

# Prepared Statements



```
PreparedStatement q3Prepared = session.prepare(  
    "SELECT * FROM sensors_by_network WHERE network = ?");  
BoundStatement q3Bound = q3Prepared.bind(network);  
ResultSet rs = session.execute(q3Bound);
```



Prepared Statements in Java

# Advantages of CQL Prepared Statements

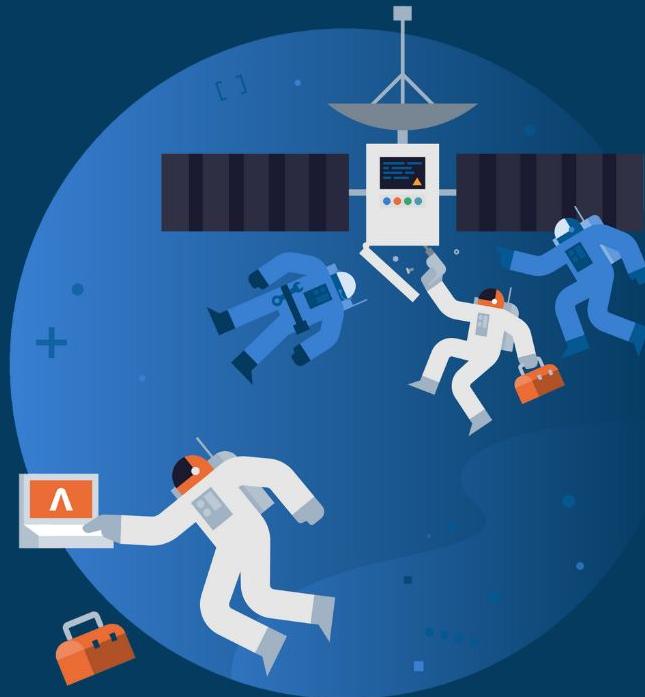


- Parse once, run many times.
- Saves network trips for result set metadata.
- Client-side type validation.
- Statements binding on partition keys compute their own cluster routing.



Advantages of Prepared Statements





# Lab 2

## Create DB and run API !

Start here: [dtsx.io/cday-ws2](https://dtsx.io/cday-ws2)



## Conclusion & next steps

# Homework (see App Dev repo!)



A couple of "theory" questions, plus ...

**Coding exercise:** *Enrich the API with a new  
GET endpoint for Q1 ("get all networks")*



DataStax Developers



## Get help about Cassandra

 stack**overflow**<sup>(\*)</sup>:

[stackoverflow.com/questions/tagged/cassandra](https://stackoverflow.com/questions/tagged/cassandra)



DBA Stack Exchange<sup>(\*)</sup>:

[dba.stackexchange.com/questions/tagged/cassandra](https://dba.stackexchange.com/questions/tagged/cassandra)



Discord:

[dtsx.io/discord](https://dtsx.io/discord)

(\*) for best results, follow the "cassandra" tag

# This is just the beginning of your journey



Roads? Where we're going we don't need roads

Enjoy the rest of today  
Do the homework ⇒ get a badge  
We have more workshops & hands-on labs!  
Get Cassandra Certified (Academy & Exam voucher)





# Thank you!



**Sponsored by DataStax**